

Pipelined CPU Design

Now, we will optimize a single cycle CPU using pipelining. Pipelining is a powerful logic design method to reduce the clock time and improve the throughput, even though it increases the latency of an individual task and adds additional logic. In a pipelined CPU, multiple instructions are overlapped in execution. This is a good example of parallelism, which is one of the great ideas in computer architecture. To obtain a pipelined CPU, we will take the following steps.

Step 1: Pipeline Registers

Pipelining starts from adding pipelining registers by dividing a large combinational logic. We have already chopped a single cycle CPU into five stages, and thus, will add pipeline registers between two stages.

Step 2: Performance Analysis

A great advantage of pipelining is the performance improvement with a shorter clock time. We will use the same timing parameters as those in the previous discussion.

| Element | Register clk-to-q | Register Setup | MUX | ALU | Mem Read | Mem Write | RegFile Read | RegFile Setup |
|-----------|-----------------------|--------------------|------------------|------------------|----------------------|-----------------------|---------------------|----------------------|
| Parameter | $t_{\text{clk-to-q}}$ | t_{setup} | t_{mux} | t_{ALU} | t_{MEMread} | t_{MEMwrite} | t_{RFread} | t_{RFsetup} |
| Delay(ps) | 30 | 20 | 25 | 200 | 250 | 200 | 150 | 20 |

Q1. What was the clock time and frequency of a single cycle CPU?

$$t_{\text{clk,single}} \geq t_{\text{PC, clk-to-q}} + t_{\text{MEMread}} + t_{\text{RFread}} + t_{\text{ALU}} + t_{\text{DMEMread}} + t_{\text{mux}} + t_{\text{RFsetup}}$$

$$= 30 + 250 + 150 + 200 + 250 + 25 + 20 = 925 \text{ ps}$$

$$f_{\text{clk,single}} = 1/t_{\text{clk,pipe}} \leq 1/(925 \text{ ps}) = 1.08 \text{ GHz}$$

Q2. What is the clock time and frequency of a pipelined CPU?

$$t_{\text{clkpipe}} \geq \max \left(\begin{array}{l} t_{\text{clk-to-q}} + t_{\text{MEMread}} + t_{\text{setup}} \text{ (Fetch)} \\ t_{\text{clk-to-q}} + t_{\text{RFread}} + t_{\text{setup}} \text{ (Decode)} \\ t_{\text{clk-to-q}} + t_{\text{ALU}} + t_{\text{mux}} + t_{\text{setup}} \text{ (Execute)} \\ t_{\text{clk-to-q}} + t_{\text{MEMread}} + t_{\text{setup}} \text{ (Memory)} \\ t_{\text{clk-to-q}} + t_{\text{mux}} + t_{\text{RFsetup}} \text{ (Writeback)} \end{array} \right) = 300\text{ps}$$

$$f_{\text{clk,pipe}} = 1/t_{\text{clk,pipe}} \leq 1/(300 \text{ ps}) = 3.33 \text{ GHz}$$

Q3. What is the speed-up? Why is it less than five?

$$\text{Speed-up} = t_{\text{clk,pipe}} / t_{\text{clk,single}} = f_{\text{clk,pipe}} / f_{\text{clk,single}} = 3.08.$$

This is because pipeline stages are not balanced evenly and there is overhead from pipeline registers ($t_{\text{clk-to-q}}$, t_{setup}). Moreover, this does not include the delays from the additional logic for hazard resolution.

Step 3: Pipeline Hazard

The performance improvement comes at a cost. Pipelining introduces pipeline hazards we have to overcome.

Structural Hazard

Structural hazards occur when more than one instruction use the same resource at the same time.

- **Register File:** One instruction reads from the register file while another writes to it. We can solve this by having separate read and write ports and writing to the register file at the falling edge of the clock.
- **Memory:** The memory is accessed not only for the instruction but also for the data. Separate caches for instructions and data solve this hazard.

Data Hazard and Forwarding

Data hazards occur due to data dependencies among instructions. Forwarding can solve many data hazards.

Q1. Spot the data dependencies in the code below and figure out how forwarding can resolve data hazards.

| Instruction | C0 | C1 | C2 | C3 | C4 | C5 | C6 |
|----------------------|----|-----|-----|-----|-----|-----|----|
| addi \$t0, \$s0, -1 | IF | REG | EX | MEM | WB | | |
| and \$s2, \$t0, \$a0 | | IF | REG | EX | MEM | WB | |
| sw \$s0, 100(\$t0) | | | IF | REG | EX | MEM | WB |

The REG step for instructions 2 and 3 depend on data in the registers only available after the WB step of instruction 1. We can forward the ALU output of the first instruction to the EX stages of future instructions

Q2. In general, under what conditions will an EX stage need to take in forwarded inputs from previous instructions? Where should those inputs come from in regards to the current cycle? Assume you have the signals ALUout(n), rt(n), rs(n), regWrite(n), and regDst(n), where n is 0 for the signal of the current instruction being executed by the EX stage, -1 for the previous, etc.

Forward ALUout(-1) if (rt(0) == regDst(-1) || rs(0) == regDst(-1)) && regWrite(-1)

Forward ALUout(-2) if (rt(0) == regDst(-2) || rs(0) == regDst(-2)) && regWrite(-2)

Forward ALUout(-3) if (rt(0) == regDst(-3) || rs(0) == regDst(-3)) && regWrite(-3)

Data Hazard and Stall

Forwarding cannot solve all data hazards. We need to stall the pipeline in some cases.

Q1. Spot the data dependencies in the code below and figure out why forwarding cannot resolve this hazard.

| Instruction | C0 | C1 | C2 | C3 | C4 | C5 |
|------------------------|----|-----|-----|-----|-----|----|
| lw \$t0, 20(\$s0) | IF | REG | EX | MEM | WB | |
| addiu \$t1, \$t0, \$t0 | | IF | REG | EX | MEM | WB |

The add instruction needs the value of \$t0 in the beginning of C3, but it is ready at the end of C3.

Q2. Now we stall the pipeline one cycle and insert nop after the lw instruction. Figure out how this can resolve the hazard.

| Instruction | C0 | C1 | C2 | C3 | C4 | C5 | C6 |
|------------------------|----|-----|-----|-----|-----|-----|----|
| lw \$t0, 20(\$s0) | IF | REG | EX | MEM | WB | | |
| nop | | IF | REG | EX | MEM | WB | |
| addiu \$t1, \$t0, \$t0 | | | IF | REG | EX | MEM | WB |

By stalling one cycle, the add instruction can start its execution stage after the \$t0 value is ready.

Q3. Under what conditions do we need to introduce a nop? Under what conditions do we need to forward the output of the MEM stage to the EX stage? Assume you have the signals memToReg(n), rt(n), rs(n), regWrite(n), and regDst(n), where n is 0 for the signal of the current instruction being executed by the EX stage, -1 for the previous, etc.

We forward if $(rt(0) == regDst(-2) \mid \mid rs(0) == regDst(-2)) \&\& memToReg(-2) \&\& regWrite(-2)$

Control Hazard and Prediction

Control hazards occur due to jumps and branches. We may solve them by stalling the pipeline. However, it is painful since the branch condition is calculated after the execution stage and the pipeline is stalled for three cycles. Instead, we add a branch comparator inside the register read stage and introduce the branch delay slot, and redefine MIPS so that the instruction after a branch statement will always be executed.

Q1. Reorder the following sets of instructions to account for the branch delay slot. You may have to insert a nop instruction.

| Set 1 | Reordered set 1 | Set 2 | Reordered Set 2 |
|-----------------------|-----------------------|-----------------------|-----------------------|
| addiu \$t0, \$t1, 5 | addiu \$t0, \$t1, 5 | addiu \$t0, \$t1, 5 | addiu \$t0, \$t1, 5 |
| ori \$t2, \$t3, 0xff | beq \$t0, \$s0, label | ori \$t2, \$t3, 0xff | ori \$t2, \$t3, 0xff |
| beq \$t0, \$s0, label | ori \$t2, \$t3, 0xff | beq \$t0, \$t2, label | beq \$t0, \$t2, label |
| lw \$t4, 0(\$t0) | lw \$t4, 0(\$t0) | lw \$t4, 0(\$t0) | nop |
| | | | lw \$t4, 0(\$t0) |

