

CS 61C: Great Ideas in Computer Architecture

Lecture 13: *Pipelining*

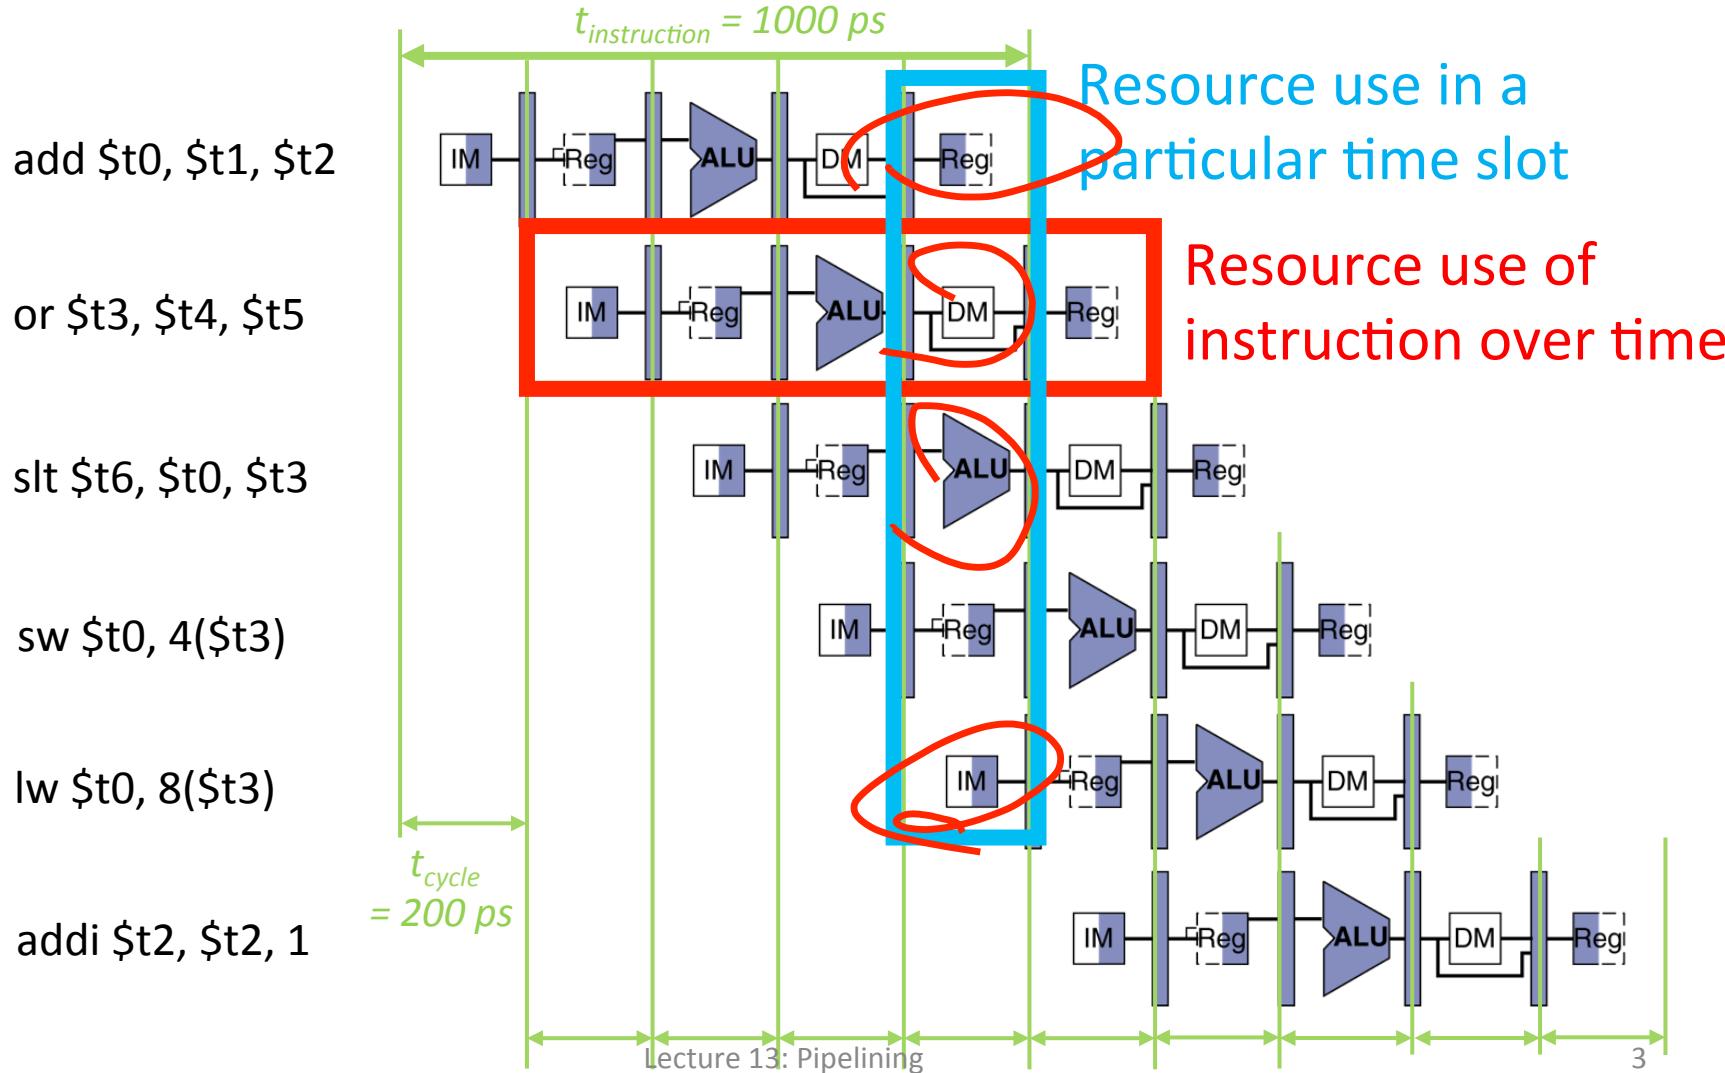
Bernhard Boser & Randy Katz

<http://inst.eecs.berkeley.edu/~cs61c>

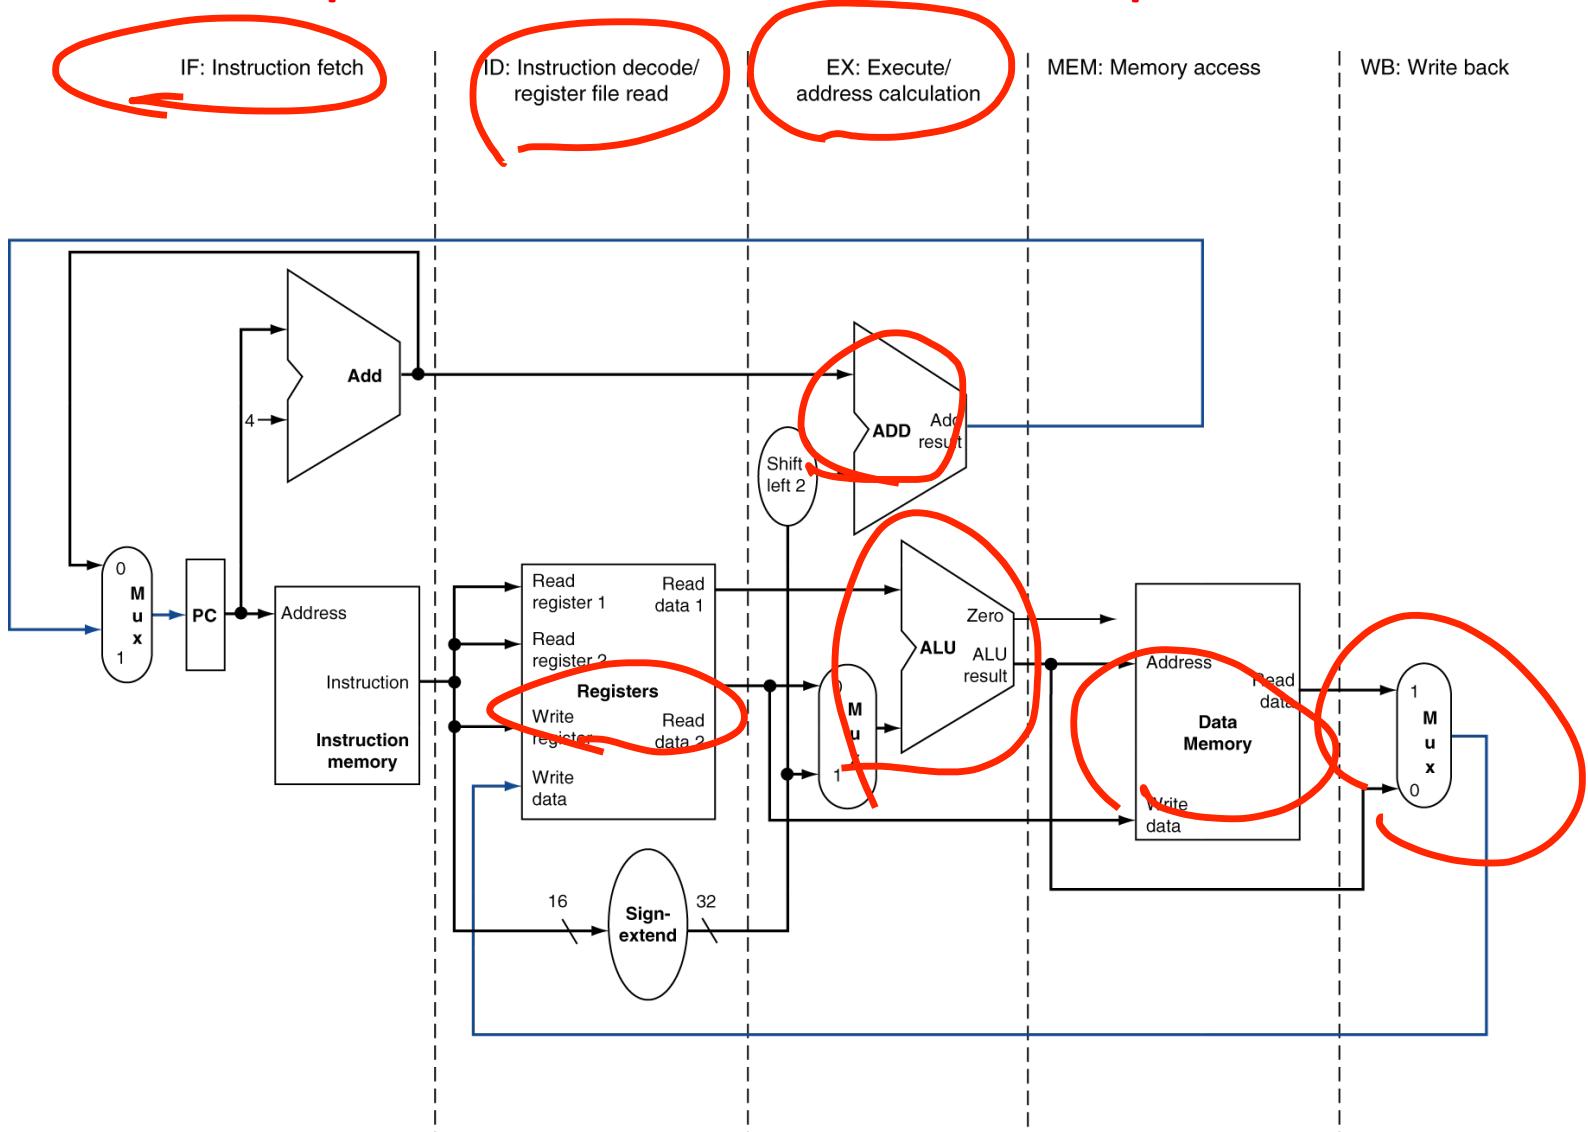
Agenda

- MIPS Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

MIPS Pipeline

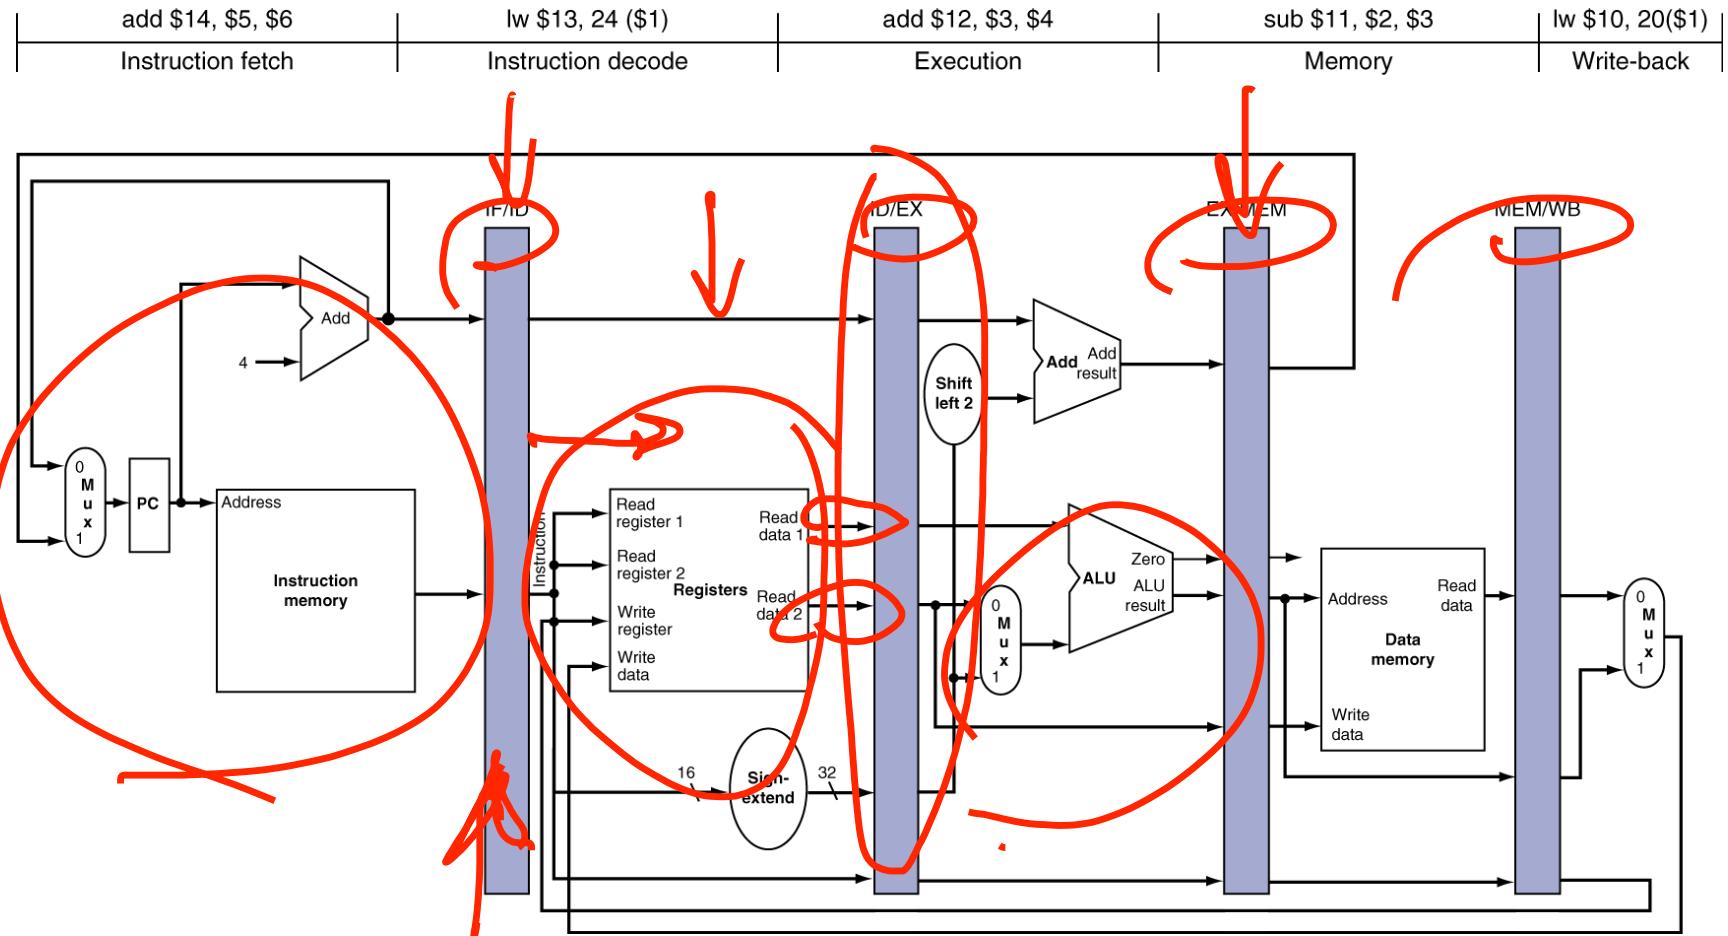


Simplified MIPS Datapath

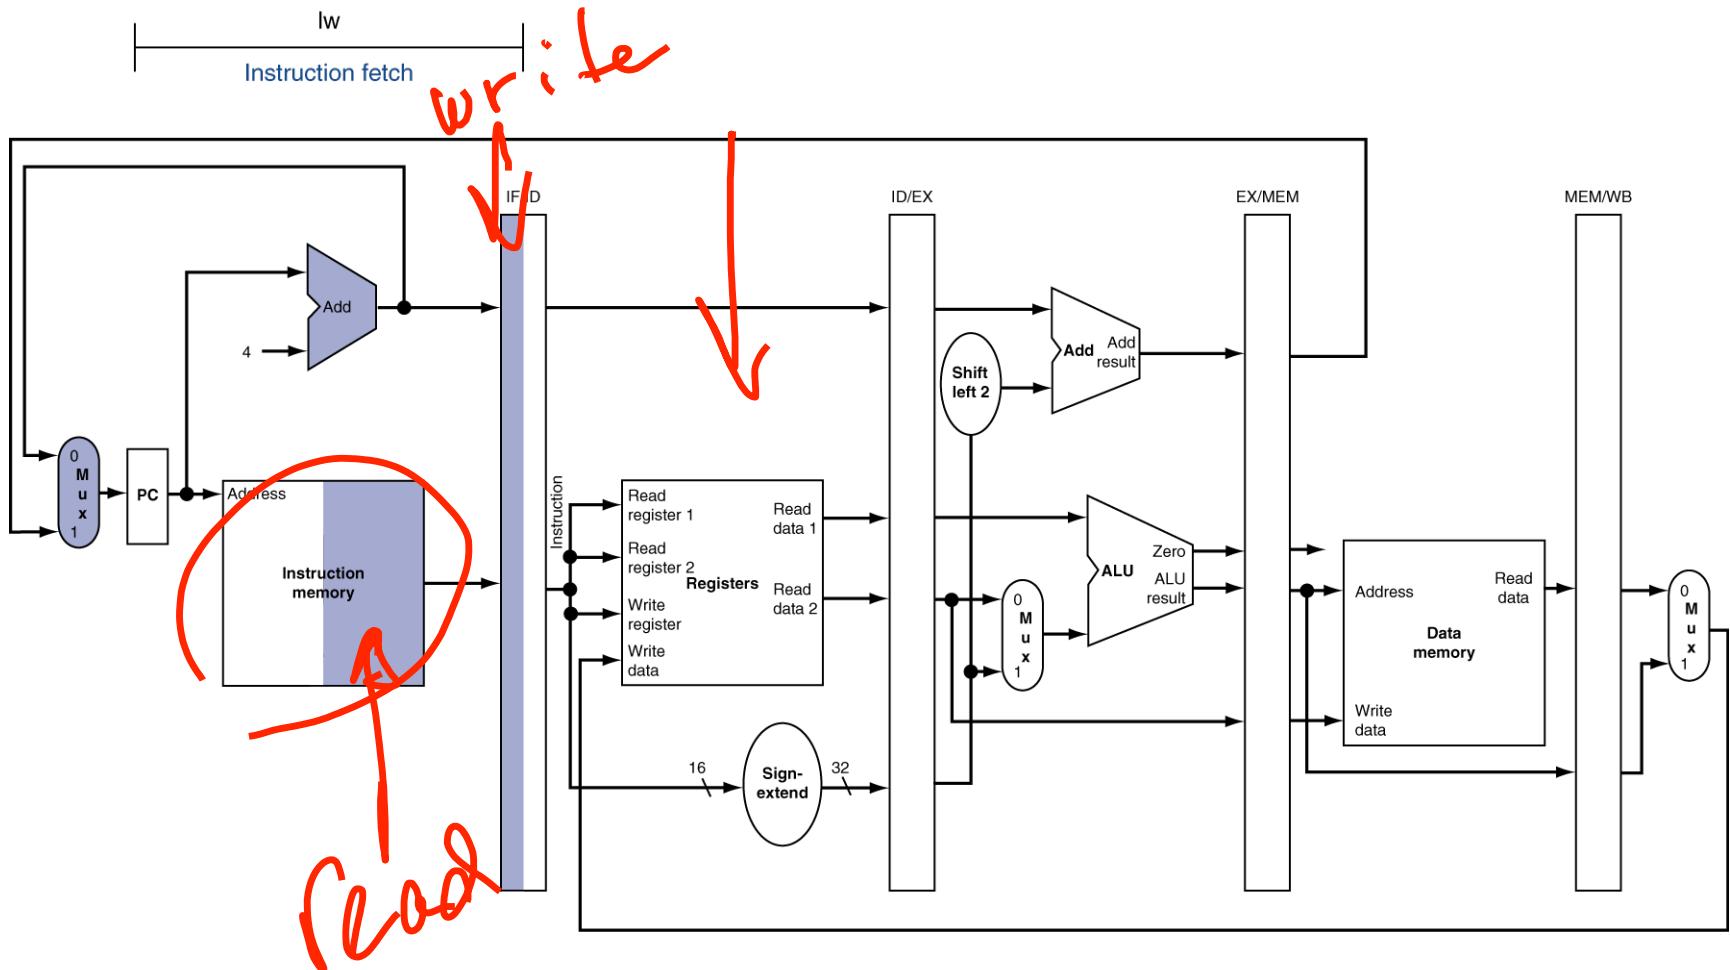


Pipeline registers

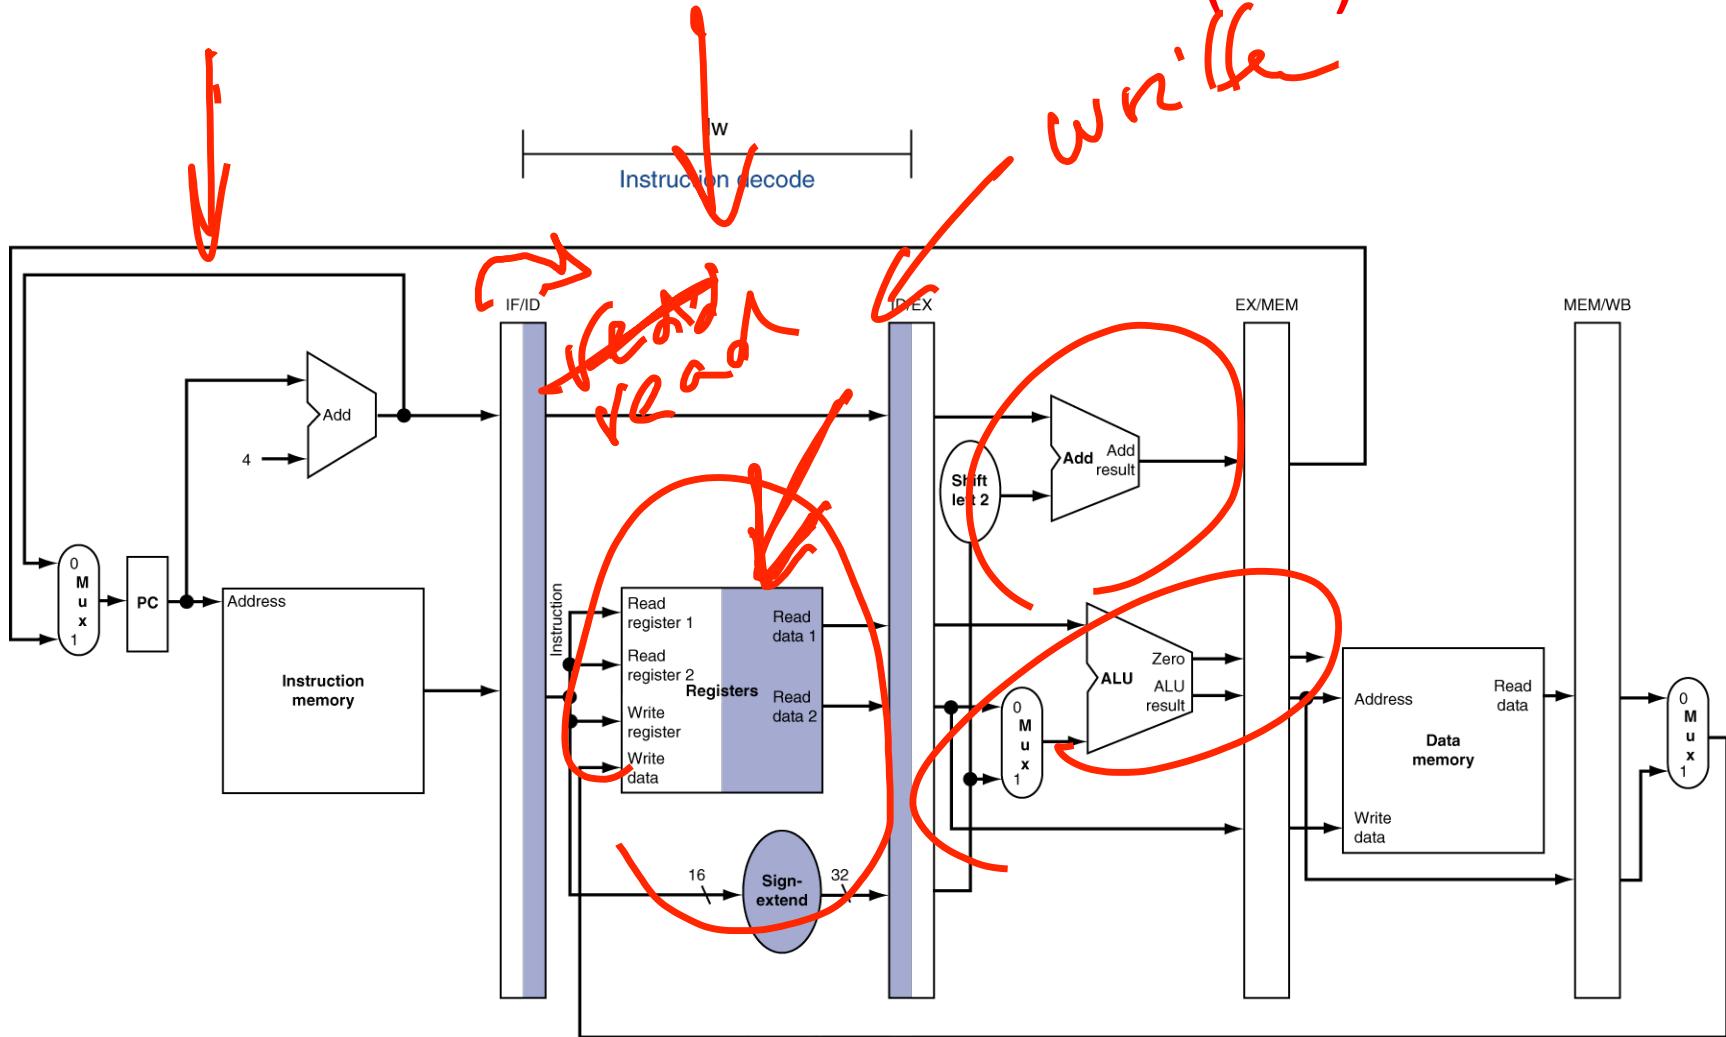
- Each stage operates on different data
 - E.g. ALU processes data fetched from registers during last cycle
 - → insert pipelining registers to hold data



Instruction Fetch (IF)



Instruction Decode (ID)

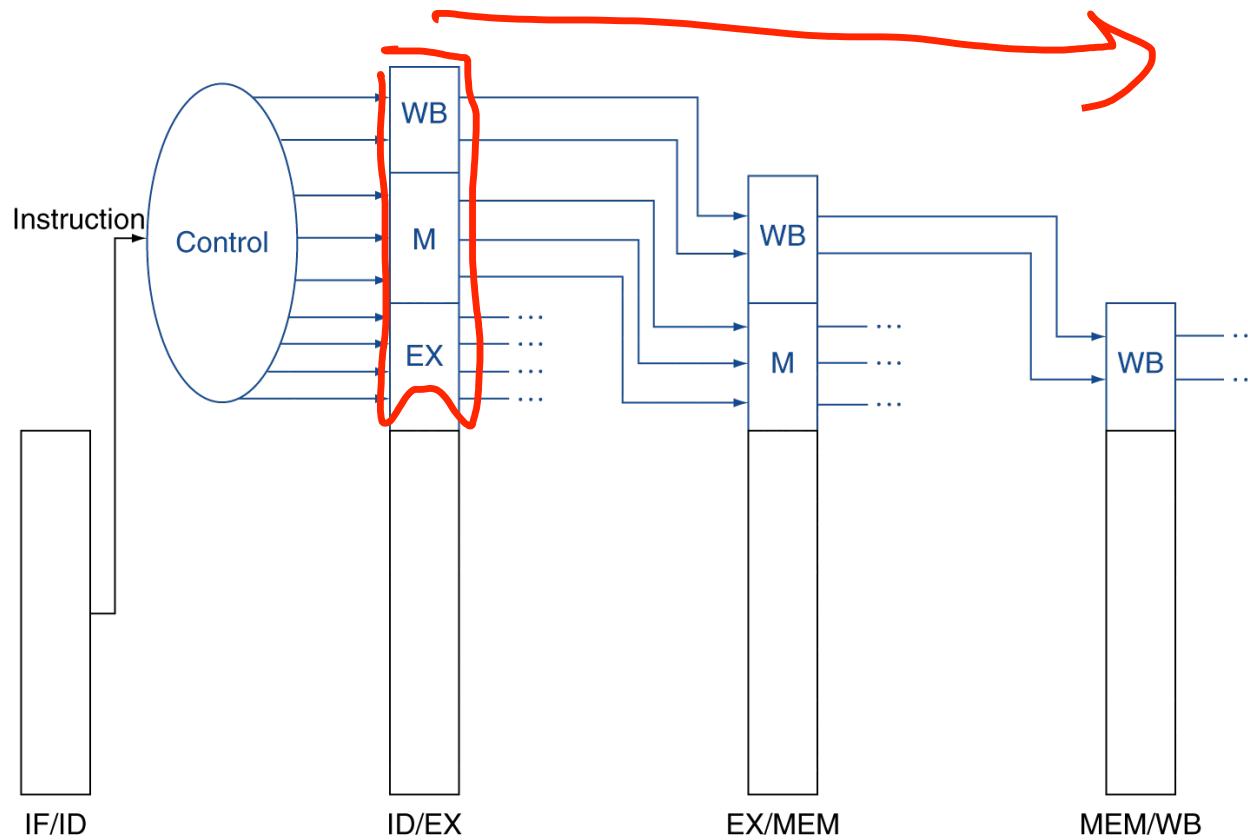


Agenda

- MIPS Pipeline
- **Pipeline Control**
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages



Hazards Ahead



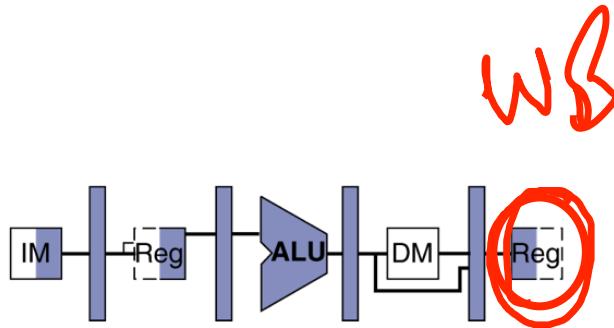
Agenda

- MIPS Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

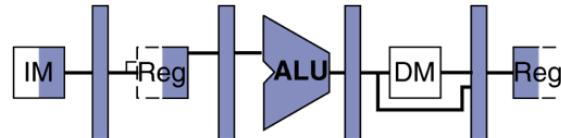
Structural Hazard: Register Access

instruction sequence ↓

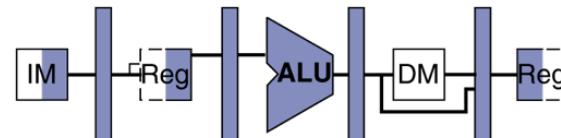
add \$t0, \$t1, \$t2



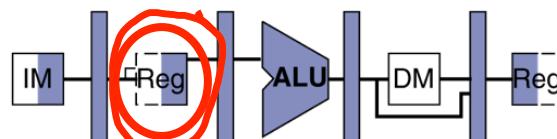
or \$t3, \$t4, \$t5



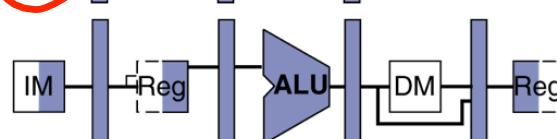
slt \$t6, \$t0, \$t3



sw \$t0, 4(\$t3)



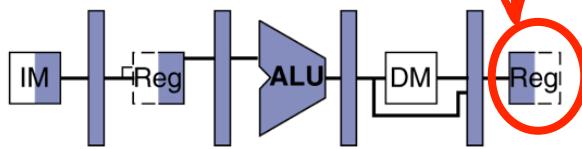
lw \$t0, 8(\$t3)



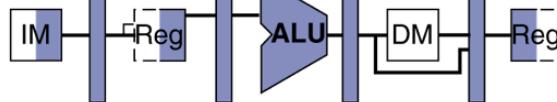
- Register file accessed simultaneously by WB and ID phases
- Does **sw** in the example fetch the old or new value?

Register Access Policy

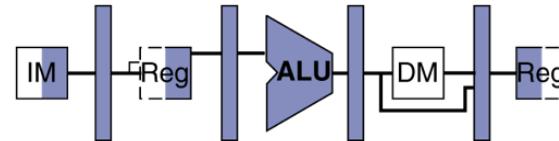
add \$t0, \$t1, \$t2



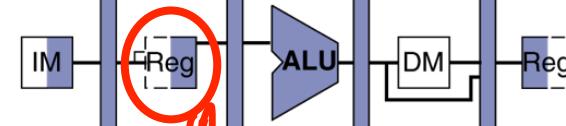
or \$t3, \$t4, \$t5



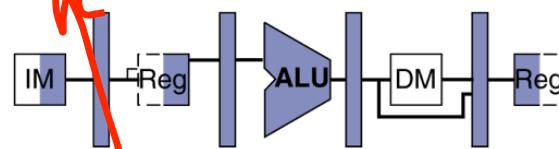
slt \$t6, \$t0, \$t3



sw \$t0, 4(\$t3)



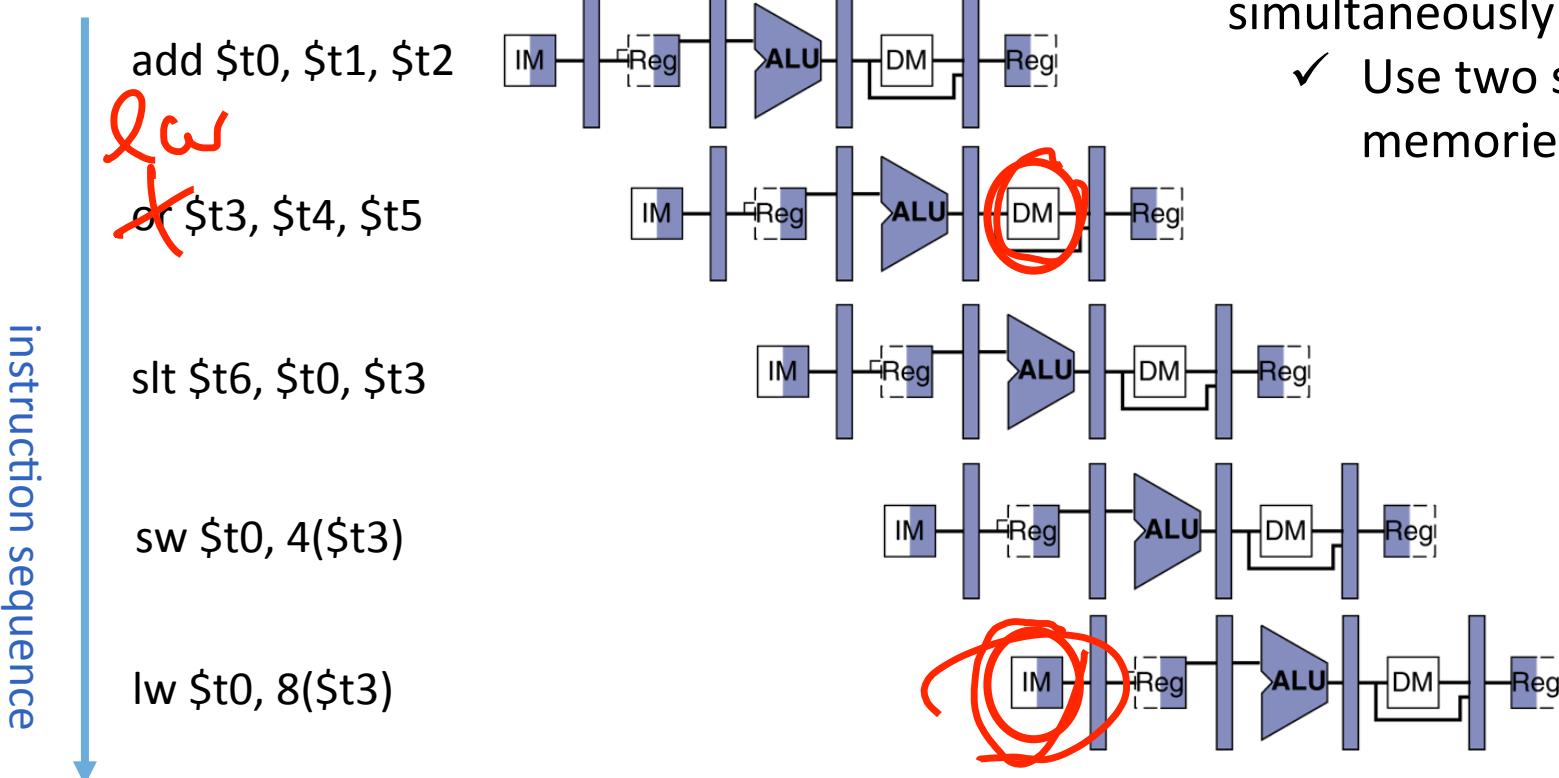
lw \$t0, 8(\$t3)



- Exploit high speed of register file (100 ps)
 - 1) WB updates value
 - 2) ID reads new value
- Indicated in diagram by shading

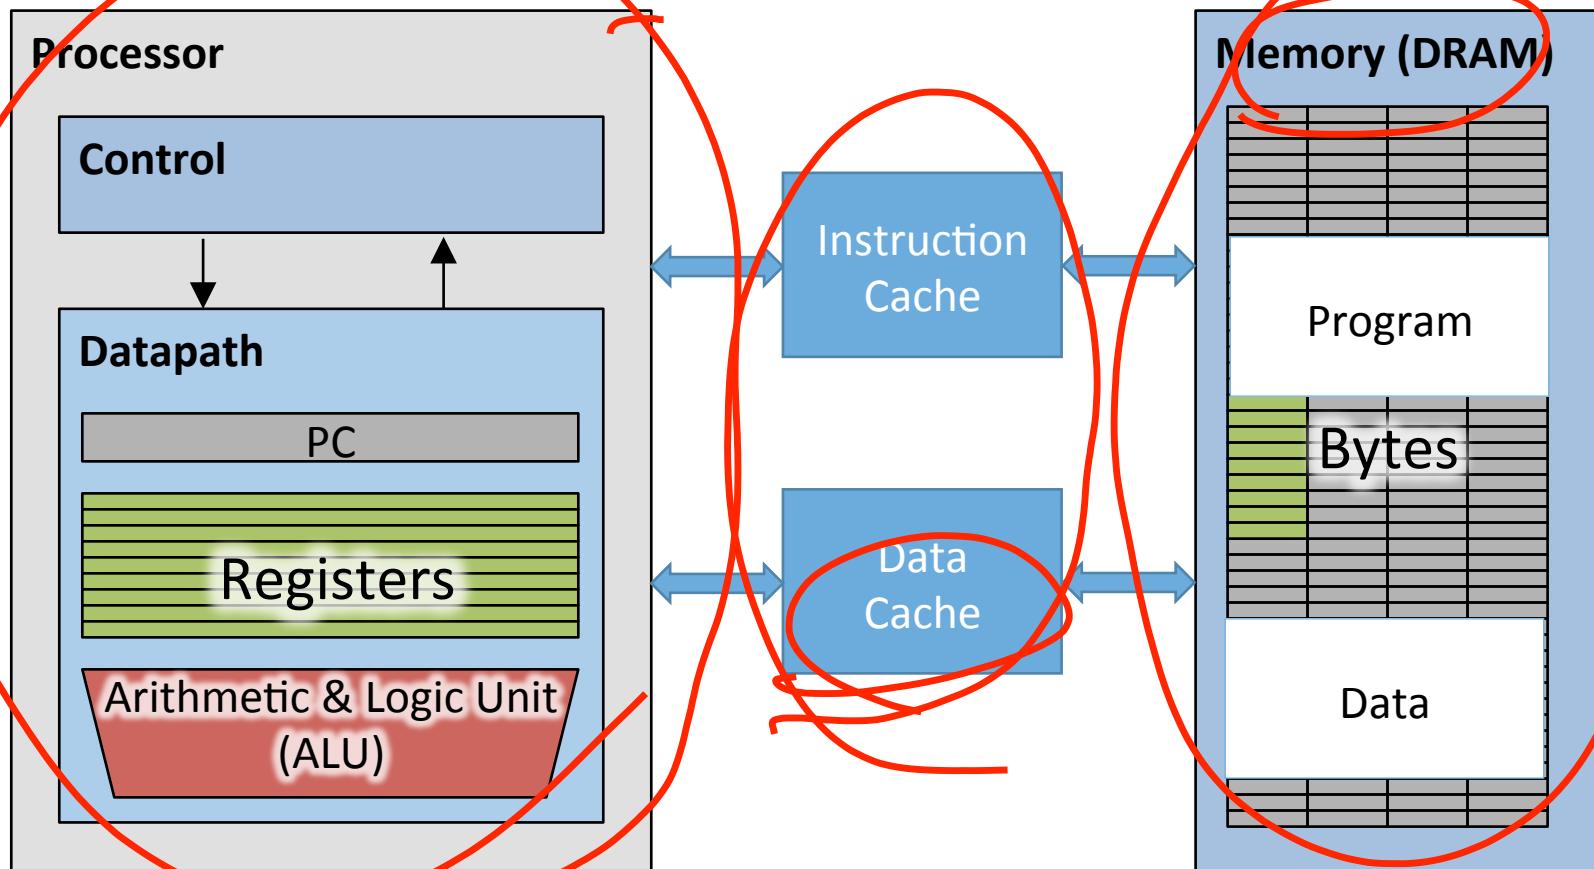
instruction sequence

Structural Hazard: Memory Access



- Instruction and data memory used simultaneously
 - ✓ Use two separate memories

Instruction and Data Caches



Caches: small and fast “buffer” memories

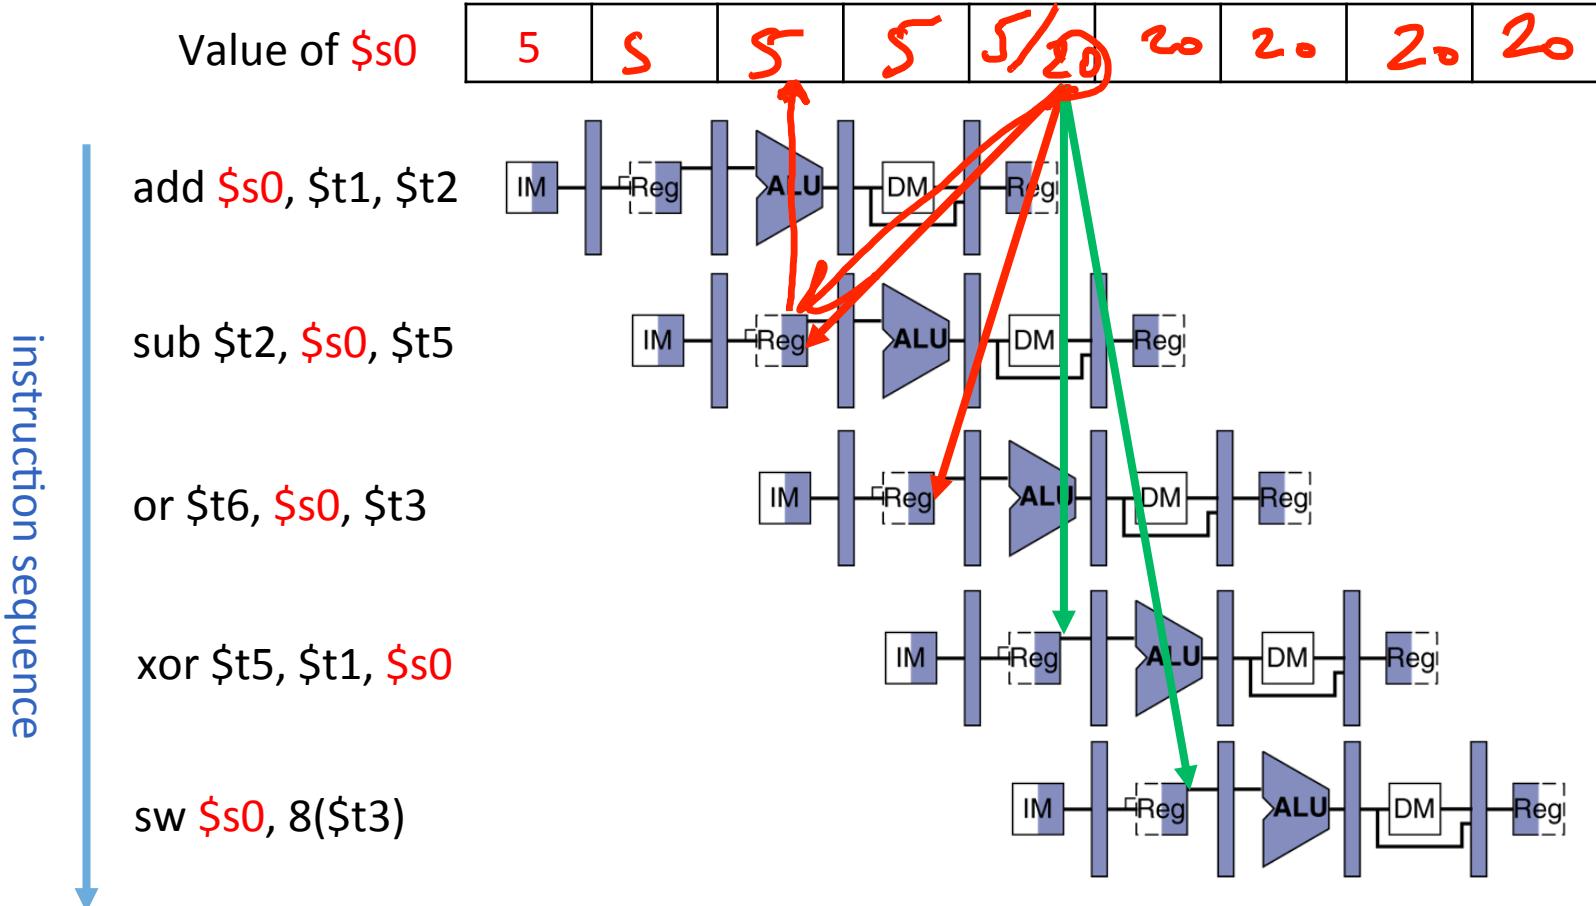
Structure Hazards – Summary

- Conflict for use of a resource
- In MIPS pipeline with a single memory
 - Load/store requires data access
 - Without separate memories, instruction fetch would have to *stall* for that cycle
 - All other operations in pipeline would have to wait
- Pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches

Agenda

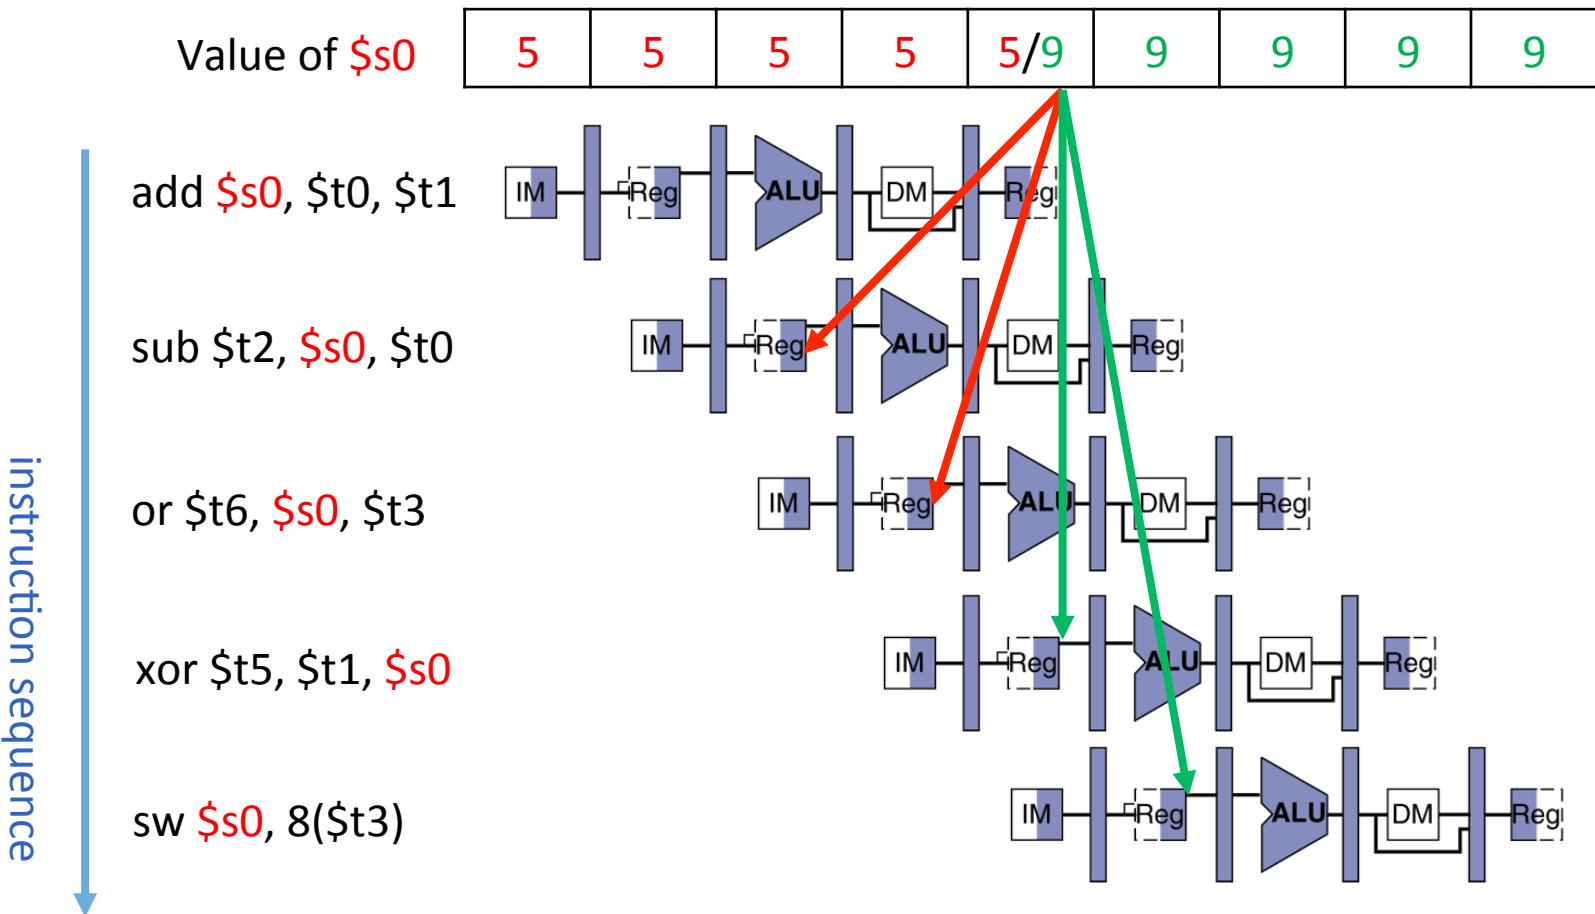
- MIPS Pipeline
- Pipeline Control
- Hazards
 - Structural
 - **Data**
 - **R-type instructions**
 - Load
 - Control
- Superscalar processors

Data Hazard: ALU Result



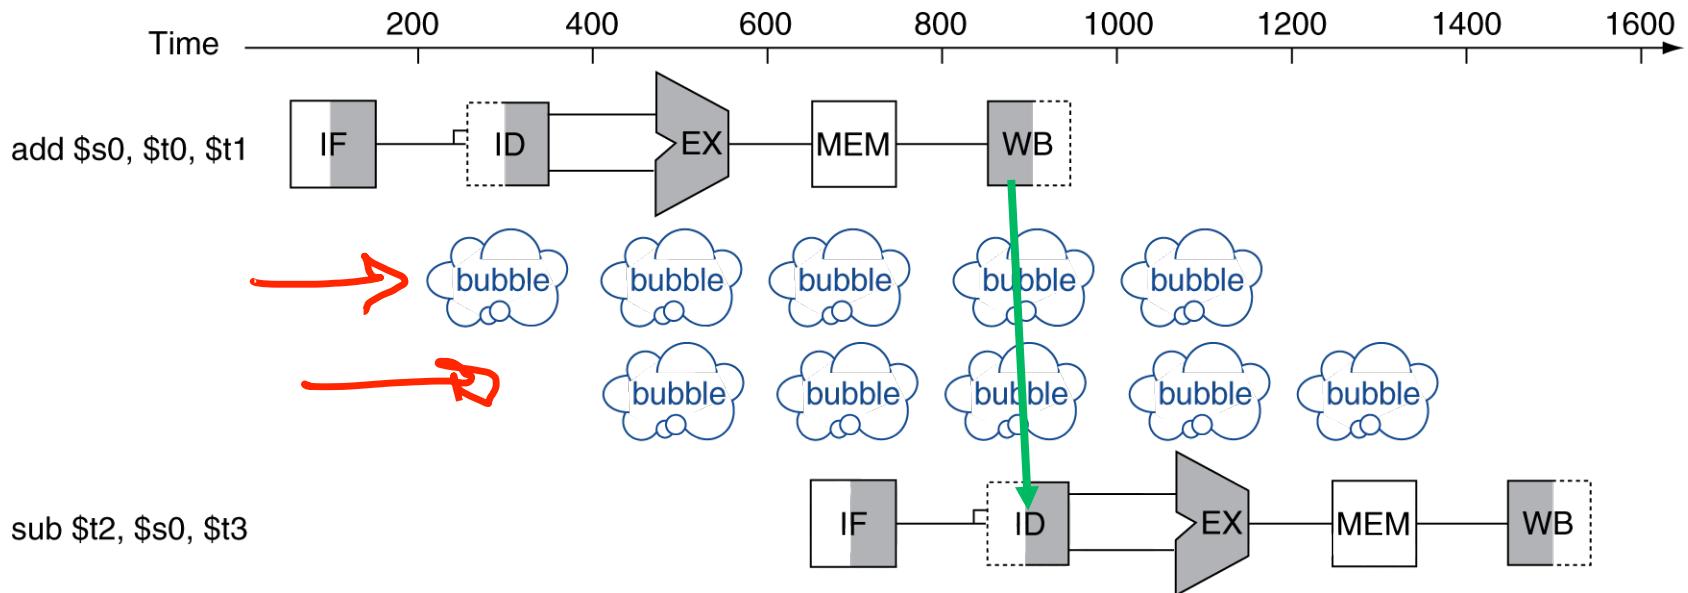
Without fix, **sub** and **or** calculate wrong result!

Data Hazard: ALU Result



Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction
 - add $\$s0, \$t0, \$t1$
 - sub $\$t2, \$s0, \$t3$

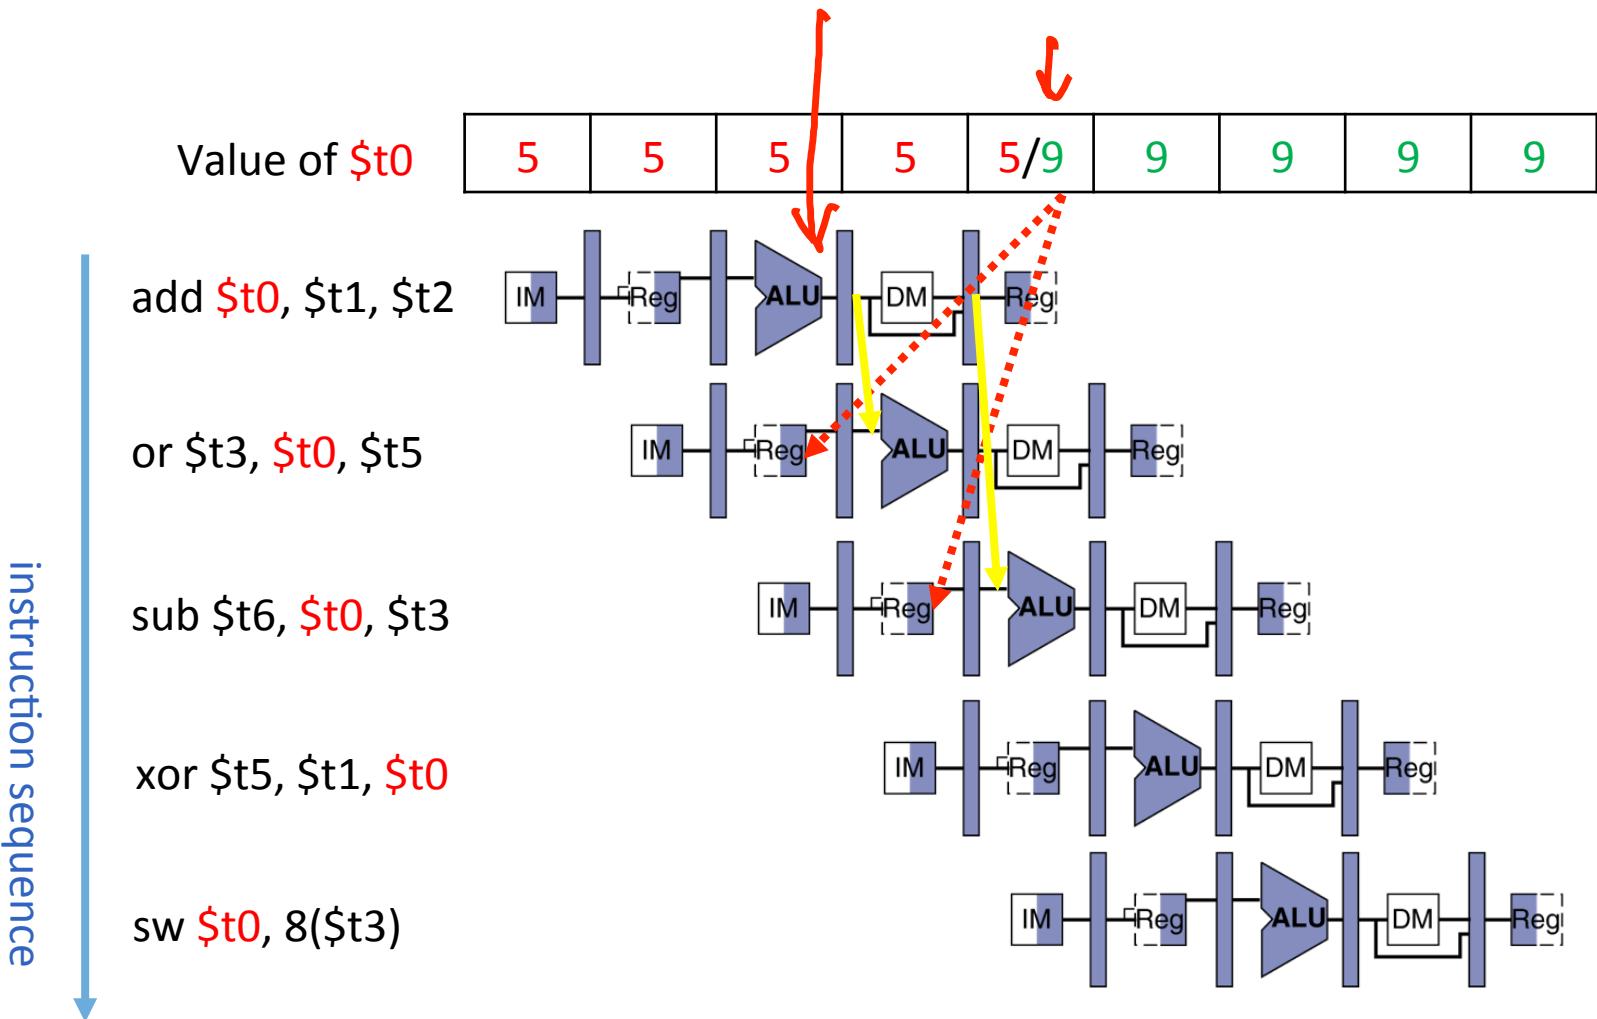


- Bubble:
 - effectively NOP: affected pipeline stages do “nothing”

Stalls and Performance

- Stalls reduce performance
 - But stalls are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

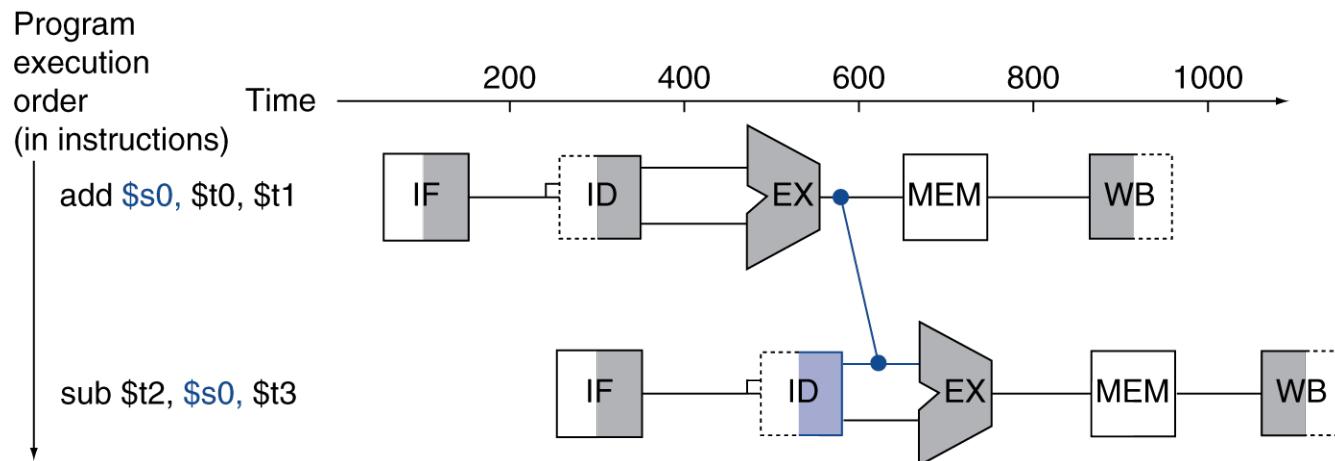
Solution 2: Forwarding



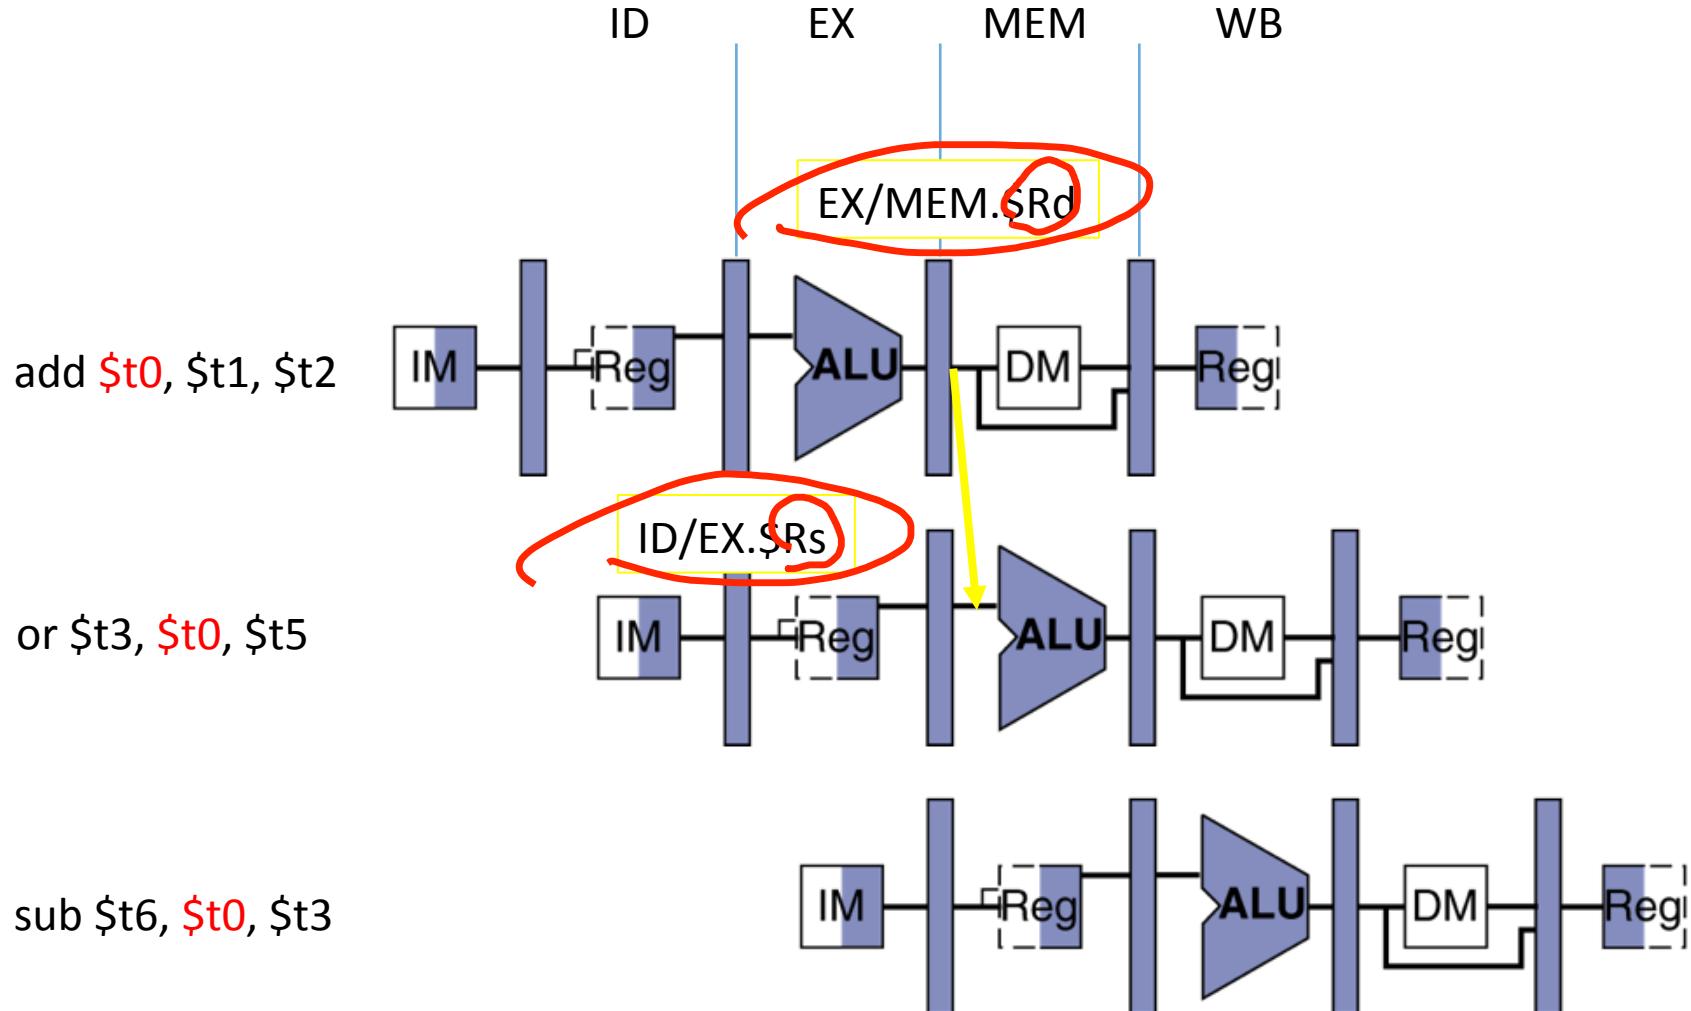
Forwarding: grab operand from pipeline register, rather than register file

Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



1) Detect Need for Forwarding (example)



Complete Forwarding Conditions

- MEM hazard

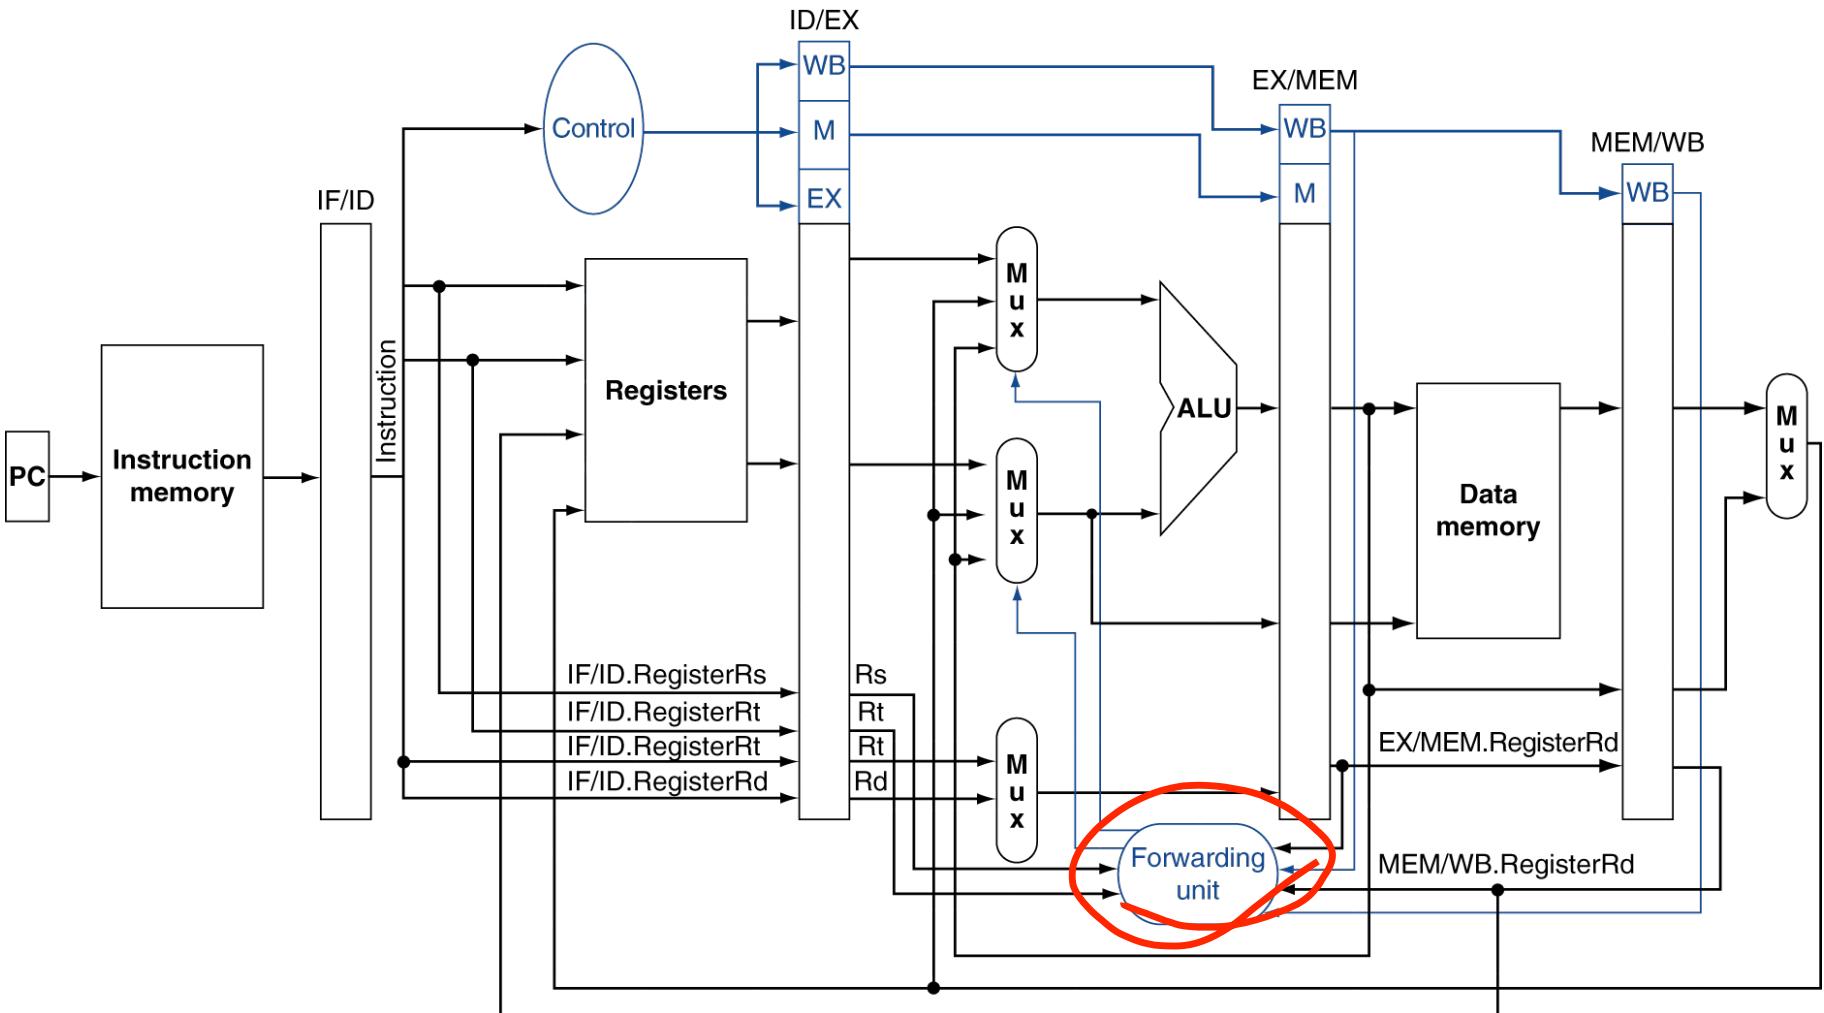
- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRs))
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))

ForwardA = 01

- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRt))
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))

ForwardB = 01

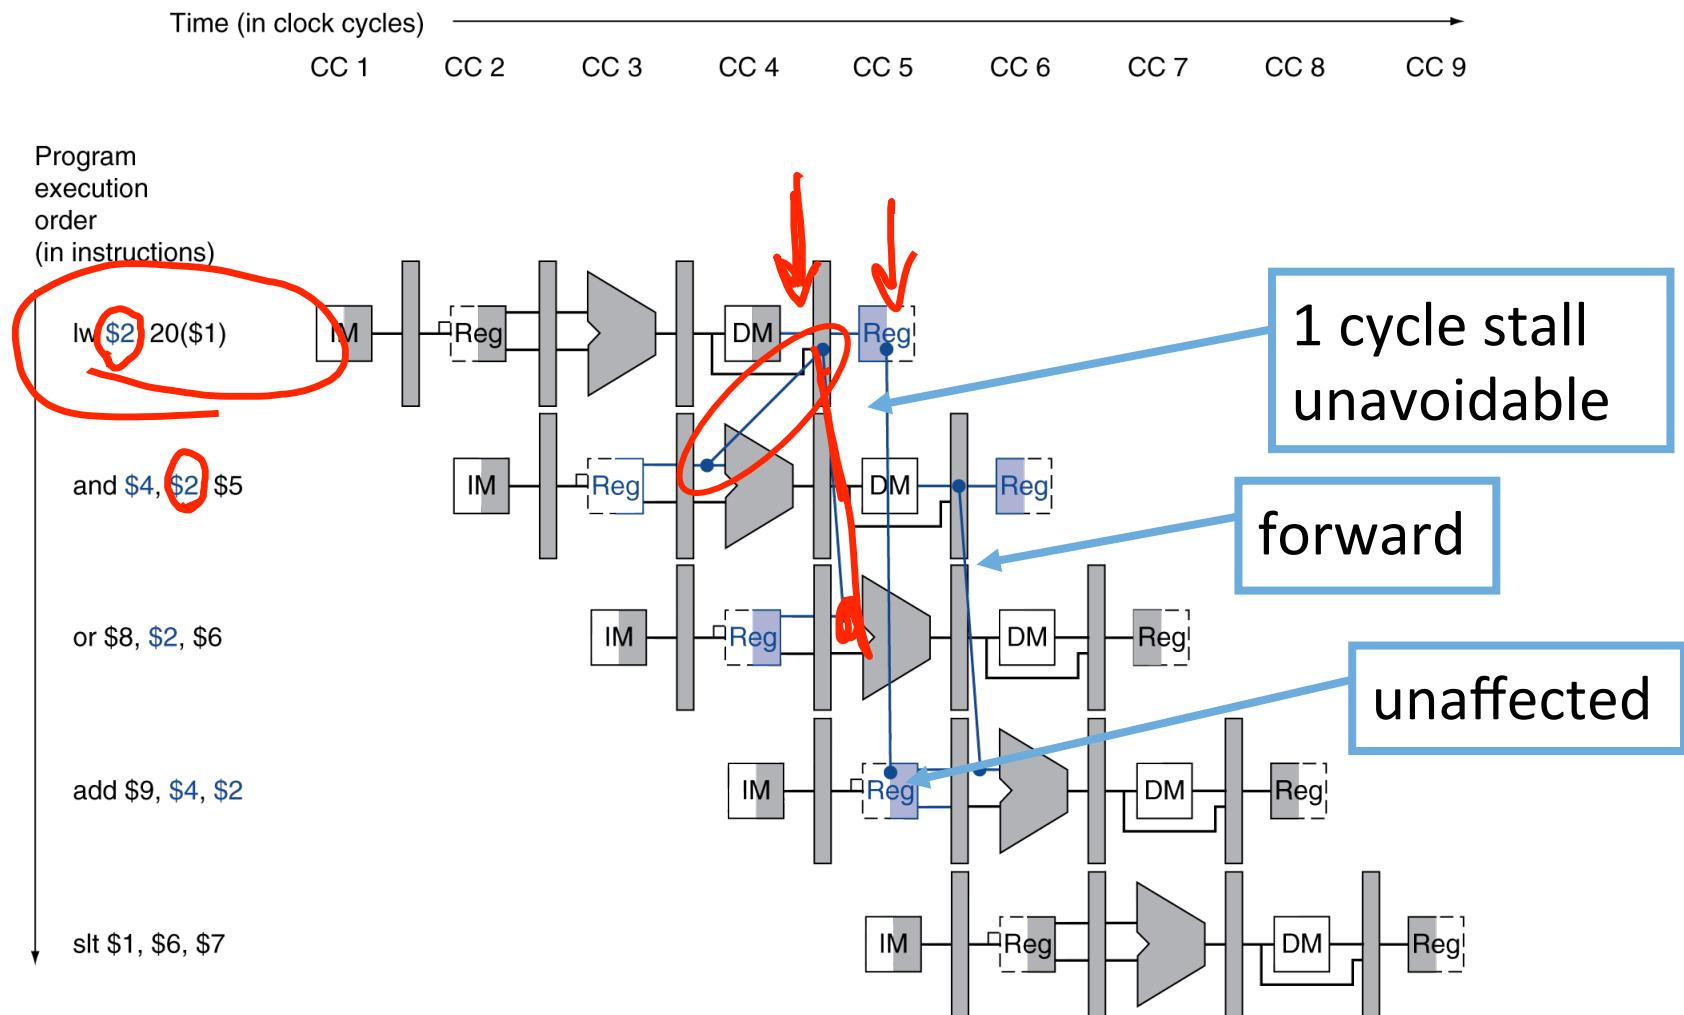
2) Datapath with Forwarding



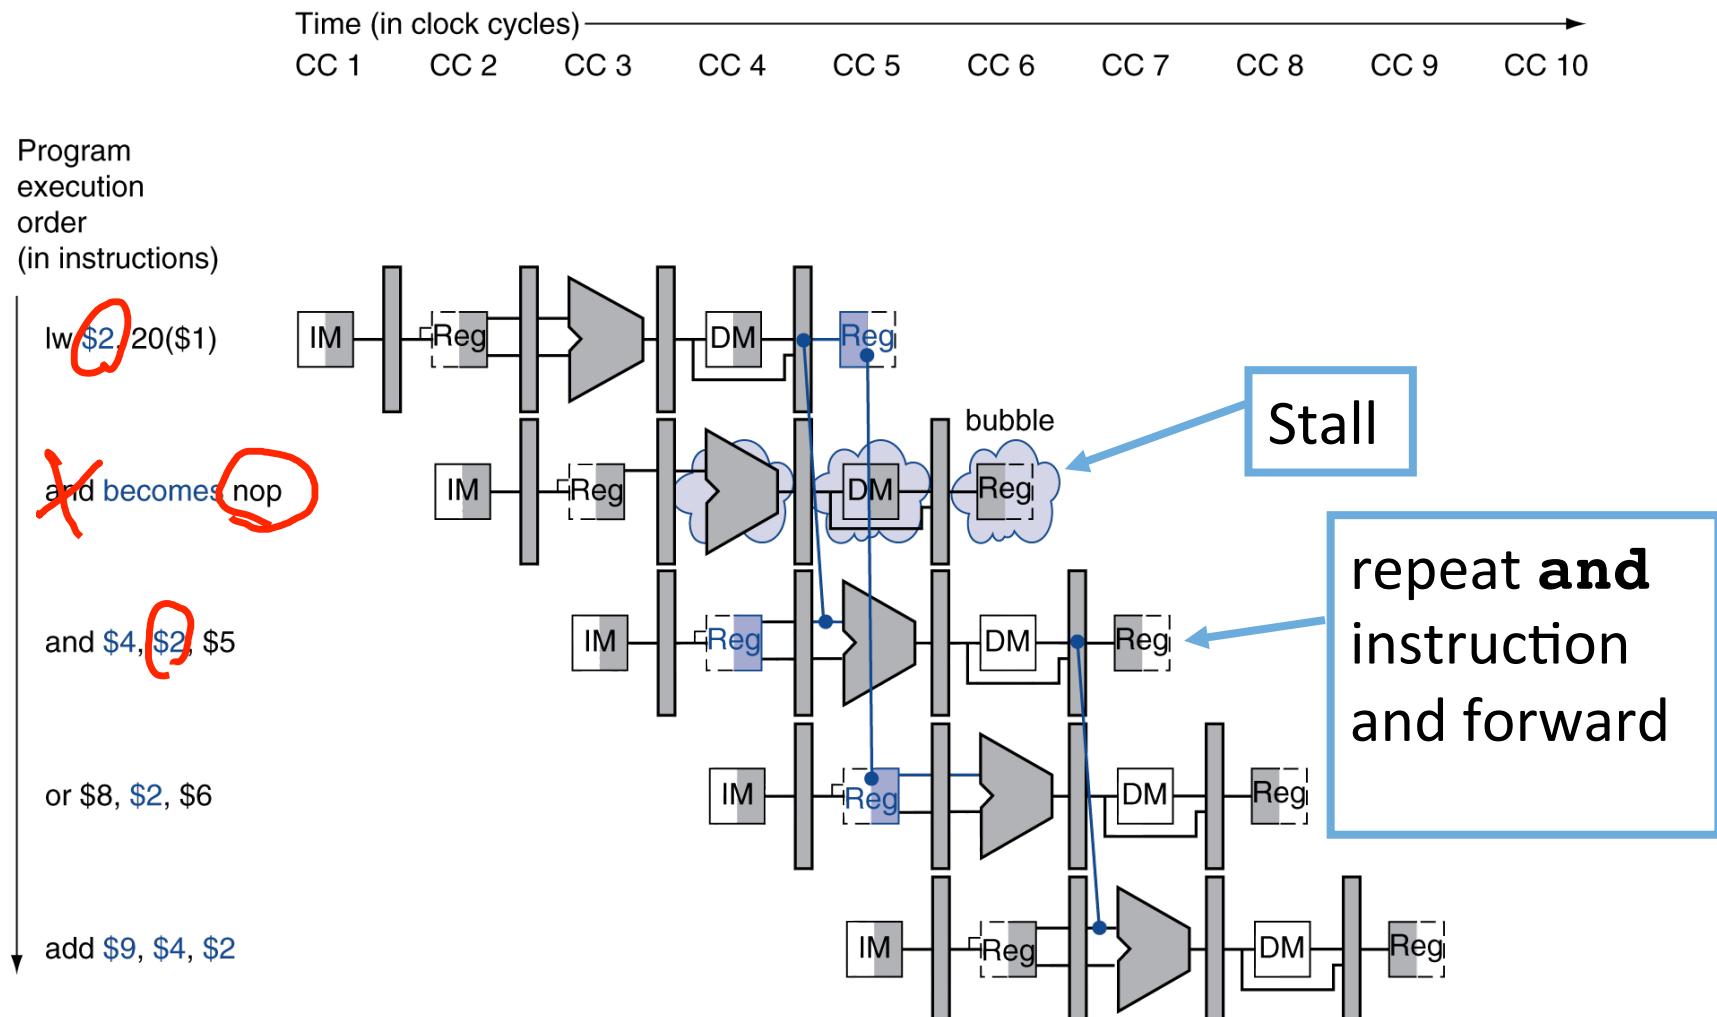
Agenda

- MIPS Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- Superscalar processors

Load Data Hazard



Stall Pipeline



Your Turn

How many cycles (pipeline fill+process+drain) does it take to *complete* the following code?

—lw \$t1, 0(\$t0)
—lw \$t2, 4(\$t0) *map*
—add \$t3, \$t1, \$t2
—sw \$t3, 12(\$t0)
—lw \$t4, 8(\$t0) *map*
—add \$t5, \$t1, \$t4
—sw \$t5, 16(\$t0)

Answer	cycles
A	7
B	9
C	11
D	13
E	14

$$7 + 2 + 4$$

Your Turn

How many cycles (pipeline fill+process+drain) does it take to *complete* the following code?

```
lw    $t1, 0($t0)
lw    $t2, 4($t0)
add  $t3, $t1, $t2
sw    $t3, 12($t0)
lw    $t4, 8($t0)
add  $t5, $t1, $t4
sw    $t5, 16($t0)
```

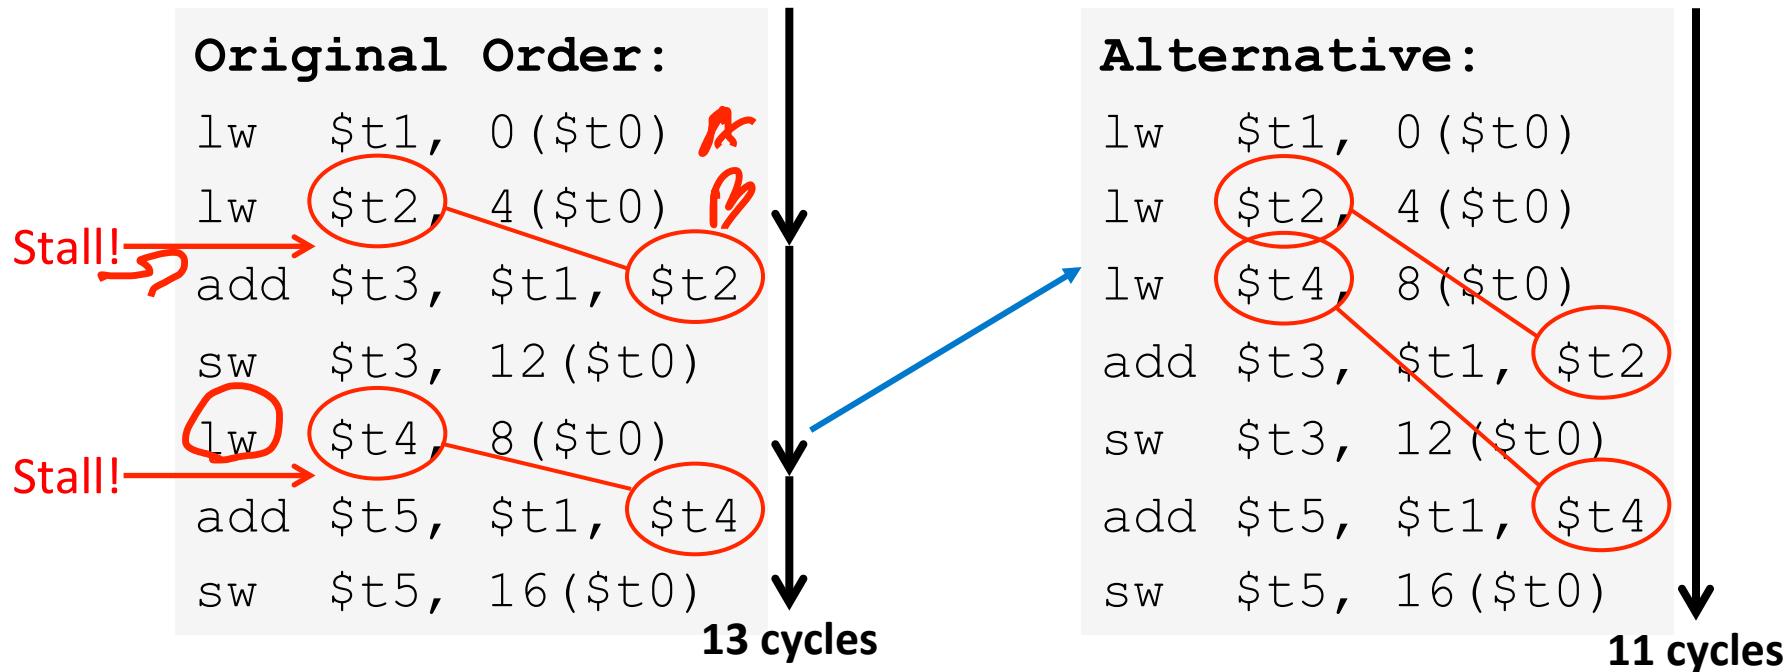
Answer	cycles
A	7
B	9
C	11
D	13
E	14

lw Data Hazard

- Slot after a load is called a *load delay slot*
 - If that instruction uses the result of the load, then the hardware will stall for one cycle
 - Equivalent to inserting an explicit **nop** in the slot
 - except the latter uses more code space
 - Performance loss
- Idea:
 - Put unrelated instruction into load delay slot
 - No performance loss!

Code Scheduling to Avoid Stalls

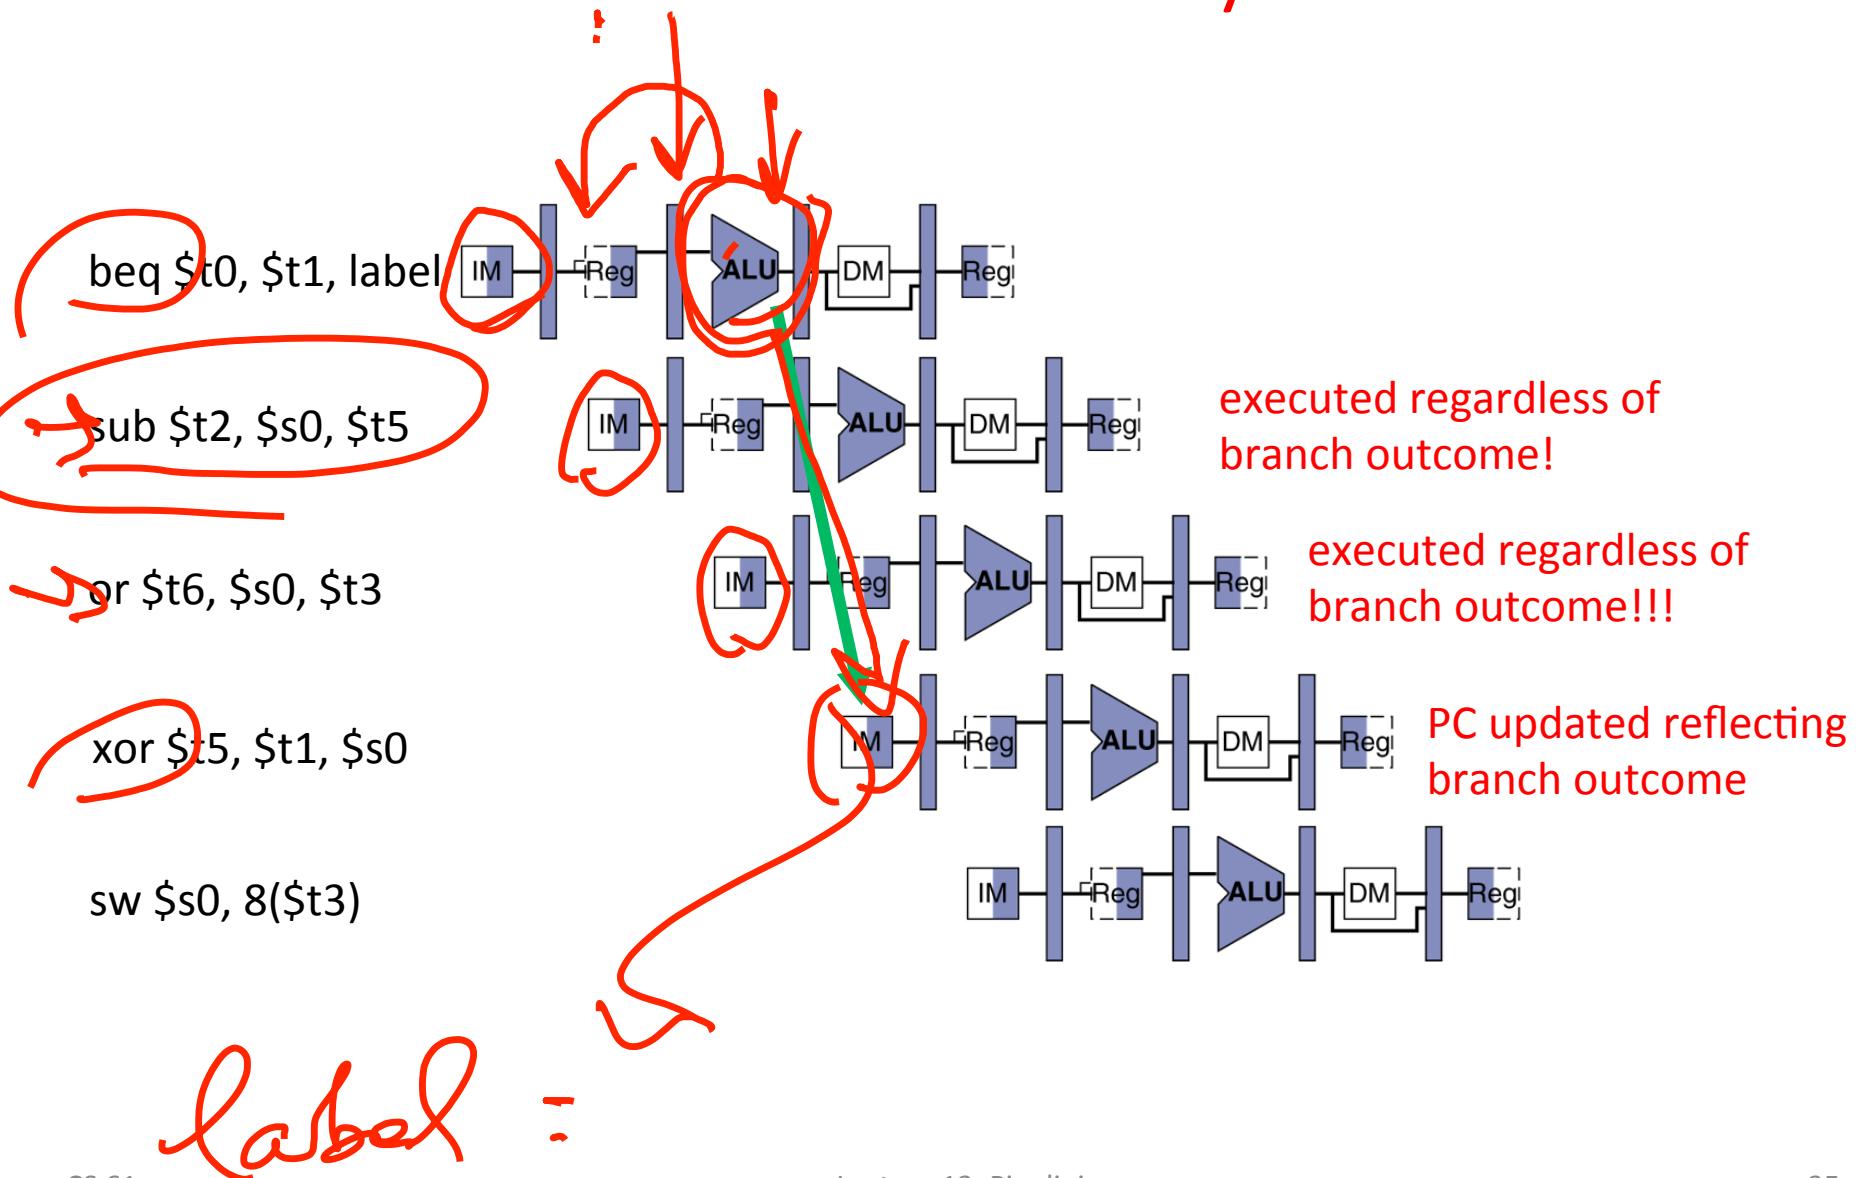
- Reorder code to avoid use of load result in the next instruction!
- MIPS code for $D=A+B; E=A+C;$



Agenda

- MIPS Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - **Control**
- Superscalar processors

Branch Delay



Dealing with Branch Control Hazard

- Without hardware modification, 2 instructions after branch are executed
 - regardless of branch outcome
- Hardware modification:
 - move branch decision to ID stage
 - requires dedicated equality comparator
- Only 1 instruction after branch executed unconditionally
 - “branch delay slot”
 - MIPS assembler/compiler try to schedule unaffected instruction
 - **nop** if none

Branch Delay Slot Scheduling Example

Original sequence:

or \$8, \$9, \$10
add \$1, \$2, \$3
sub \$4, \$5, \$6
beq \$1, \$4, Exit
nop
xor \$10, \$1, \$11

Improved sequence:

add \$1, \$2, \$3
sub \$4, \$5, \$6
beq \$1, \$4, Exit
or \$8, \$9, \$10
xor \$10, \$1, \$11

Alternatives to Branch Delay Slot

- Predict (“guess”) branch outcome
 - Transparent to programmer
 - Cancel instructions in pipeline if guess was wrong
 - “flush pipeline”
- Typical solutions:
 - Always guess that branch is not taken
 - Do same as last time branch instruction was executed
 - requires special “book-keeping hardware”
 - Solution used by most current processors
 - except MIPS

Agenda

- MIPS Pipeline
- Pipeline Control
- Hazards
 - Structural
 - Data
 - R-type instructions
 - Load
 - Control
- **Superscalar processors**

Increasing Processor Performance

1. Clock rate

- Limited by technology and power dissipation

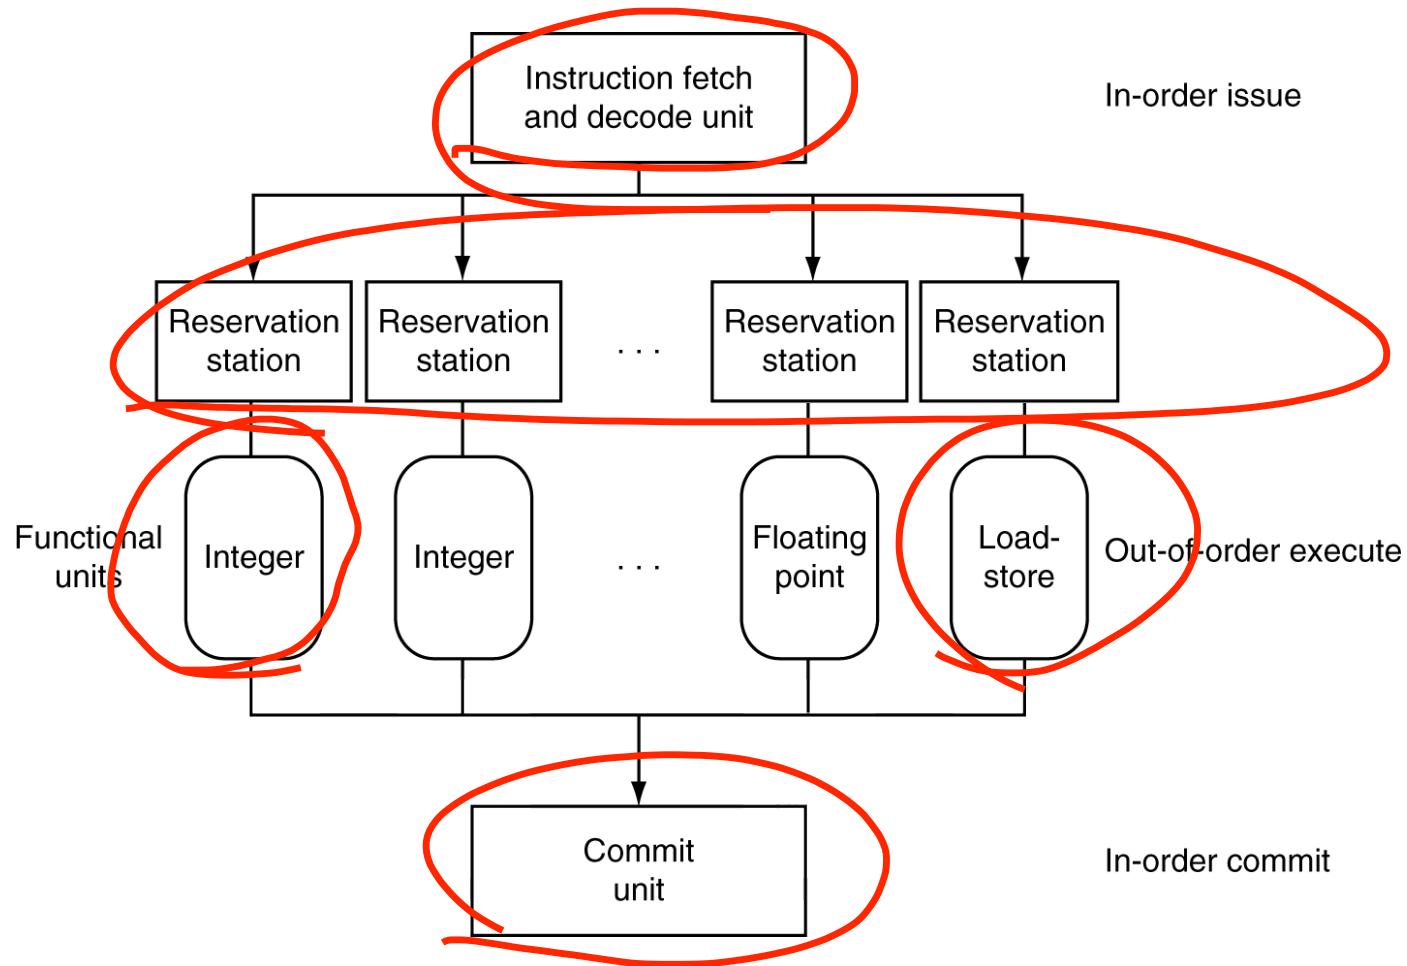
2. Pipelining

- “Overlap” instruction execution
- Deeper pipeline: 5 => 10 => 15 stages
 - Less work per stage → shorter clock cycle
 - But more potential for hazards (CPI > 1)

3. Multi-issue “super-scalar” processor

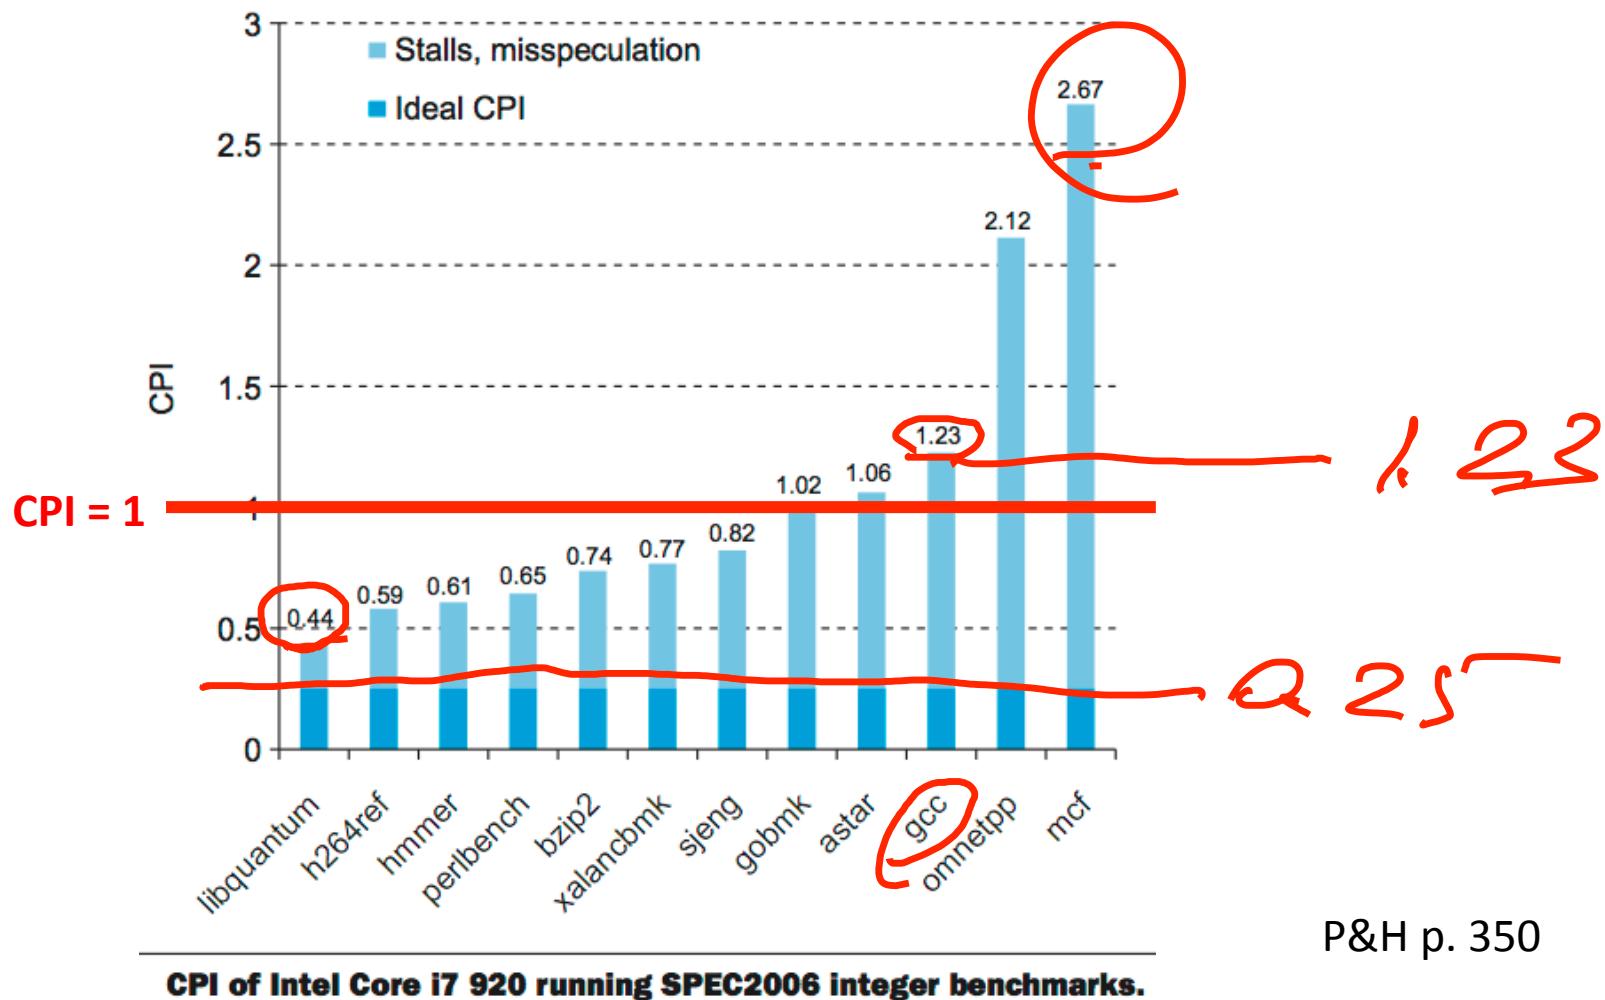
- Multiple execution units (ALUs)
 - Several instructions executed simultaneously
 - CPI < 1 (ideally)

Superscalar Processor



P&H p. 340

Benchmark: CPI of Intel Core i7



In Conclusion

- Pipelining increases throughput by overlapping execution of multiple instructions
- All pipeline stages have same duration
 - Choose partition that accommodates this constraint
- Hazards ~~potentially~~ limit performance
 - Maximizing performance requires programmer/compiler assistance
 - E.g. Load and Branch delay slots
- Superscalar processors use multiple execution units for additional instruction level parallelism
 - Performance benefit highly code dependent

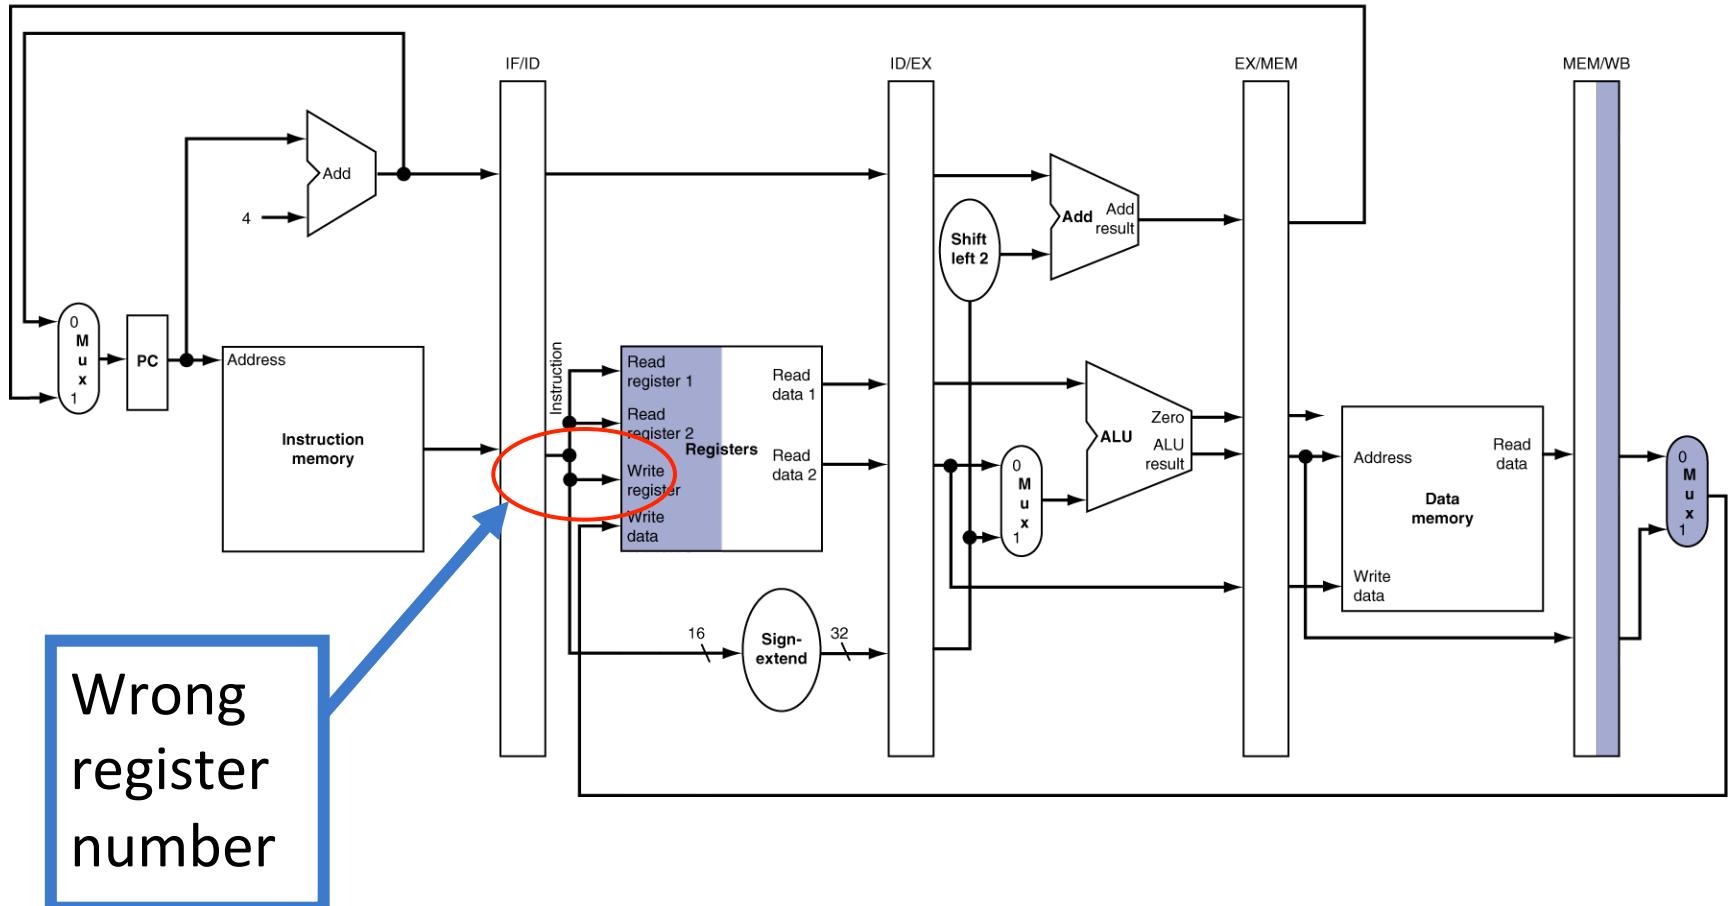
Extra Slides

Pipelining and ISA Design

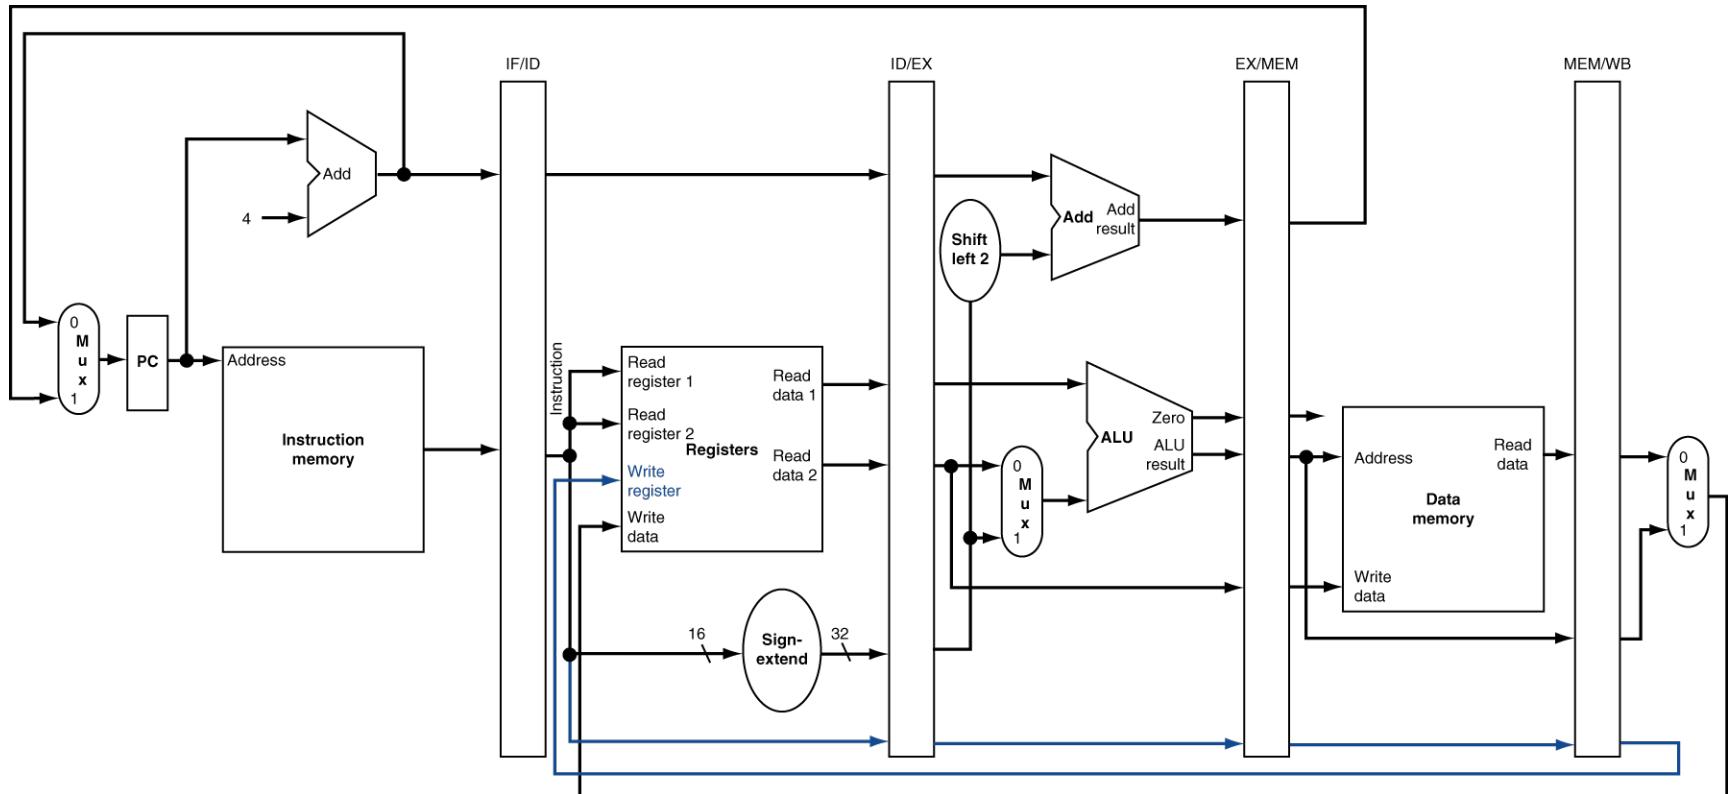
- MIPS ISA designed for pipelining
 - All instructions are 32-bits
 - Easy to fetch and decode in one cycle
 - Versus x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Decode and read registers in one step
 - Load/store addressing
 - Calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle

Write-Back (WB, **1w**)

lw
Write back



Corrected Datapath for WB



Superscalar Processor

- Multiple issue “superscalar”
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - CPI < 1, so use Instructions Per Cycle (IPC)
 - E.g., 4GHz 4-way multiple-issue
 - 16 BIPS, peak CPI = 0.25, peak IPC = 4
 - Dependencies reduce this in practice
- “Out-of-Order” execution
 - Reorder instructions dynamically in hardware to reduce impact of hazards
- *CS152 discusses these techniques!*