

# CS 61C: Great Ideas in Computer Architecture (Machine Structures)

## Caches Part 2

Instructors:

Bernhard Boser & Randy H. Katz

<http://inst.eecs.berkeley.edu/~cs61c/>

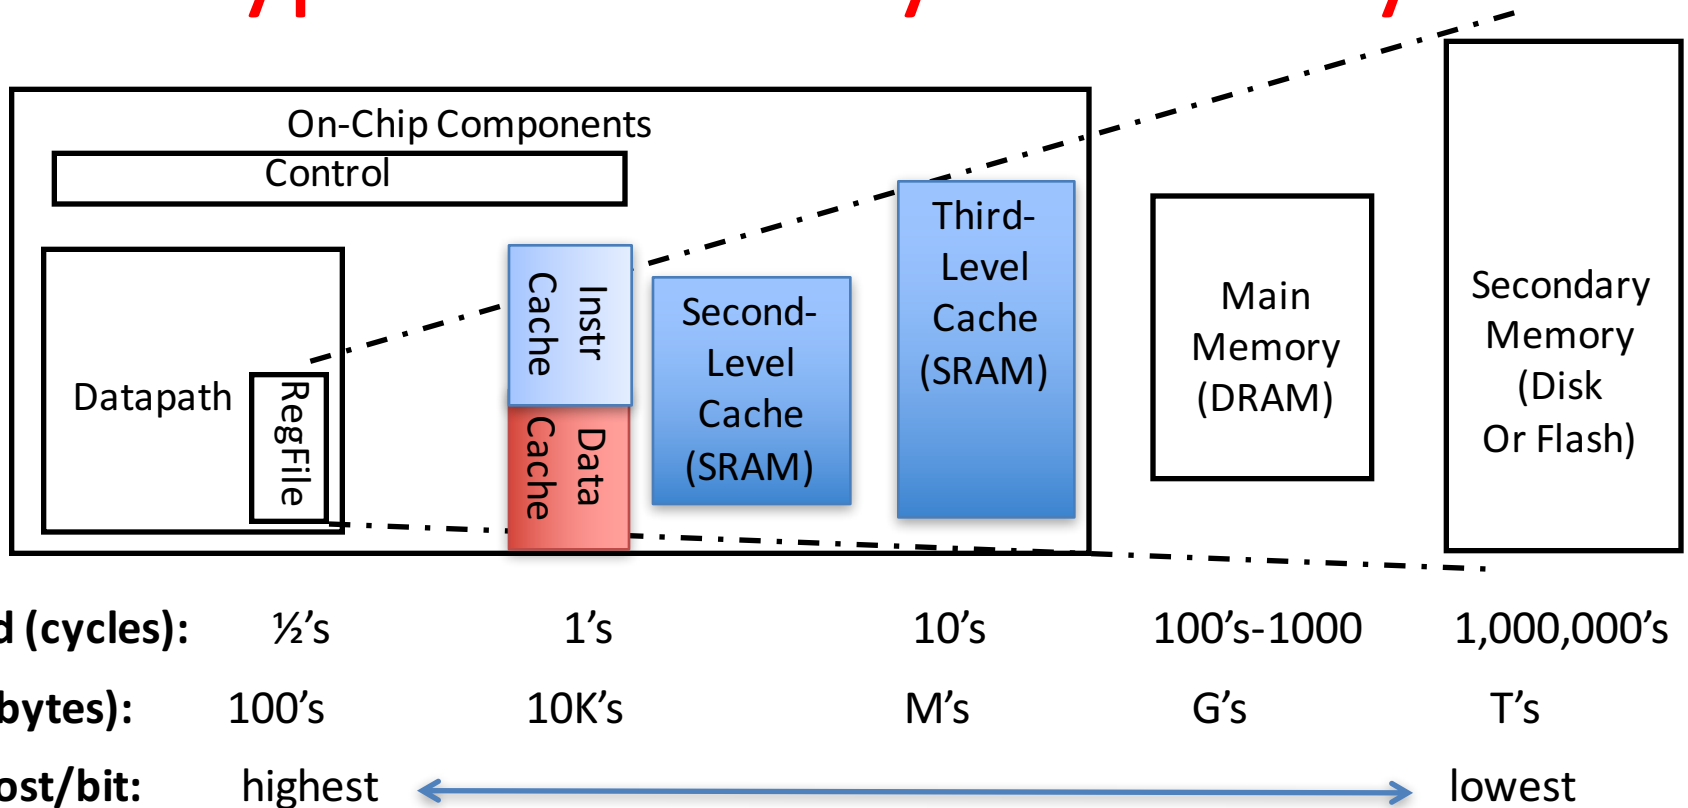
# Outline

- Cache Organization and Principles
- Write Back vs. Write Through
- Cache Performance
- Cache Design Tradeoffs
- And in Conclusion ...

# Outline

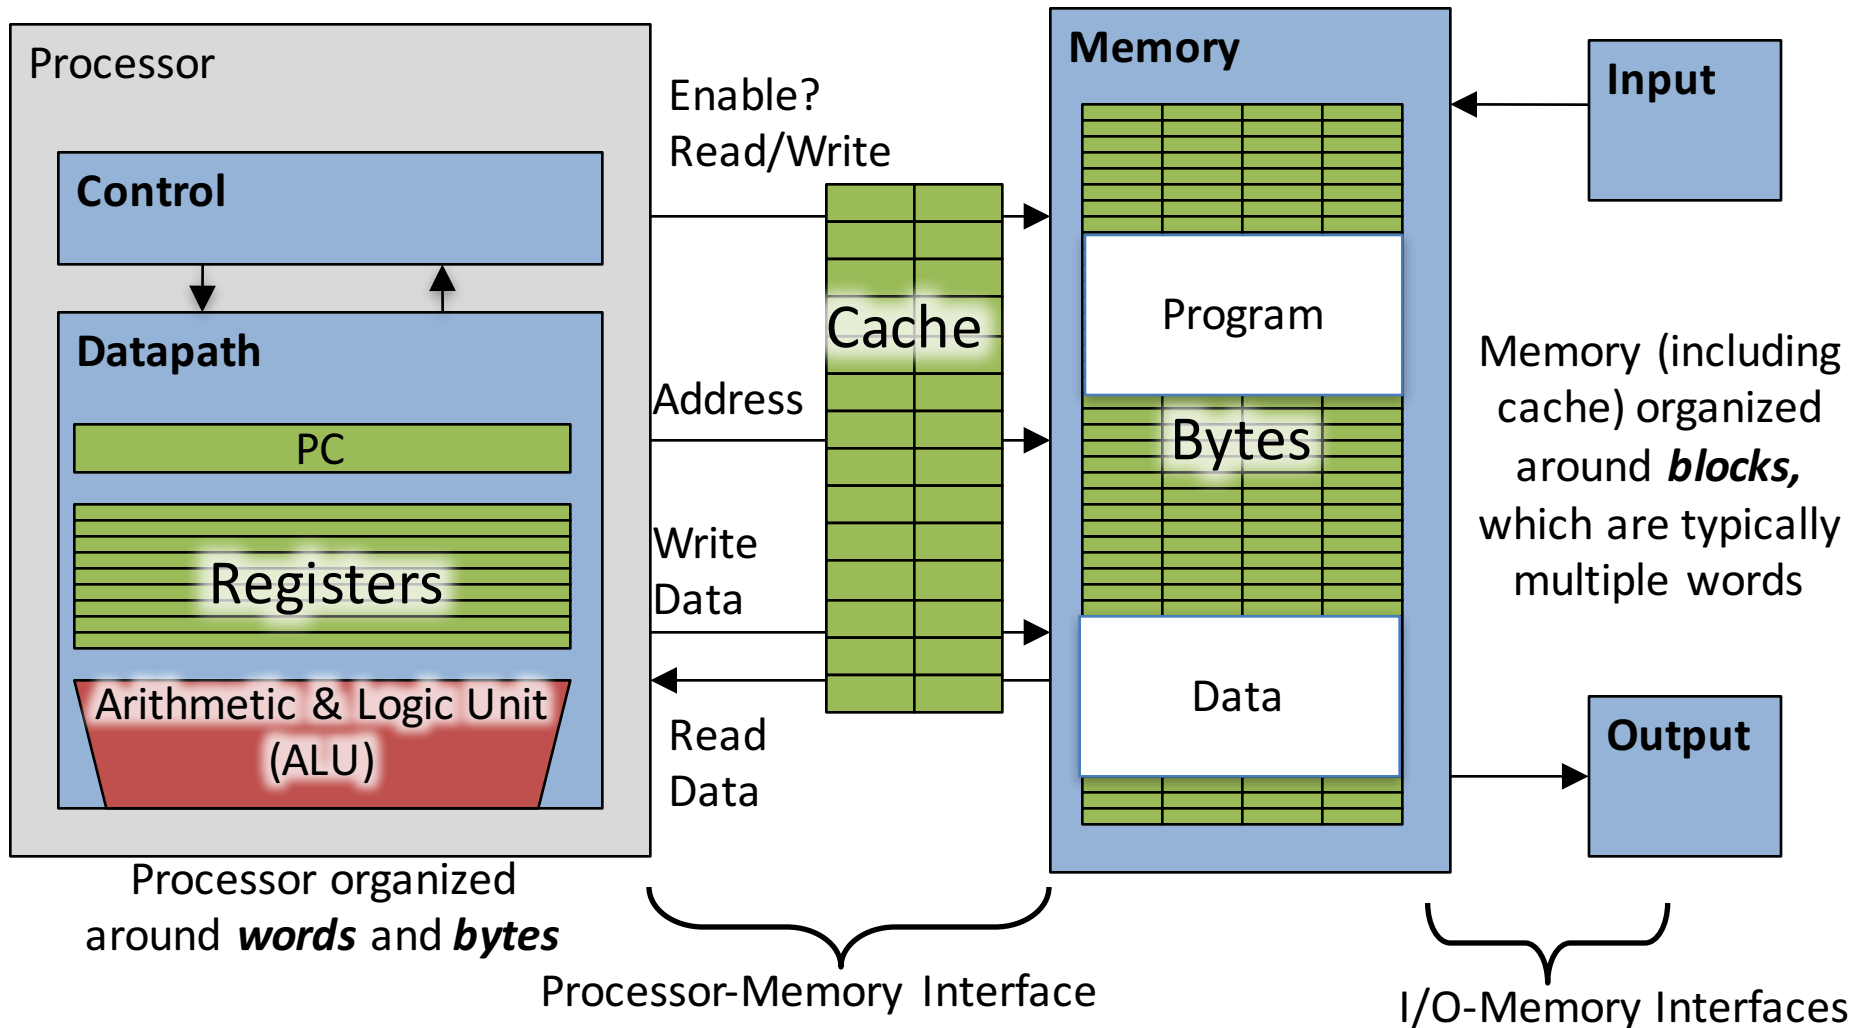
- Cache Organization and Principles
- Write Back vs. Write Through
- Cache Performance
- Cache Design Tradeoffs
- And in Conclusion ...

# Typical Memory Hierarchy



- **Principle of locality + memory hierarchy** presents programmer with  $\approx$  as much memory as is available in the *cheapest* technology at the  $\approx$  speed offered by the *fastest* technology

# Adding Cache to Computer



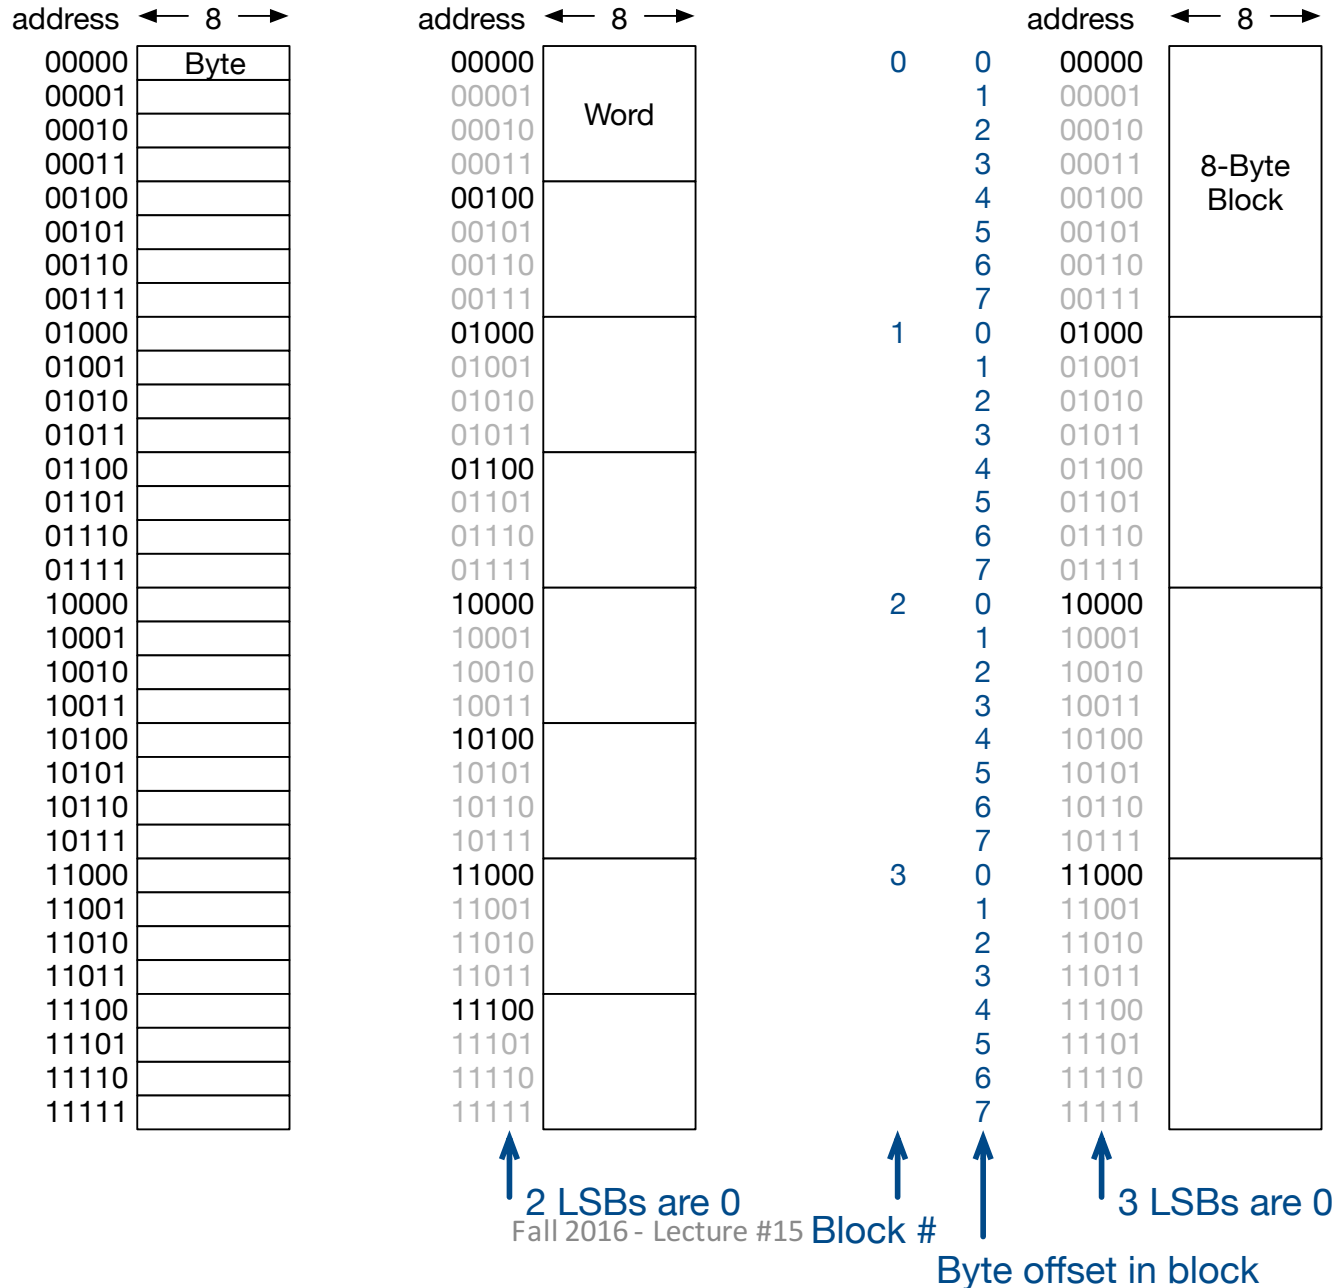
# Key Cache Concepts

- Principle of Locality
  - Temporal Locality and Spatial Locality
- Hierarchy of Memories (speed/size/cost per bit) to exploit locality
- Cache – copy of data in lower level of memory hierarchy
- Direct Mapped to find block in cache using Tag field and Valid bit for Hit
- Cache Design Organization Choices:
  - Fully Associative, Set-Associative, Direct-Mapped

# Cache Organizations

- “Fully Associative”: Block placed anywhere in cache
  - First design last lecture
  - Note: No Index field, but one comparator/block
- “Direct Mapped”: Block goes *only one place* in cache
  - Note: Only one comparator
  - Number of sets = number blocks
- “N-way Set Associative”: N places for block in cache
  - Number of sets = Number of Blocks / N
  - N comparators
  - ***Fully Associative:  $N = \text{number of blocks}$***
  - ***Direct Mapped:  $N = 1$***

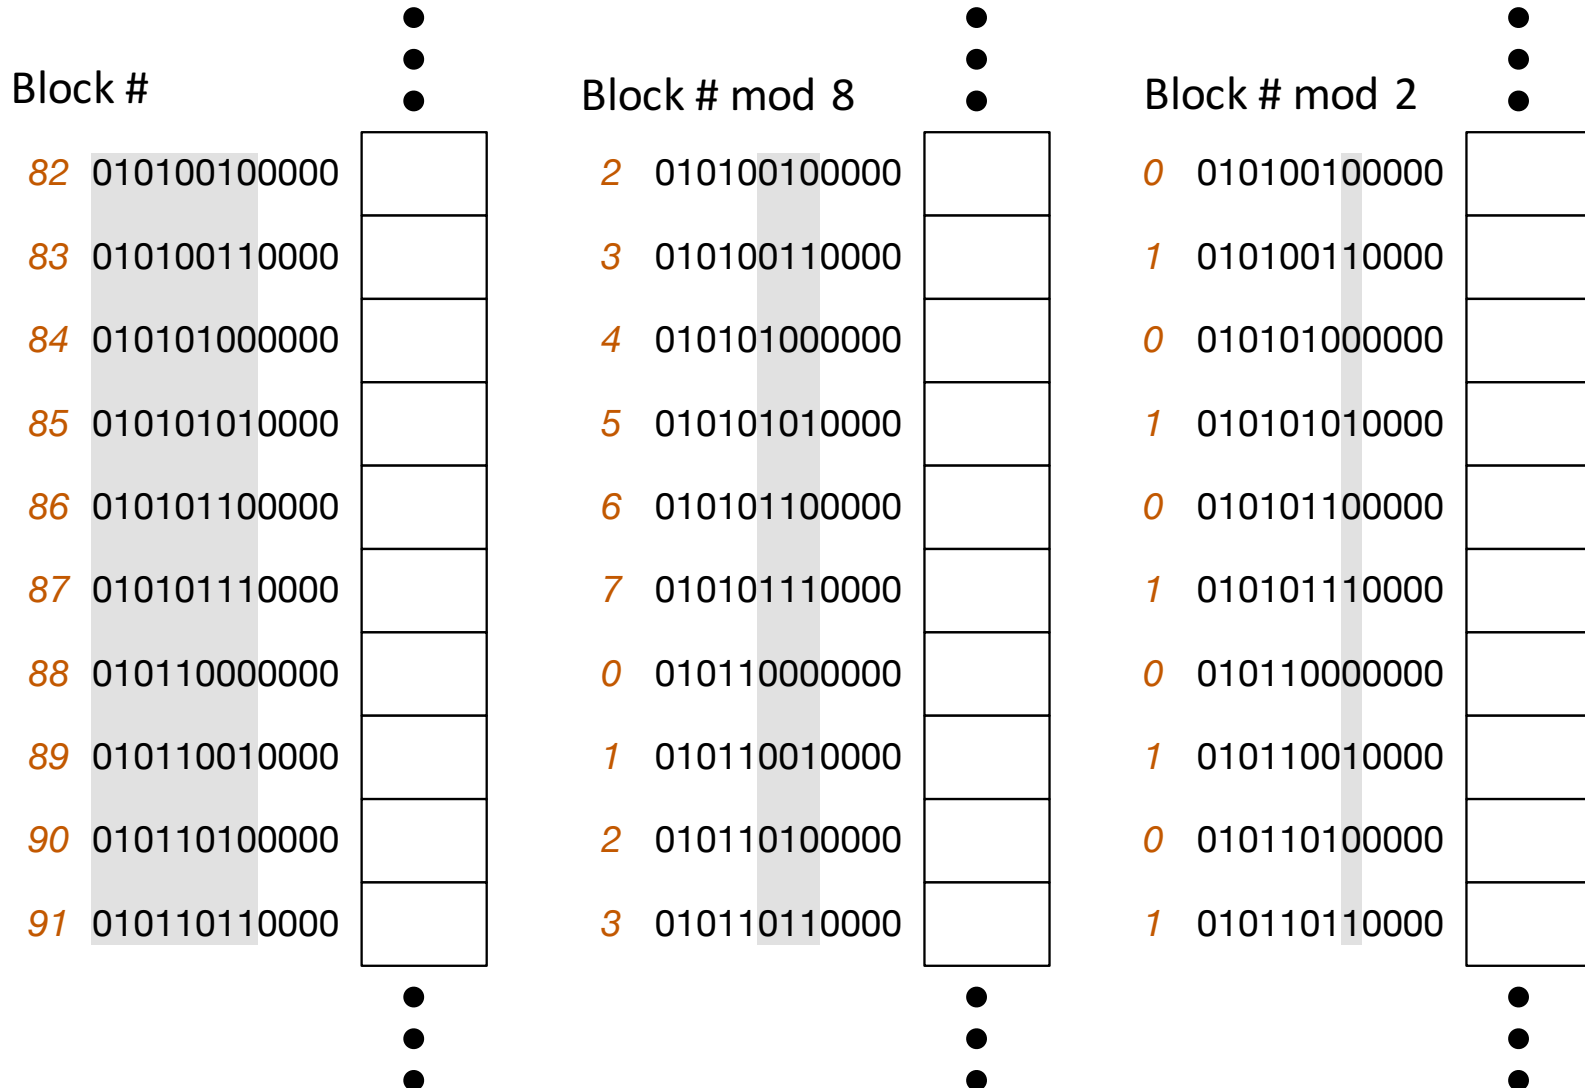
# Memory Block vs. Word Addressing






# Memory Block Number Aliasing

*12-bit memory addresses, 16 Byte blocks*



# Processor Address Fields used by Cache Controller

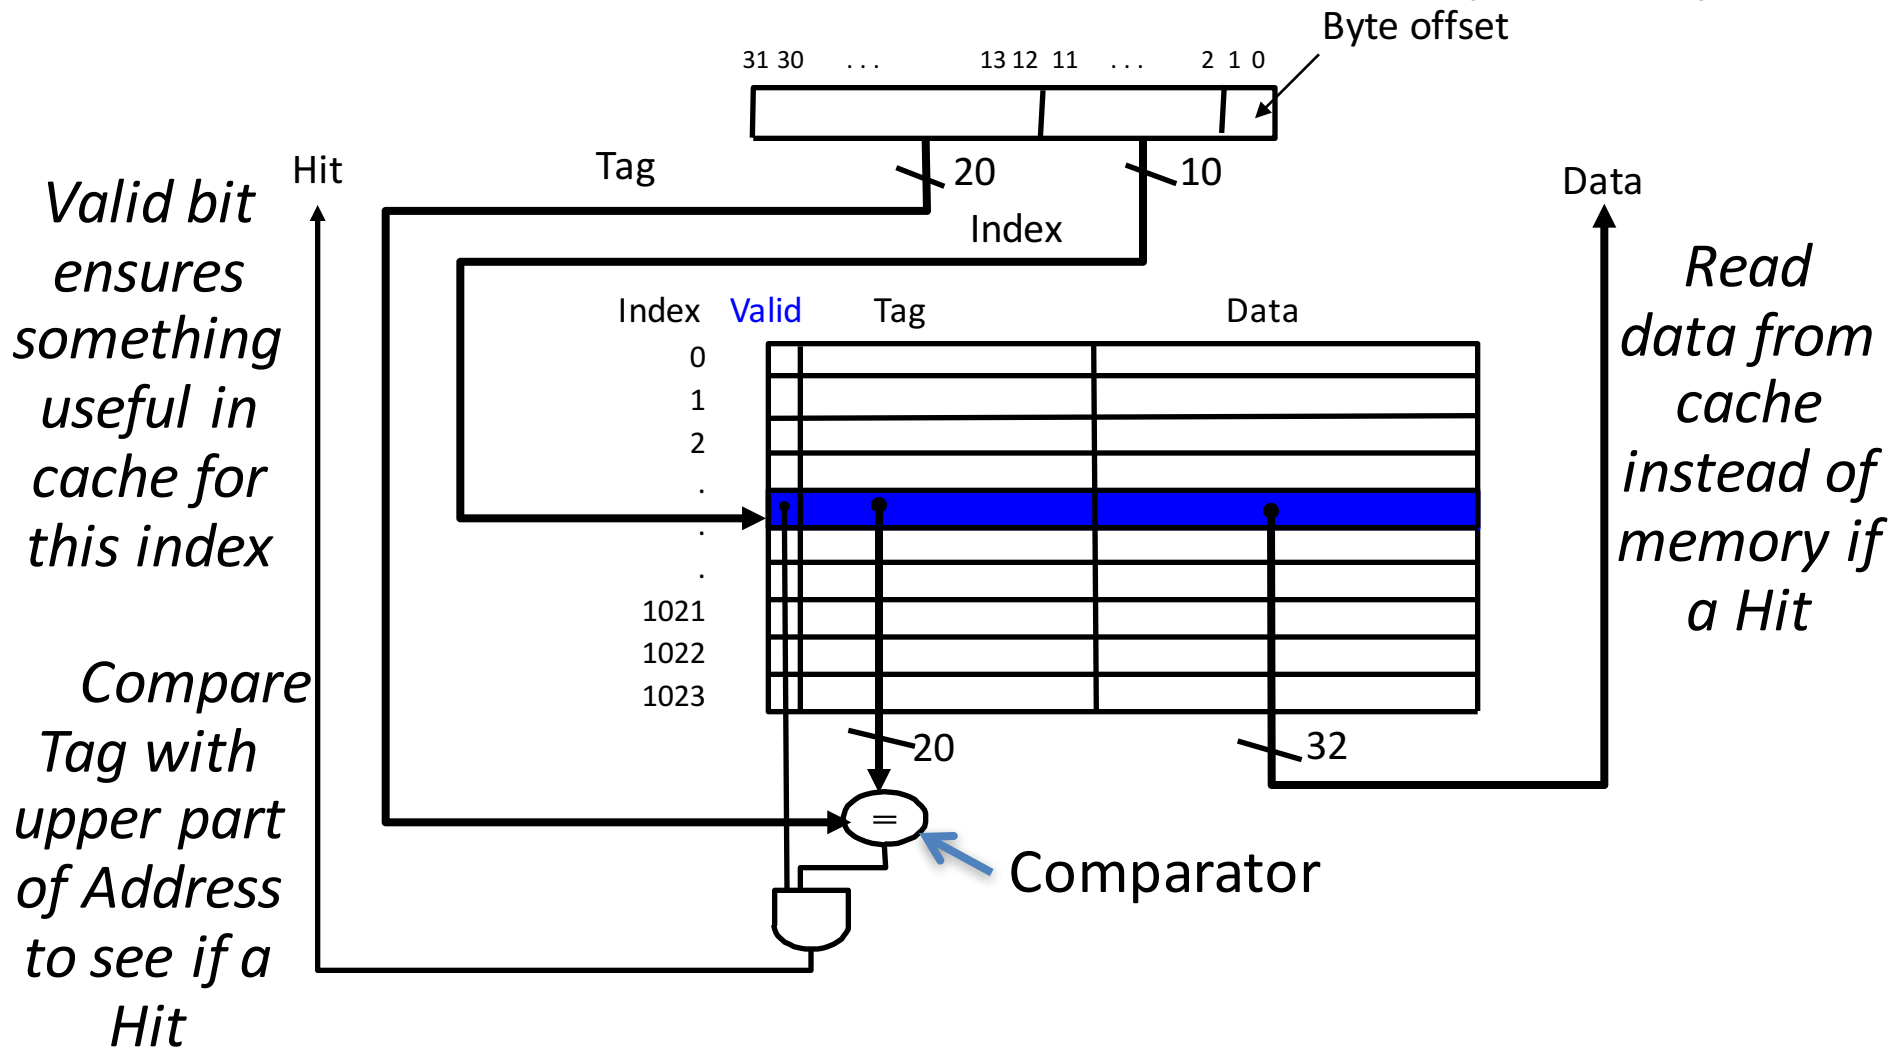
- **Block Offset**: Byte address within block
- **Set Index**: Selects which set
- **Tag**: Remaining portion of processor address



- Size of Index =  $\log_2(\text{number of sets})$
- Size of Tag = Address size – Size of Index –  $\log_2(\text{number of bytes/block})$

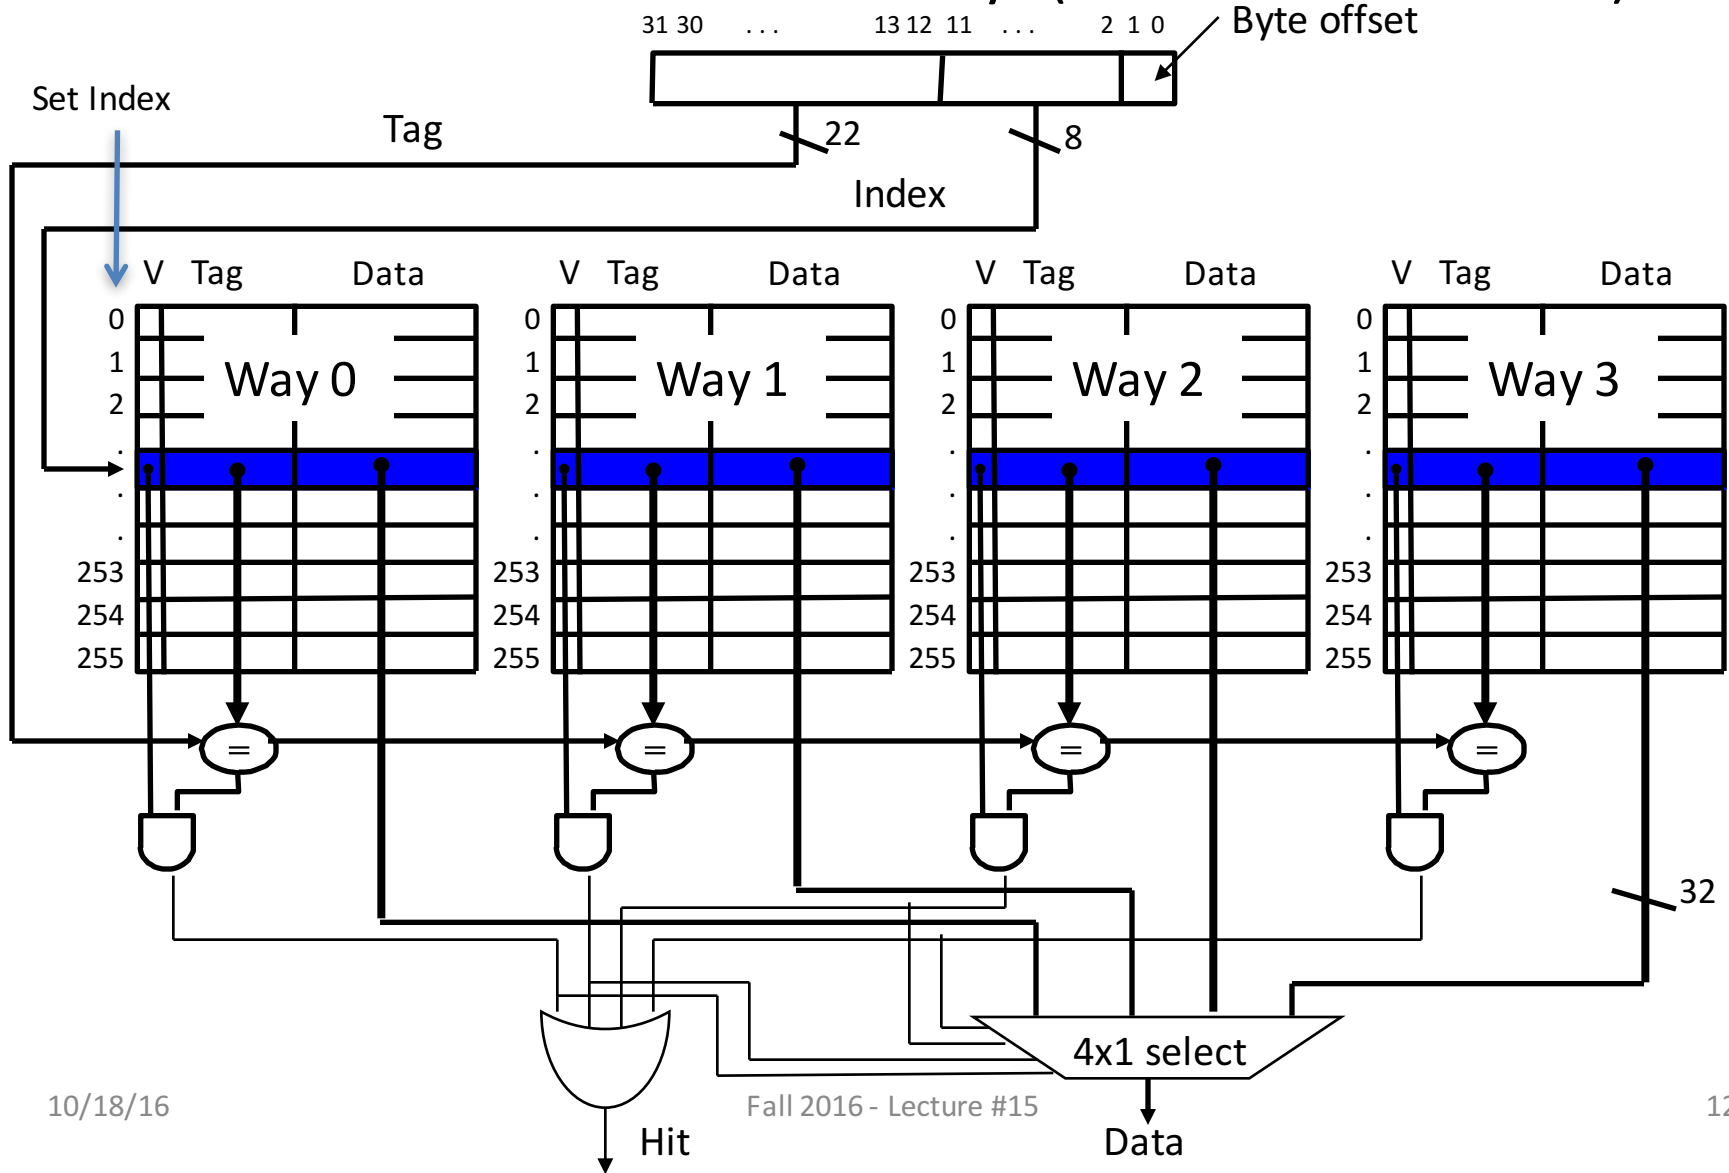
# Direct-Mapped Cache Revisted

- One word blocks, cache size = 1K words (or 4KB)



# Four-Way Set-Associative Cache

- $2^8 = 256$  sets each with four ways (each with one block)



# Outline

- Cache Organization and Principles
- **Write Back vs. Write Through**
- Cache Performance
- Cache Design Tradeoffs
- And in Conclusion ...

# Handling Stores with Write-Through

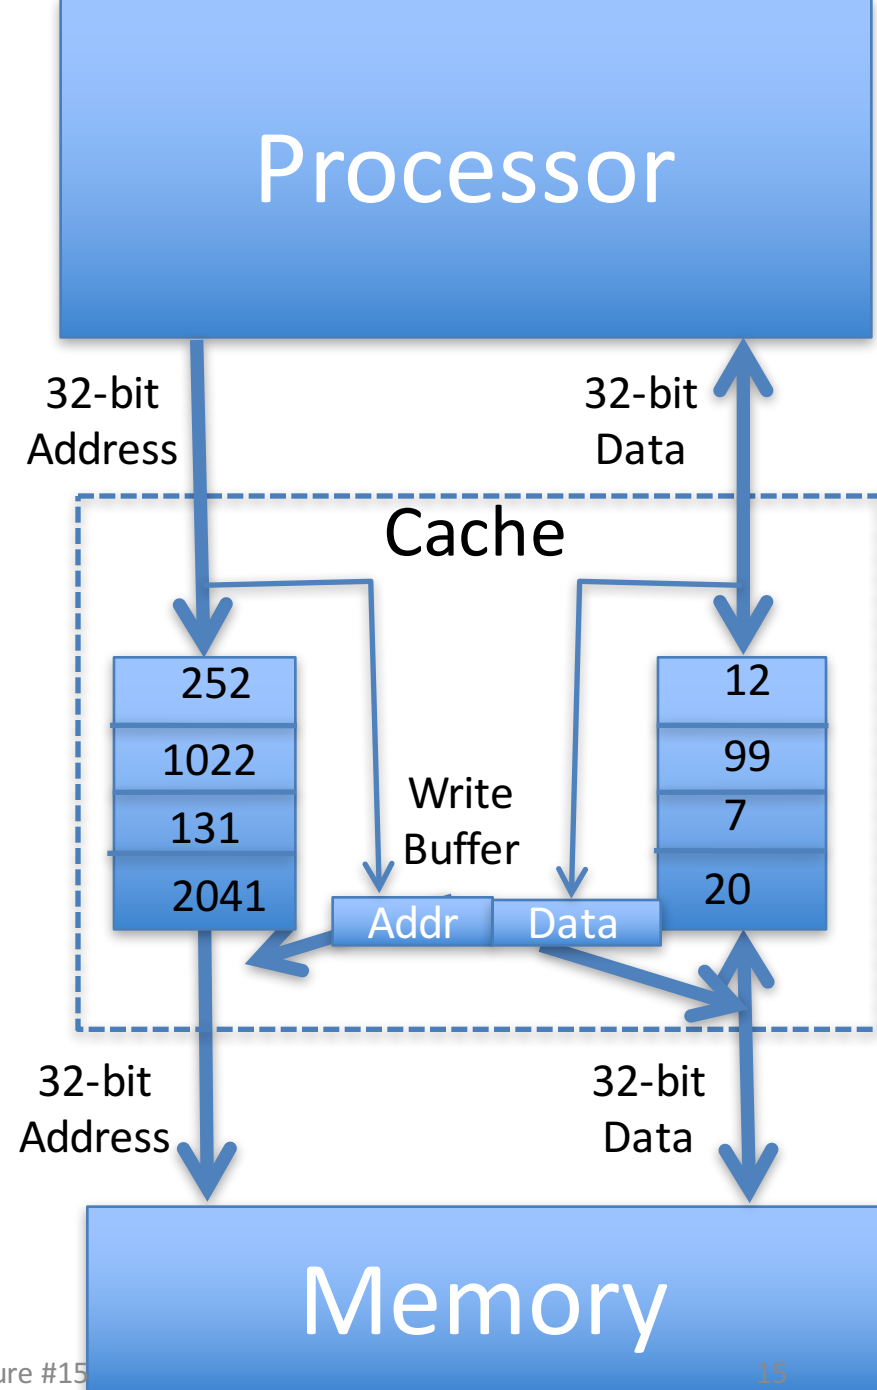
- Store instructions write to memory, changing values
- Need to make sure cache and memory have same values on writes: two policies

1) **Write-Through Policy**: write cache and write *through* the cache to memory

- Every write eventually gets to memory
- Too slow, so include Write Buffer to allow processor to continue once data in Buffer
- Buffer updates memory in parallel to processor

# Write-Through Cache

- Write both values in cache and in memory
- Write buffer stops CPU from stalling if memory cannot keep up
- Write buffer may have multiple entries to absorb bursts of writes
- What if store misses in cache?



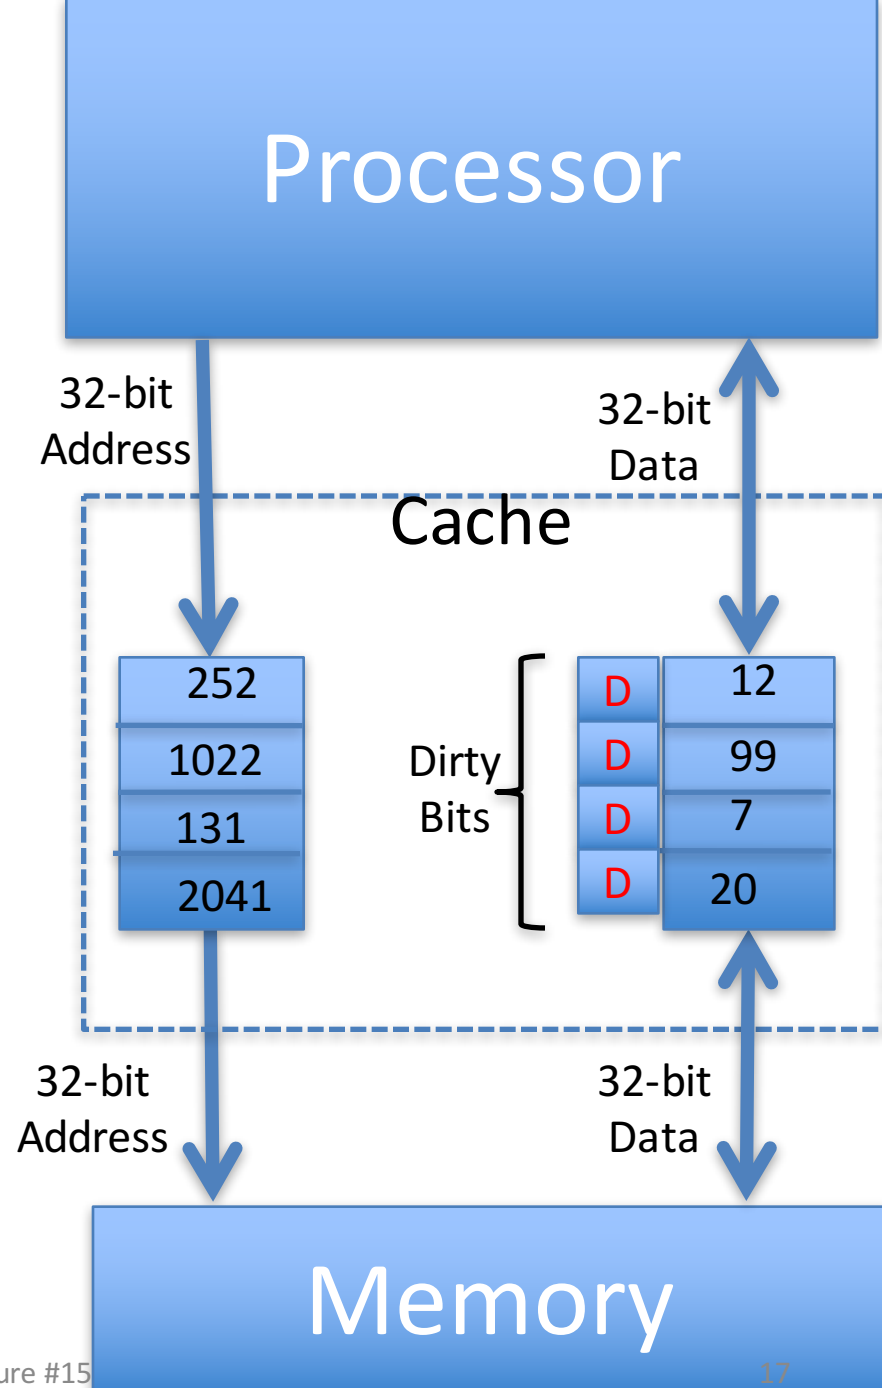
# Handling Stores with Write-Back

- 2) **Write-Back Policy**: write only to cache and then write cache block *back* to memory when evict block from cache
- Writes collected in cache, only single write to memory per block
  - Include bit to see if wrote to block or not, and then only write back if bit is set
    - Called “**Dirty**” bit (writing makes it “dirty”)



# Write-Back Cache

- Store/cache hit, write data in cache *only* and set dirty bit
  - Memory has stale value
- Store/cache miss, read data from memory, then update and set dirty bit
  - “Write-allocate” policy
- Load/cache hit, use value from cache
- On any miss, write back evicted block, only if dirty. Update cache with new block and clear dirty bit



# Write-Through vs. Write-Back

- Write-Through:
  - Simpler control logic
  - More predictable timing  
simplifies processor control logic
  - Easier to make reliable, since memory always has copy of data (big idea: Redundancy!)
- Write-Back
  - More complex control logic
  - More variable timing (0,1,2 memory accesses per cache access)
  - Usually reduces write traffic
  - Harder to make reliable, sometimes cache has only copy of data

# Administrivia

- Midterm #2 2 weeks away!  
November 1!
  - In class! 3:40-5 PM
  - Synchronous digital design and Project 3 (processor design) included
  - Pipelines and Caches
  - ONE Double sided Crib sheet
  - Review Session, Sunday, 10/30, 1-3 PM, 10 Evans



# iClicker Saga



# iClicker and EPA

- No longer taking attendance in lecture – but we hope you will continue to come anyway
- Continue to use Clicker questions in lecture to help you test your understanding
- E**P**A will be based on a “holistic” assessment of lecture, piazza, guerrilla and tutoring sessions, office hours, discussion, and lab **participation**
- EPA will be calculated so as to only help your course grade, never hurt it

# Outline

- Cache Organization and Principles
- Write Back vs. Write Through
- **Cache Performance**
- Cache Design Tradeoffs
- And in Conclusion ...

# Cache (*Performance*) Terms

- **Hit rate**: fraction of accesses that hit in the cache
- **Miss rate**:  $1 - \text{Hit rate}$
- **Miss penalty**: time to replace a block from lower level in memory hierarchy to cache
- **Hit time**: time to access cache memory (including tag comparison)
- Abbreviation: “\$” = cache (a Berkeley innovation!)

# Average Memory Access Time (AMAT)

- Average Memory Access Time (AMAT) is the average time to access memory considering both hits and misses in the cache

$$\text{AMAT} = \text{Time for a hit} + \text{Miss rate} \times \text{Miss penalty}$$

Important Equation!



# Clickers/Peer Instruction

**AMAT = Time for a hit + Miss rate x Miss penalty**

- Given a 200 psec clock, a miss penalty of 50 clock cycles, a miss rate of 0.02 misses per instruction and a cache hit time of 1 clock cycle, what is AMAT?
  - A:  $\leq 200$  psec
  - B: 400 psec
  - C: 600 psec
  - D:  $\geq 800$  psec

# Clickers/Peer Instruction

**AMAT = Time for a hit + Miss rate x Miss penalty**

- Given a 200 psec clock, a miss penalty of 50 clock cycles, a miss rate of 0.02 misses per instruction and a cache hit time of 1 clock cycle, what is AMAT?

A:  $\leq 200$  psec

**B: 400 psec**

C: 600 psec

D:  $\geq 800$  psec

1 clock cycle + .02 \* 50 clock cycles = 2 clock cycles

# Ping Pong Cache Example: Direct-Mapped Cache w/4 Single-Word Blocks, Worst-Case Reference String

- Consider the main memory address reference string of word numbers:

0 4 0 4 0 4 0 4

Start with an empty cache - all blocks initially marked as not valid

**0**


**4**


**0**


**4**


**0**


**4**

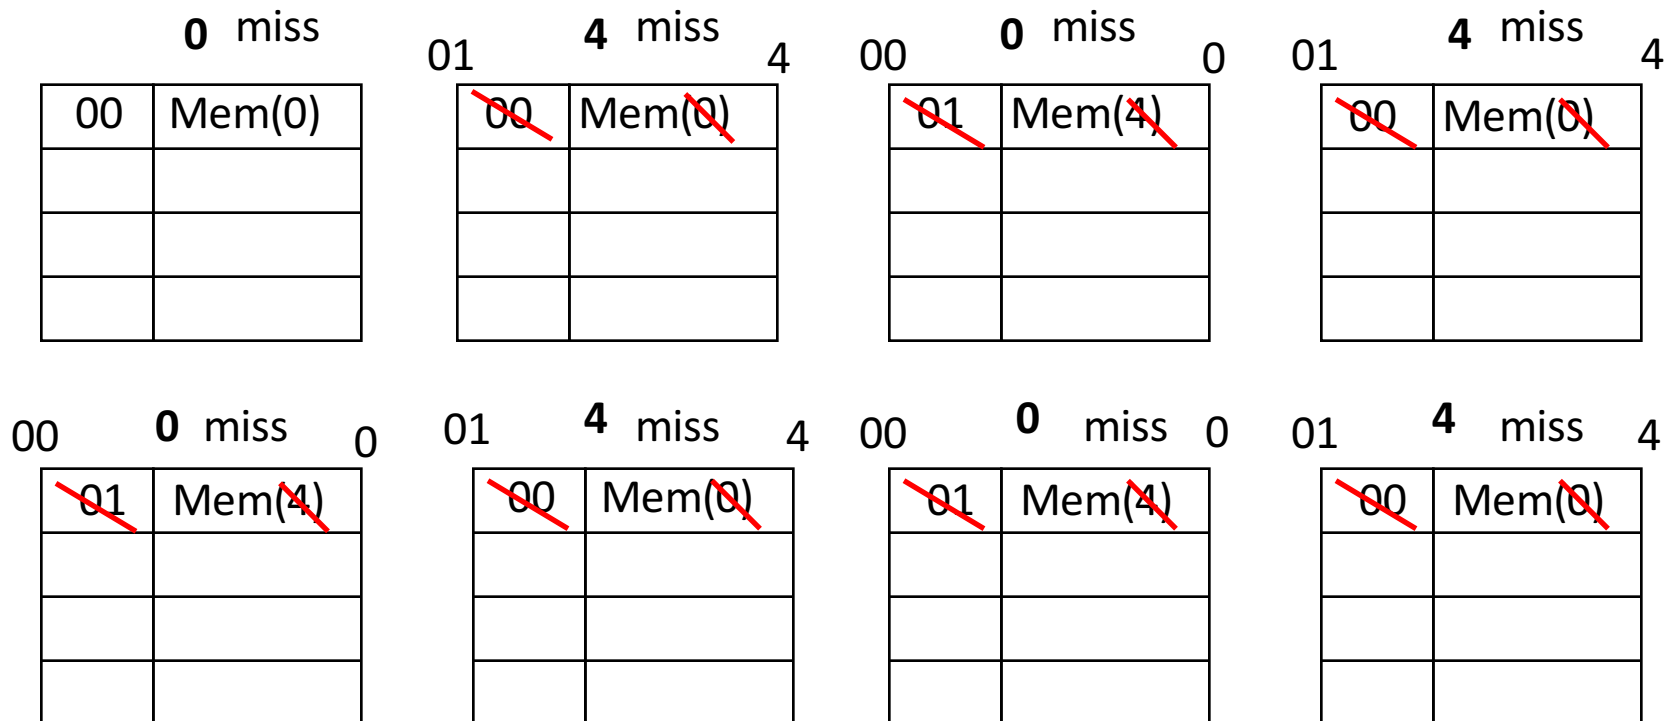

**0**


**4**


# Ping Pong Cache Example: Direct-Mapped Cache w/4 Single-Word Blocks, Worst-Case Reference String

- Consider the main memory address reference string of word numbers:  
0 4 0 4 0 4 0 4

Start with an empty cache - all blocks initially marked as not valid



- 8 requests, 8 misses


Ping-pong effect due to conflict misses - two memory locations that map into the same cache block

# Outline


- Cache Organization and Principles
- Write Back vs. Write Through
- Cache Performance
- **Cache Design Tradeoffs**
- And in Conclusion ...

# Alternative Block Placement Schemes


**Direct mapped**



**Set associative**



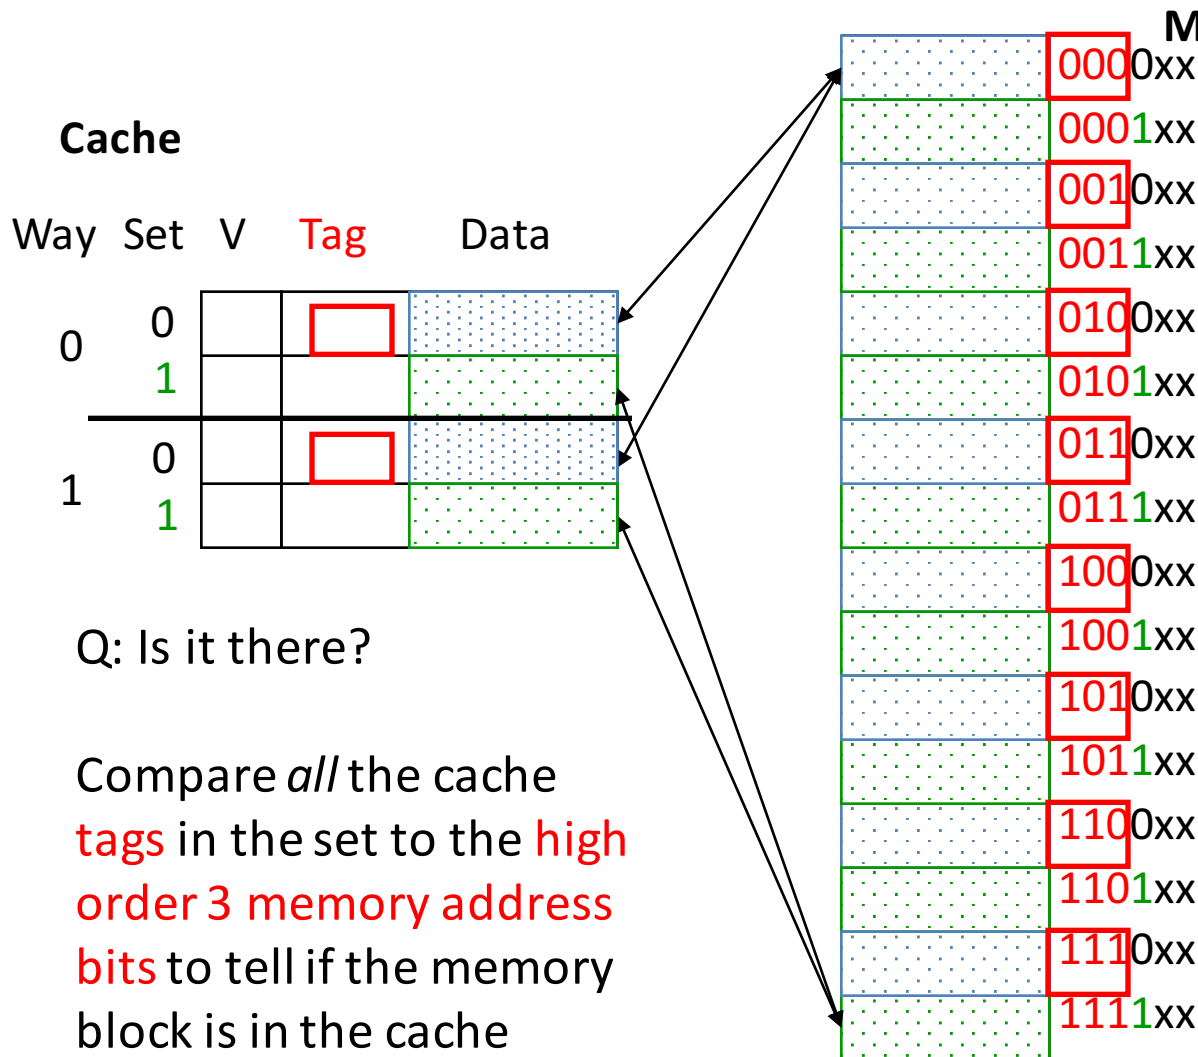
**Fully associative**



- DM placement: mem block 12 in 8 block cache: only one cache block where mem block 12 can be found— $(12 \text{ modulo } 8) = 4$
- SA placement: four sets x 2-ways (8 cache blocks), memory block 12 in set  $(12 \text{ mod } 4) = 0$ ; either element of the set
- FA placement: mem block 12 can appear in any cache blocks

# Example: 2-Way Set Associative \$

(4 words = 2 sets x 2 ways per set)



One word blocks  
Two low order bits  
define the byte in the  
word (32b words)


Q: How do we find it?

Use next 1 low order  
memory address bit to  
determine which cache  
set (i.e., modulo the  
number of sets in the  
cache)

# Ping Pong Cache Example: 4 Word 2-Way SA \$, Same Reference String

- Consider the main memory word reference string

Start with an empty cache - all blocks  
initially marked as not valid      0   4   0   4   0   4   0   4





# Ping Pong Cache Example: 4-Word 2-Way SA \$, Same Reference String

- Consider the main memory address reference string

Start with an empty cache - all blocks initially marked as not valid      0 4 0 4 0 4 0 4

0 miss

000	Mem(0)

4 miss

000	Mem(0)
010	Mem(4)

0 hit

000	Mem(0)
010	Mem(4)

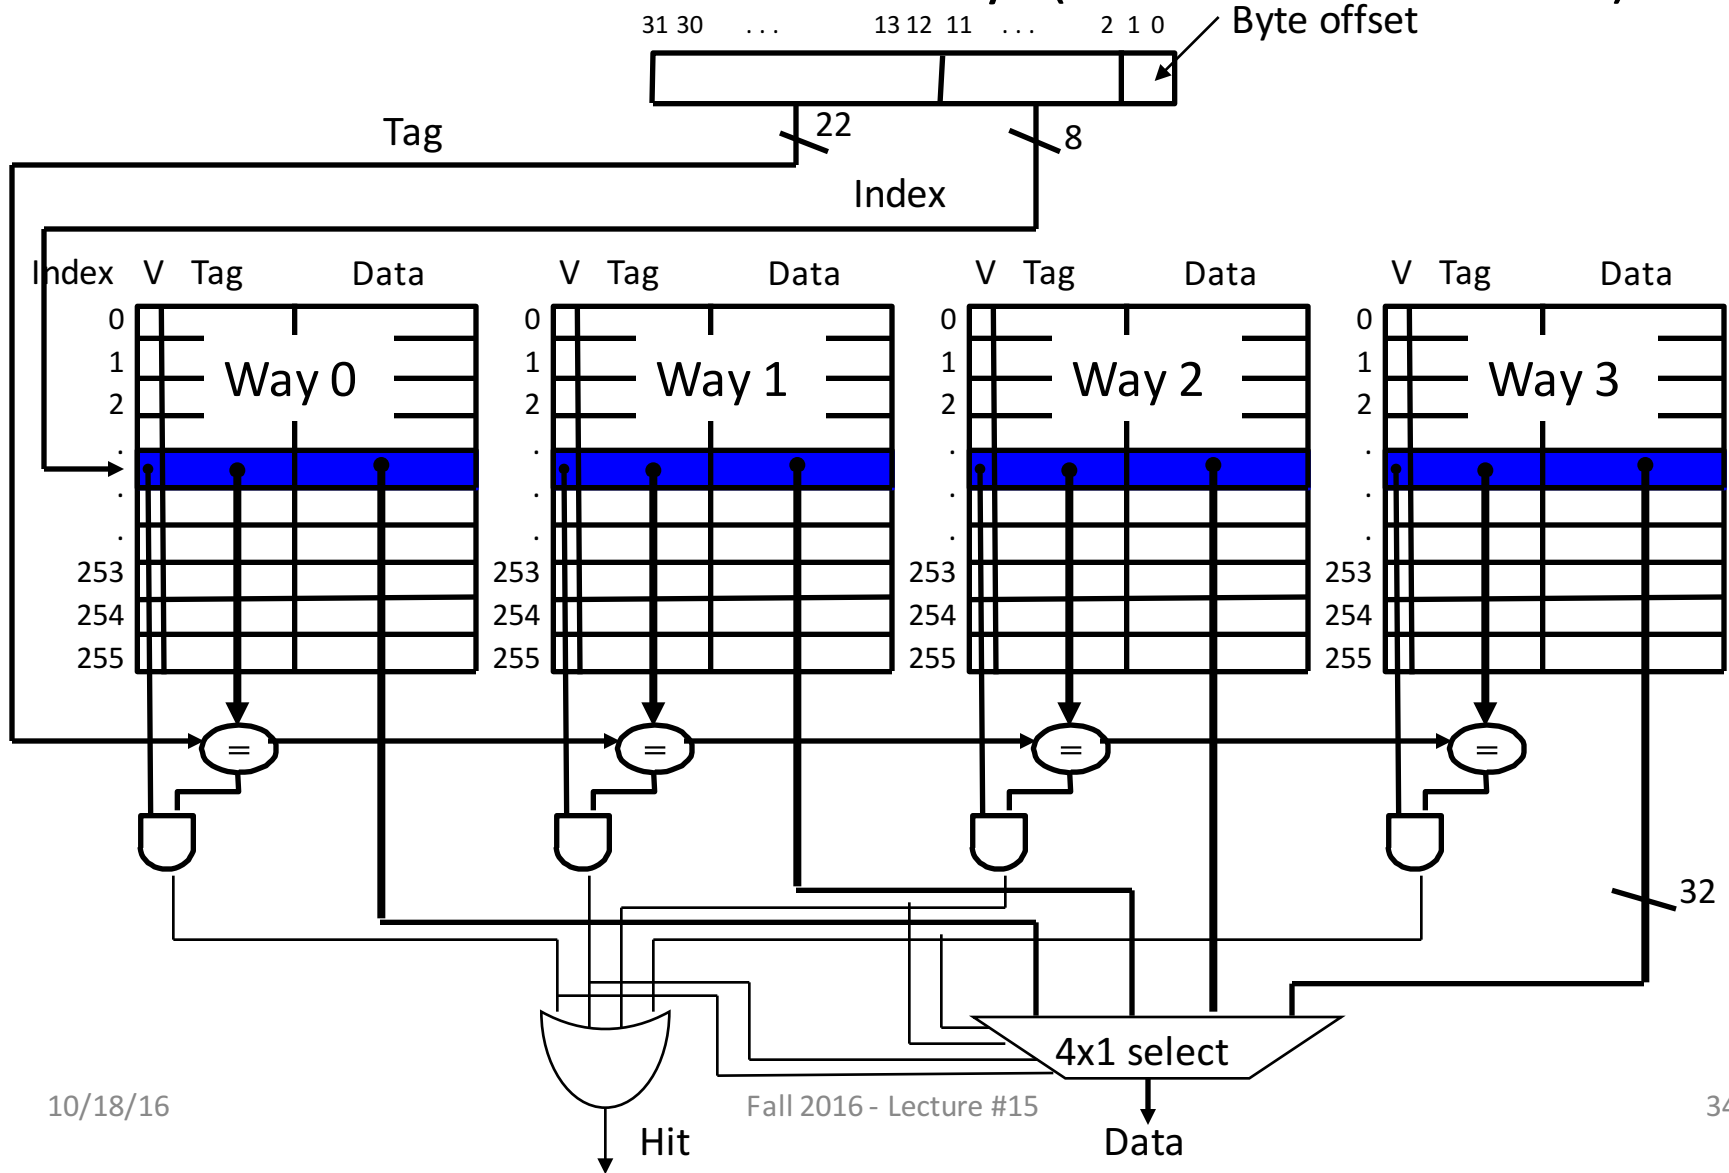
4 hit

000	Mem(0)
010	Mem(4)

- 8 requests, 2 misses
- Solves the ping-pong effect in a direct-mapped cache due to conflict misses since now two memory locations that map into the same cache set can co-exist!

# Four-Way Set-Associative Cache

- $2^8 = 256$  sets each with four ways (each with one block)



# Alternative Organizations of an Eight-Block Cache

## One-way set associative (direct mapped)

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

## Two-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

## Four-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

## Eight-way set associative (fully associative)

Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data

Total size of \$ in blocks is equal to *number of sets*  $\times$  *associativity*. For fixed \$ size and fixed block size, increasing associativity decreases number of sets while increasing number of elements per set. With eight blocks, an 8-way set-associative \$ is same as a fully associative \$.

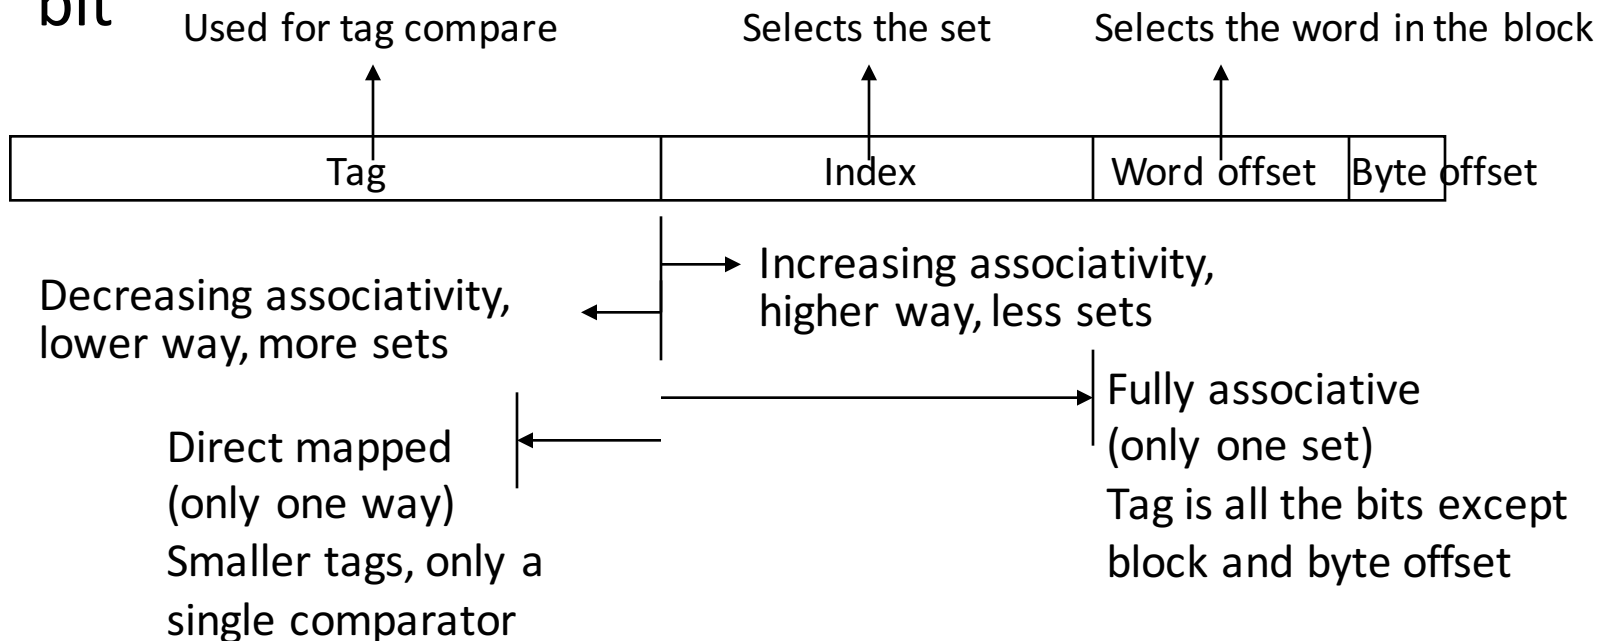
# Range of Set-Associative Caches

- For a *fixed-size* cache and fixed block size, each increase by a factor of two in associativity doubles the number of blocks per set (i.e., the number or ways) and halves the number of sets – decreases the size of the index by 1 bit and increases the size of the tag by 1 bit

Tag	Index	Word offset	Byte offset
-----	-------	-------------	-------------

# Range of Set-Associative Caches

- For a *fixed-size* cache and fixed block size, each increase by a factor of two in associativity doubles the number of blocks per set (i.e., the number or ways) and halves the number of sets – decreases the size of the index by 1 bit and increases the size of the tag by 1 bit



# Total Cache Capacity =

Associativity  $\times$  # of sets  $\times$  block\_size

*Bytes = blocks/set  $\times$  sets  $\times$  Bytes/block*

$$C = N \times S \times B$$

<i>Tag</i>	<i>Index</i>	<i>Byte Offset</i>
------------	--------------	--------------------

$$\begin{aligned}\text{address\_size} &= \text{tag\_size} + \text{index\_size} + \text{offset\_size} \\ &= \text{tag\_size} + \log_2(S) + \log_2(B)\end{aligned}$$

# Total Cache Capacity =

Associativity \* # of sets \* block\_size

*Bytes = blocks/set \* sets \* Bytes/block*

$$C = N * S * B$$

<i>Tag</i>	<i>Index</i>	<i>Byte Offset</i>
------------	--------------	--------------------

address\_size = tag\_size + index\_size + offset\_size

= tag\_size +  $\log_2(S)$  +  $\log_2(B)$

Double the Associativity: Number of sets?

*tag\_size? index\_size? # comparators?*

Double the Sets: Associativity?

*tag\_size? index\_size? # comparators?*

# Total Cache Capacity =

Associativity \* # of sets \* block\_size

*Bytes = blocks/set \* sets \* Bytes/block*

$$C = N * S * B$$

<i>Tag</i>	<i>Index</i>	<i>Byte Offset</i>
------------	--------------	--------------------

address\_size = tag\_size + index\_size + offset\_size

$$= \text{tag\_size} + \log_2(S) + \log_2(B)$$

Double the Associativity: Halve the number of sets

*tag\_size + 1 while index\_size - 1, 2 x comparators*

Double the Sets: Halve the associativity

*tag\_size - 1 while index\_size + 1, 1/2 x comparators*



# Your Turn

- For a cache of 64 blocks, each block four bytes in size:
  1. The capacity of the cache is: 256 bytes.
  2. Given a 2-way Set Associative organization, there are 32 sets, each of 2 blocks, and 2 places a block from memory could be placed.
  3. Given a 4-way Set Associative organization, there are 16 sets each of 4 blocks and 4 places a block from memory could be placed.
  4. Given an 8-way Set Associative organization, there are 8 sets each of 8 blocks and 8 places a block from memory could be placed.

# Clicker/Peer Instruction

- For S sets, N ways, B blocks, which statements hold?
  - (i) The cache has B tags
  - (ii) The cache needs N comparators
  - (iii)  $B = N \times S$
  - (iv) Size of Index =  $\log_2(S)$

A: (i) only

B: (i) and (ii) only

C: (i), (ii), (iii) only

D: All four statements are true


E: None are true

# Costs of Set-Associative Caches

- N-way set-associative cache costs
  - N comparators (delay and area)
  - MUX delay (set selection) before data is available
  - Data available after set selection (and Hit/Miss decision).  
DM \$: block is available before the Hit/Miss decision
    - In Set-Associative, not possible to just assume a hit and continue and recover later if it was a miss
- When miss occurs, which way's block selected for replacement?
  - **Least Recently Used** (LRU): one that has been unused the longest (principle of temporal locality)
    - Must track when each way's block was used relative to other blocks in the set
    - For 2-way SA \$, one bit per set → set to 1 when a block is referenced; reset the other way's bit (i.e., "last used")


# Cache Replacement Policies

- Random Replacement
  - Hardware randomly selects a cache evict
- Least-Recently Used
  - Hardware keeps track of access history
  - Replace the entry that has not been used for the longest time
  - For 2-way set-associative cache, need one bit for LRU replacement
- Example of a Simple “Pseudo” LRU Implementation
  - Assume 64 Fully Associative entries
  - Hardware replacement pointer points to one cache entry
  - Whenever access is made to the entry the pointer points to:
    - Move the pointer to the next entry
  - Otherwise: do not move the pointer
  - (example of “not-most-recently used” replacement policy)



# Benefits of Set-Associative Caches

- Choice of DM \$ versus SA \$ depends on the cost of a miss versus the cost of implementation



- Largest gains are in going from direct mapped to 2-way (20%+ reduction in miss rate)

# Outline

- Cache Organization and Principles
- Write Back vs. Write Through
- Cache Performance
- Cache Design Tradeoffs
- **And in Conclusion ...**

## And in Conclusion ...

- Name of the Game: Reduce AMAT
  - Reduce Hit Time
  - Reduce Miss Rate
  - Reduce Miss Penalty
- Balance cache parameters (Capacity, associativity, block size)