

CS 61C:
Great Ideas in Computer Architecture

Lecture 18:
Parallel Processing – SIMD

Bernhard Boser & Randy Katz

<http://inst.eecs.berkeley.edu/~cs61c>

61C Survey

It would be nice to have a review lecture every once in a while, actually showing us how things fit in the bigger picture

Agenda

- **61C – the big picture**
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

61C Topics so far ...

- What we learned:
 1. Binary numbers
 2. C
 3. Pointers
 4. Assembly language
 5. Datapath architecture
 6. Pipelining
 7. Caches
 8. Performance evaluation
 9. Floating point
- What does this buy us?
 - Promise: execution speed
 - Let's check!

Reference Problem

- **Matrix multiplication**
 - Basic operation in many engineering, data, and imaging processing tasks
 - Image filtering, noise reduction, ...
 - Many closely related operations
 - E.g. stereo vision (project 4)
- **dgemm**
 - double precision floating point matrix multiplication

Application Example: Deep Learning


- Image classification (cats ...)
- Pick “best” vacation photos
- Machine translation
- Clean up accent
- Fingerprint verification
- Automatic game playing

Matrices

- Square (or rectangular) $N \times N$ array of numbers
 - Dimension N


$$C = A \cdot B$$

$$c_{ij} = \sum_k a_{ik} b_{kj}$$



Matrix Multiplication

$$\mathbf{C} = \mathbf{A} \cdot \mathbf{B}$$
$$c_{ij} = \sum_k a_{ik} b_{kj}$$



Reference: Python

- Matrix multiplication in Python

```
def dgemm(N, a, b, c):  
    for i in range(N):  
        for j in range(N):  
            c[i+j*N] = 0  
            for k in range(N):  
                c[i+j*N] += a[i+k*N] * b[k+j*N]
```

| N | Python [Mflops] |
|-----|-----------------|
| 32 | 5.4 |
| 160 | 5.5 |
| 480 | 5.4 |
| 960 | 5.3 |

- 1 Mflop = 1 Million floating point operations per second (fadd, fmul)
- dgemm(N ...) takes $2*N^3$ flops

C

- $c = a \times b$
- a, b, c are $N \times N$ matrices

```
// Scalar; P&H p. 226
void dgemm_scalar(int N, double *a, double *b, double *c) {
    for (int i=0; i<N; i++)
        for (int j=0; j<N; j++) {
            double cij = 0;
            for (int k=0; k<N; k++)
                //      a[i][k] * b[k][j]
                cij += a[i+k*N] * b[k+j*N];
            // c[i][j]
            c[i+j*N] = cij;
        }
    }
```

Timing Program Execution

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(void) {
    // start time
    // Note: clock() measures execution time, not real time
    //       big difference in shared computer environments
    //       and with heavy system load
    clock_t start = clock();

    // task to time goes here:
    // dgemm(N, ...);

    // "stop" the timer
    clock_t end = clock();

    // compute execution time in seconds
    double delta_time = (double)(end-start)/CLOCKS_PER_SEC;
}
```

C versus Python

| N | C [G flops] | Python [Gflops] |
|-----|---------------------|-----------------|
| 32 | 1.30 | 0.0054 |
| 160 | 1.30 | 0.0055 |
| 480 | 1.32 | 0.0054 |
| 960 | 0.91 | 0.0053 |



Which class gives you this kind of power?
We could stop here ... but why? Let's do better!

Agenda

- 61C – the big picture
- **Parallel processing**
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

Why Parallel Processing?

- CPU Clock Rates are no longer increasing
 - Technical & economic challenges
 - Advanced cooling technology too expensive or impractical for most applications
 - Energy costs are prohibitive
- Parallel processing is only path to higher speed
 - Compare airlines:
 - Maximum speed limited by speed of sound and economics
 - Use more and larger airplanes to increase throughput
 - And smaller seats ...

Using Parallelism for Performance

- Two basic ways:
 - Multiprogramming
 - run multiple independent programs in parallel
 - “Easy”
 - Parallel computing
 - run one program faster
 - “Hard”
- We’ll focus on parallel computing in the next few lectures

New-School Machine Structures (It's a bit more complicated!)

Software

Hardware

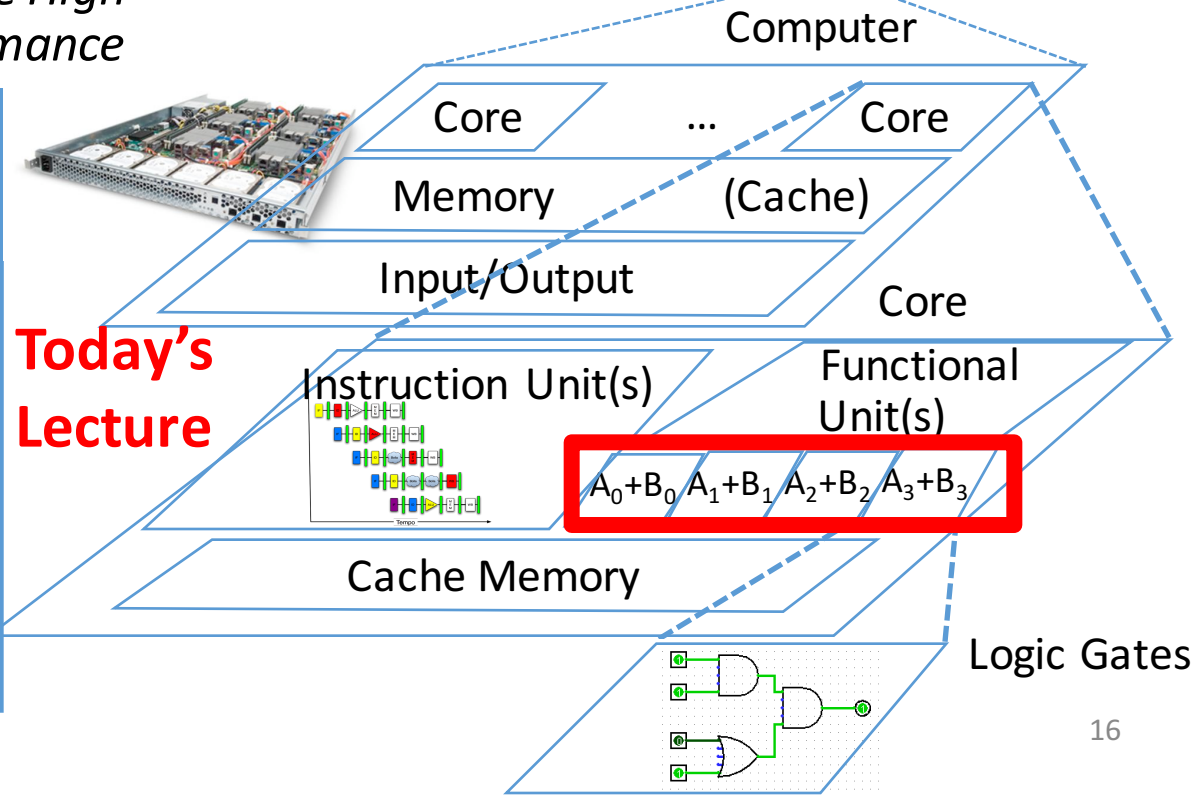
- Parallel Requests
Assigned to computer
e.g., Search “Katz”
- Parallel Threads
Assigned to core
e.g., Lookup, Ads
- Parallel Instructions
>1 instruction @ one time
e.g., 5 pipelined instructions
- Parallel Data
>1 data item @ one time
e.g., Add of 4 pairs of words
- Hardware descriptions
All gates @ one time
- Programming Languages

*Harness
Parallelism &
Achieve High
Performance*

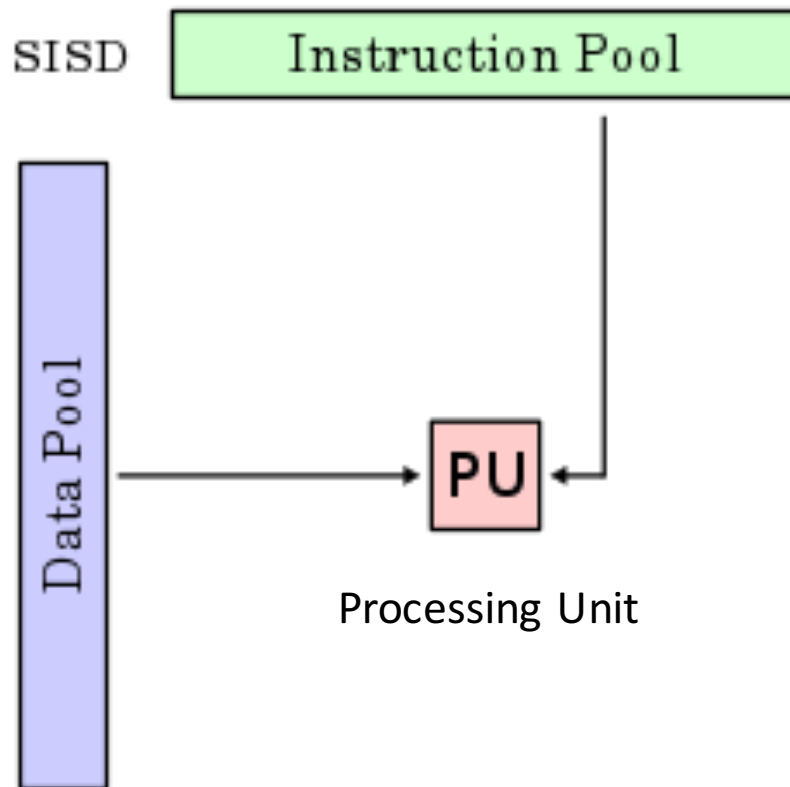
Warehouse
Scale
Computer



Smart
Phone




Single-Instruction/Single-Data Stream (SISD)



- Sequential computer that exploits no parallelism in either the instruction or data streams. Examples of SISD architecture are traditional uniprocessor machines
 - E.g. our trusted MIPS

This is what we did up to now in 61C

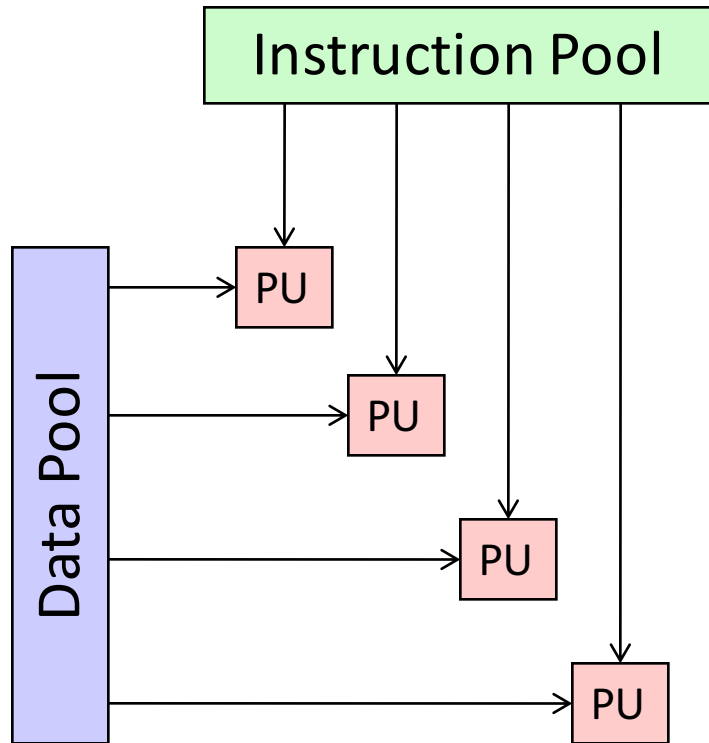
Single-Instruction/Multiple-Data Stream (SIMD or “sim-dee”)



- SIMD computer exploits multiple data streams against a single instruction stream to operations that may be naturally parallelized, e.g., Intel SIMD instruction extensions or NVIDIA Graphics Processing Unit (GPU)

Today's topic.

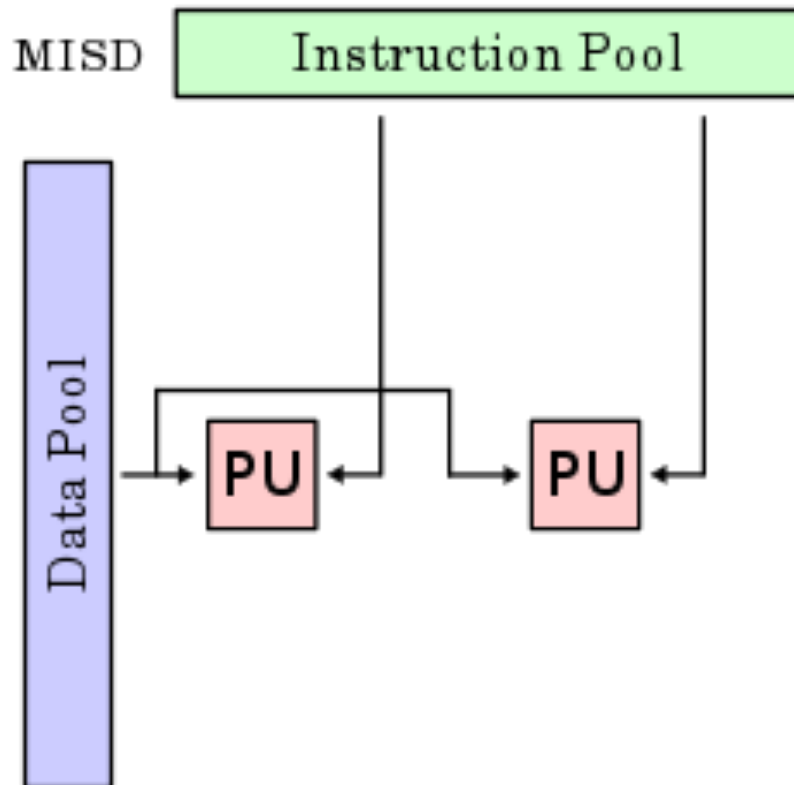
Multiple-Instruction/Multiple-Data Streams (MIMD or “mim-dee”)



- Multiple autonomous processors simultaneously executing different instructions on different data.
 - MIMD architectures include multicore and Warehouse-Scale Computers

Topic of Lecture 19 and beyond.

Multiple-Instruction/Single-Data Stream (MISD)



- Multiple-Instruction, Single-Data stream computer that exploits multiple instruction streams against a single data stream.
 - Historical significance

This has few applications. Not covered in 61C.

Flynn* Taxonomy, 1966


| | | Data Streams | |
|---------------------|----------|-------------------------|-------------------------------------|
| | | Single | Multiple |
| Instruction Streams | Single | SISD: Intel Pentium 4 | SIMD: SSE instructions of x86 |
| | Multiple | MISD: No examples today | MIMD: Intel Xeon e5345 (Clovertown) |

- SIMD and MIMD are currently the most common parallelism in architectures – usually both in same system!
- Most common parallel processing programming style: Single Program Multiple Data (“SPMD”)
 - Single program that runs on all processors of a MIMD
 - Cross-processor execution coordination using synchronization primitives

Agenda

- 61C – the big picture
- Parallel processing
- **Single instruction, multiple data**
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...


SIMD – “Single Instruction Multiple Data”



SIMD Applications & Implementations

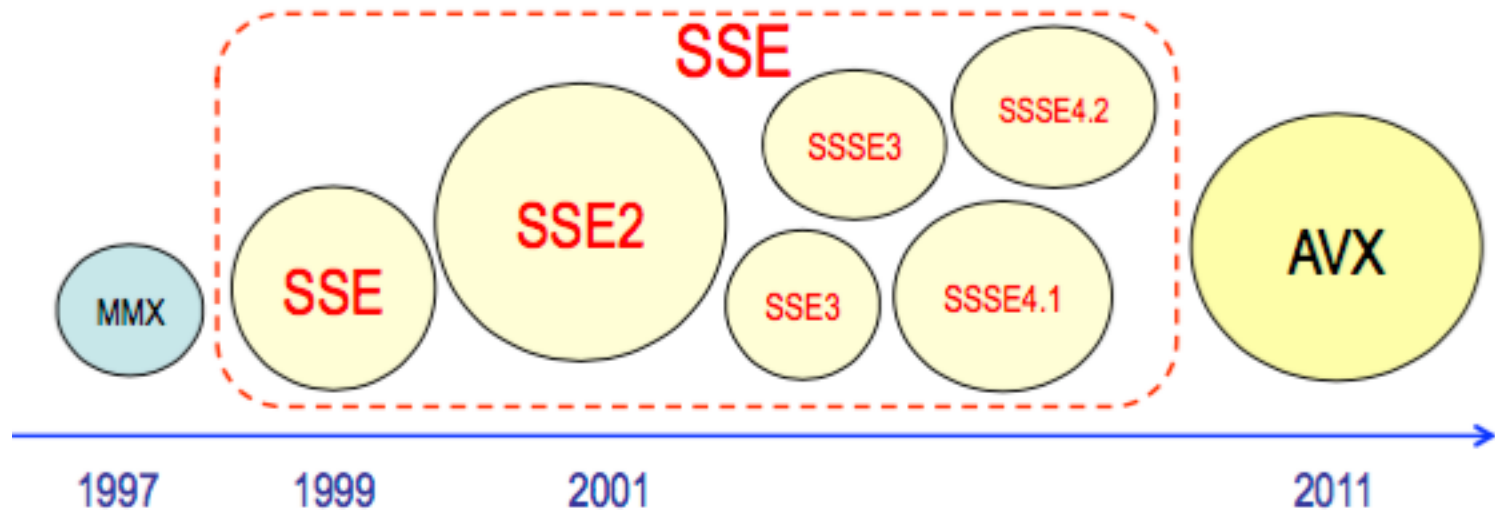
- Applications
 - Scientific computing
 - Matlab, NumPy
 - Graphics and video processing
 - Photoshop, ...
 - Big Data
 - Deep learning
 - Gaming
 - ...
- Implementations
 - x86
 - ARM
 - ...

First SIMD Extensions: MIT Lincoln Labs TX-2, 1957



x86 SIMD Evolution

- New instructions
- New, wider, more registers
- More parallelism



<http://svmoore.pbworks.com/w/file/fetch/70583970/VectorOps.pdf>

CPU Specs (Bernhard's Laptop)

```
$ sysctl -a | grep cpu
```

```
hw.physicalcpu: 2
```

```
hw.logicalcpu: 4
```

```
machdep.cpu.brand_string:
```

```
Intel(R) Core(TM) i7-5557U CPU @ 3.10GHz
```

```
machdep.cpu.features: FPU VME DE PSE TSC MSR PAE
```

```
MCE CX8 APIC SEP MTRR PGE MCA CMOV PAT PSE36
```

```
CLFSH DS ACPI MMX FXSR SSE SSE2 SS HTT TM PBE
```

```
SSE3 PCLMULQDQ DTES64 MON DSCPL VMX EST TM2
```

```
SSSE3 FMA CX16 TPR PDCM SSE4.1 SSE4.2 x2APIC
```

```
MOVBE POPCNT AES PCID XSAVE OSXSAVE SEGLIM64
```


```
TSCTMR AVX1.0 RDRAND F16C
```

```
machdep.cpu.leaf7_features: SMEP ERMS RDWRFSGS
```


```
TSC_THREAD_OFFSET BMI1 AVX2 BMI2 INVPCID SMAP
```

```
RDSEED ADX IPT FPU_CSDS
```

SIMD Registers




SIMD Data Types




SIMD Vector Mode

SIMD Mode



Scalar Mode



Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- **SIMD matrix multiplication**
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

Problem

- Today's compilers (largely) do not generate SIMD code
- Back to assembly ...
- x86
 - Over 1000 instructions to learn ...
 - Green Book
- Can we use the compiler to generate all non-SIMD instructions?

x86 Intrinsics AVX Data Types

Intrinsics: Direct access to registers & assembly from C

Register

| Type | Meaning |
|----------------------|---|
| <code>__m256</code> | 256-bit as eight single-precision floating-point values, representing a YMM register or memory location |
| <code>__m256d</code> | 256-bit as four double-precision floating-point values, representing a YMM register or memory location |
| <code>__m256i</code> | 256-bit as integers, (bytes, words, etc.) |
| <code>__m128</code> | 128-bit single precision floating-point (32 bits each) |
| <code>__m128d</code> | 128-bit double precision floating-point (64 bits each) |

Intrinsics AVX Code Nomenclature

| Marking | Meaning |
|------------|---|
| [s/d] | Single- or double-precision floating point |
| [i/u]nnn | Signed or unsigned integer of bit size <i>nnn</i> , where <i>nnn</i> is 128, 64, 32, 16, or 8 |
| [ps/pd/sd] | Packed single, packed double, or scalar double |
| epi32 | Extended packed 32-bit signed integer |
| si256 | Scalar 256-bit integer |

x86 SIMD “Intrinsics”



mul_pd

Technologies

☐ MMX

☐ SSE

☐ SSE2

☐ SSE3

☐ SSSE3

☐ SSE4.1

☐ SSE4.2

☒ AVX

☐ AVX2

☐ FMA

☐ AVX-512

☐ KNC

☐ SVM

☐ Other

Categories

☐ Application-Targeted

☐ Arithmetic

☐ Bit Manipulation

☐ Cast

☐ Compare

`__m256d _mm256_mul_pd (__m256d a, __m256d b)`

Synopsis

`__m256d _mm256_mul_pd (__m256d a, __m256d b)`

`#include "immintrin.h"`

Instruction: `vmulpd ymm, ymm, ymm`

CPUID Flags: AVX

assembly instruction

Description

Multiply packed double-precision (64-bit) floating-point elements in `a` and `b`, and store the results in `dst`.

Operation

```
FOR j := 0 to 3
  i := j*64
  dst[i+63:i] := a[i+63:i] * b[i+63:i]
ENDFOR
dst[MAX:256] := 0
```

4 parallel multiplies

Performance


| Architecture | Latency | Throughput |
|--------------|---------|------------|
| Haswell | 5 | 0.5 |
| Ivy Bridge | 5 | 1 |
| Sandy Bridge | 5 | 1 |

2 instructions per clock cycle (CPI = 0.5)

<https://software.intel.com/sites/landingpage/IntrinsicsGuide/>

Raw Double Precision Throughput (Bernhard's Powerbook Pro)

| Characteristic | Value |
|-------------------------------------|--------------------|
| CPU | i7-5557U |
| Clock rate (sustained) | 3.1 GHz |
| Instructions per clock (mul_pd) | 2 |
| Parallel multiplies per instruction | 4 |
| Peak double flops | 24.8 Gflops |



<https://www.karlsruher.net/2013/06/cpu-gpu-and-mic-hardware-characteristics-over-time/>




Actual performance is lower because of overhead

Vectorized Matrix Multiplication


```
for i ...; i+=4  
  for j ...
```

Inner Loop:

```
__m256d c0 = {0,0,0,0};  
for (int k=0; k<N; k++) {  
  c0 = _mm256_fmadd_pd(  
    _mm256_load_pd(a+i+k*N),  
    _mm256_broadcast_sd(b+k+j*N),  
    c0);  
}  
_mm256_store_pd(c+i+j*N, c0);
```



“Vectorized” dgemm

```
// AVX intrinsics; P&H p. 227
void dgemm_avx(int N, double *a, double *b, double *c) {
    // avx operates on 4 doubles in parallel
    for (int i=0; i<N; i+=4) { 
        for (int j=0; j<N; j++) {
            // c0 = c[i][j]
            __m256d c0 = {0,0,0,0};
            for (int k=0; k<N; k++) {
                c0 = _mm256_add_pd(
                    c0, // c0 += a[i][k] * b[k][j]
                    _mm256_mul_pd(
                        _mm256_load_pd(a+i+k*N),
                        _mm256_broadcast_sd(b+k+j*N)));
            }
            _mm256_store_pd(c+i+j*N, c0); // c[i,j] = c0
        }
    }
}
```

Performance

| N | Gflops | |
|-----|--------|------|
| | scalar | avx |
| 32 | 1.30 | 4.56 |
| 160 | 1.30 | 5.47 |
| 480 | 1.32 | 5.27 |
| 960 | 0.91 | 3.64 |

- 4x faster
- But still << theoretical 25 Gflops!

We are flying ...

- Survey:



- But ... there is so much material to cover!
 - Solution: targeted reading
 - Weekly homework with integrated reading & lecture review

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- **Amdahl's law**
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

A trip to LA

Commercial airline:

| Get to SFO & check-in | SFO → LAX | Get to destination |
|-----------------------|-----------|--------------------|
| 3 hours | 1 hour | 3 hours |

Total time: 7 hours

Supersonic aircraft:

| Get to SFO & check-in | SFO → LAX | Get to destination |
|-----------------------|-----------|--------------------|
| 3 hours | 6 min | 3 hours |

Total time: 6.1 hours

Speedup:

Flying time

$$S_{flight} = 60 / 6 = 10x$$

Trip time

$$S_{trip} = 7 / 6.1 = 1.15x$$

Amdahl's Law

- Get enhancement E for your new PC
 - E.g. floating point rocket booster
- E
 - Speeds up some task (e.g. arithmetic) by factor S_E
 - F is fraction of program that uses this "task"

Execution Time:

| | no speedup | speedup section |
|-------------------|------------|-----------------|
| T_0 (no E) | $1-F$ | F |
| T_E (with E) | $1-F$ | F / S_E |

Speedup:

$$S = \frac{T_0}{T_E} = \frac{1}{(1 - F) + \frac{F}{S_E}}$$

Big Idea: Amdahl's Law

$$S = \frac{T_0}{T_E} = \frac{1}{(1 - F) + \frac{F}{S_E}}$$

Part not sped up

Part sped up

Example: The execution time of **half** of a program can be accelerated by a factor of **2**.
What is the program speed-up overall?

$$S = \frac{T_0}{T_E} = \frac{1}{(1 - 0.5) + \frac{0.5}{2}} = \underline{1.33} \ll 2$$

Maximum “Achievable” Speed-Up

$$S_{max} = \frac{1}{(1 - F) + \frac{F}{S_E}} \bigg|_{S_E \Rightarrow \infty} = \frac{1}{1 - F}$$

Question: What is a reasonable # of parallel processors to speed up an algorithm with $F = 95\%$? (i.e. $19/20^{\text{th}}$ can be sped up)

a) Maximum speedup:

$$F = 95\% \Rightarrow S_{max} = 20 \quad \text{but } S_E \rightarrow \infty \text{ !?}$$


b) Reasonable “engineering” compromise:

Equal time in sequential and parallel code

$$(1 - F) = \frac{F}{S_E} \Rightarrow S_E = \frac{F}{1 - F} = \frac{0.95}{0.05} = 19$$

$$\text{Then } S = \frac{S_{max}}{2} = 10$$

Amdahl's Law



Strong and Weak Scaling

- To get good speedup on a parallel processor while keeping the problem size fixed is harder than getting good speedup by increasing the size of the problem.
 - *Strong scaling*: when speedup can be achieved on a parallel processor without increasing the size of the problem
 - *Weak scaling*: when speedup is achieved on a parallel processor by increasing the size of the problem proportionally to the increase in the number of processors
- *Load balancing* is another important factor: every processor doing same amount of work
 - Just one unit with twice the load of others cuts speedup almost in half

Clickers/Peer Instruction

Suppose a program spends 80% of its time in a square root routine. How much must you speedup square root to make the program run 5 times faster?

$$S = \frac{T_0}{T_E} = \frac{1}{(1 - F) + \frac{F}{S_E}}$$

| Answer | S_E |
|--------|-------------------|
| A | 5 |
| B | 16 |
| C | 20 |
| D | 100 |
| E | None of the above |

Clickers/Peer Instruction

Suppose a program spends 80% of its time in a square root routine. How much must you speedup square root to make the program run 5 times faster?

$$S = \frac{T_0}{T_E} = \frac{1}{(1 - F) + \frac{F}{S_E}}$$

| Answer | S_E |
|--------|-------------------|
| A | 5 |
| B | 16 |
| C | 20 |
| D | 100 |
| E | None of the above |

Administrivia

- MT2 is
 - Tuesday, November 1,
 - 3:30-5pm
 - see web for room assignments
- TA Review Session:
 - Sunday 10/30, 3:30 – 5 PM in 10 Evans
 - See Piazza

MT 2 Topics

- Covers lecture material up to 10/20
 - Caches
 - not floating point
- Combinatorial logic including synthesis and truth tables
- FSMs
- Timing and timing diagrams
- Pipelining
- Datapath, hazards, stalls
- Performance (e.g. CPI, instructions per second, latency)
- Caches
- All topics covered in MT 1
 - Focus is new material, but do not be surprised by e.g. MIPS assembly

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- **Loop unrolling**
- Memory access strategy - blocking
- And in Conclusion, ...

Amdahl's Law applied to **dgemm**

- Measured **dgemm** performance
 - Peak 5.5 Gflops
 - Large matrices 3.6 Gflops
 - Processor 24.8 Gflops
- Why are we not getting (close to) 25 Gflops?
 - Something else (not floating point ALU) is limiting performance!
 - But what? Possible culprits:
 - Cache
 - Hazards
 - Let's look at both!

Pipeline Hazards – **dgemm**



Technologies

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☒ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX-512
- ☐ KNC
- ☐ SVML
- ☐ Other

Categories

- ☐ Application-Targeted
- ☐ Arithmetic
- ☐ Bit Manipulation
- ☐ Cast
- ☐ Compare

mul_pd

```
__m256d _mm256_mul_pd (__m256d a, __m256d b)
```

Synopsis

```
__m256d _mm256_mul_pd (__m256d a, __m256d b)
#include "immintrin.h"
Instruction: vmulpd ymm, ymm, ymm
CPUID Flags: AVX
```

Description

Multiply packed double-precision (64-bit) floating-point elements in **a** and **b**, and store the results in **dst**.

Operation

```
FOR j := 0 to 3
  i := j*64
  dst[i+63:i] := a[i+63:i] * b[i+63:i]
ENDFOR
dst[MAX:256] := 0
```

Performance

| Architecture | Latency | Throughput |
|--------------|---------|------------|
| Haswell | 5 | 0.5 |
| Ivy Bridge | 5 | 1 |
| Sandy Bridge | 5 | 1 |

Loop Unrolling

```
// Loop unrolling; P&H p. 352
const int UNROLL = 4;
```

```
void dgemm_unroll(int n, double *A, double *B, double *C) {
    for (int i=0; i<n; i+= UNROLL*4) {
        for (int j=0; j<n; j++) {
            __m256d c[4];
            for (int x=0; x<UNROLL; x++)
                c[x] = _mm256_load_pd(C+i+x*4+j*n);
            for (int k=0; k<n; k++) {
                __m256d b = _mm256_broadcast_sd(B+k+j*n);
                for (int x=0; x<UNROLL; x++)
                    c[x] = _mm256_add_pd(c[x],
                                         _mm256_mul_pd(_mm256_load_pd(A+n*k+x*4+i), b));
            }
            for (int x=0; x<UNROLL; x++)
                _mm256_store_pd(C+i+x*4+j*n, c[x]);
        }
    }
}
```

← 4 registers

← Compiler does the unrolling

How do you verify that the generated code is actually unrolled?

Performance

| N | Gflops | | |
|-----|--------|------|--------|
| | scalar | avx | unroll |
| 32 | 1.30 | 4.56 | 12.95 |
| 160 | 1.30 | 5.47 | 19.70 |
| 480 | 1.32 | 5.27 | 14.50 |
| 960 | 0.91 | 3.64 | 6.91 |

WOW!

?

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- **Memory access strategy - blocking**
- And in Conclusion, ...

FPU versus Memory Access

- How many floating point operations does matrix multiply take?
 - $F = 2 \times N^3$ (N^3 multiplies, N^3 adds)
- How many memory load/stores?
 - $M = 3 \times N^2$ (for A, B, C)
- Many more floating point operations than memory accesses
 - $q = F/M = 2/3 * N$
 - Good, since arithmetic is faster than memory access
 - Let's check the code ...

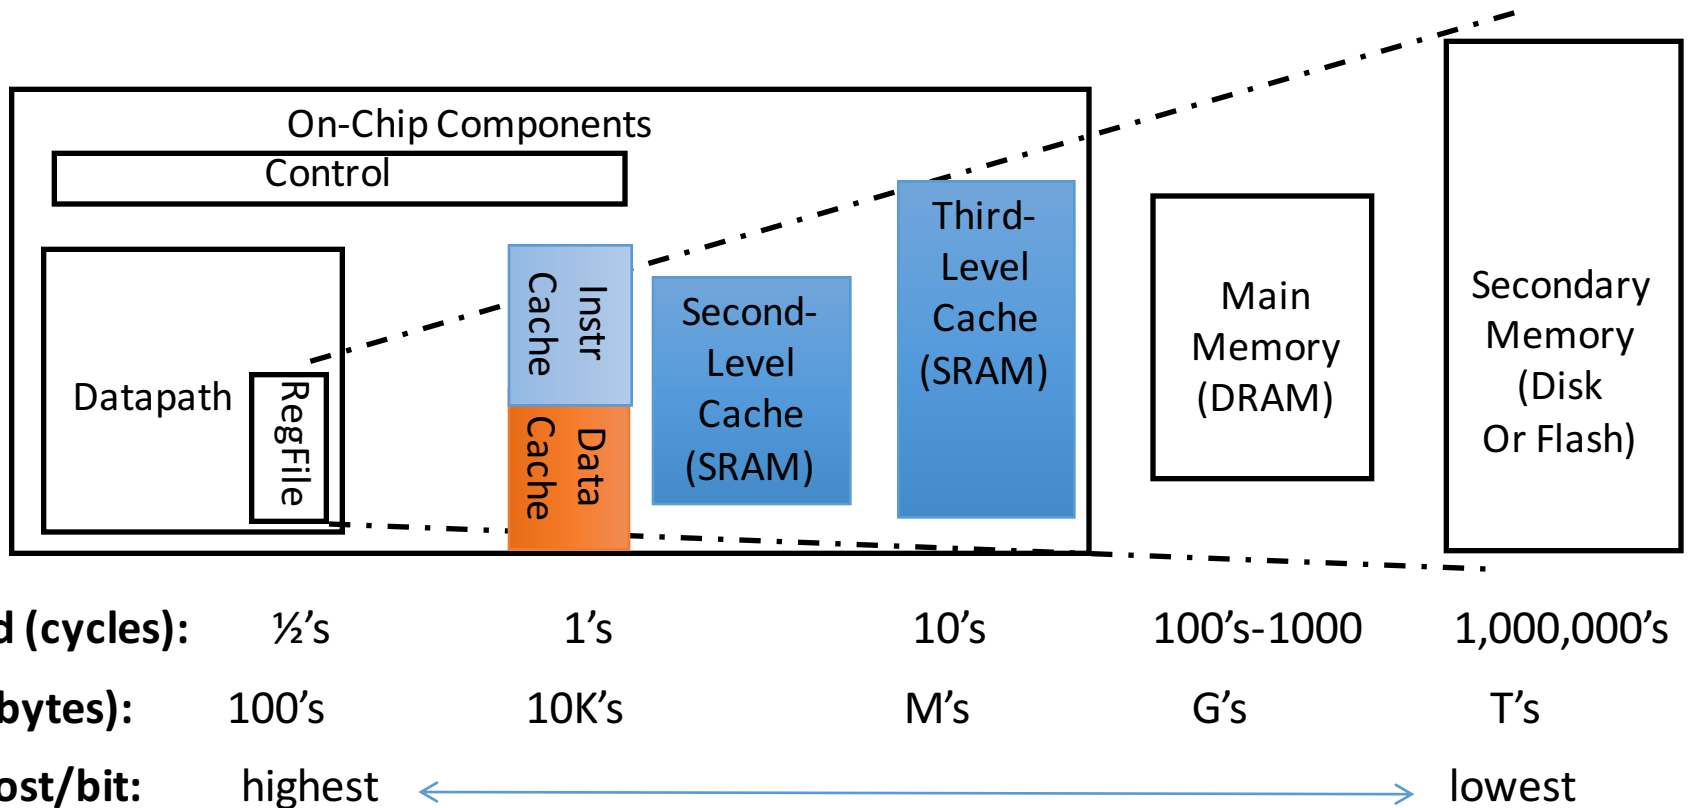
But memory is accessed repeatedly

Inner loop:

```
for (int k=0; k<N; k++) {  
    c0 = _mm256_add_pd(  
        c0, // c0 += a[i][k] * b[k][j]  
        _mm256_mul_pd(  
            _mm256_load_pd(a+i+k*N),  
            _mm256_broadcast_sd(b+k+j*N)));  
}
```




- $q = F/M = 1!$ (2 loads and 2 floating point operations)

Typical Memory Hierarchy



- Where are the operands (A, B, C) stored?
- What happens as N increases?
- **Idea:** arrange that most accesses are to fast cache!

Sub-Matrix Multiplication *or: Beating Amdahl's Law*



Blocking

- Idea:

- Rearrange code to use values loaded in cache many times
- Only “few” accesses to slow main memory (DRAM) per floating point operation
- → throughput limited by FP hardware and cache, not slow DRAM
- P&H p. 556

Memory Access Blocking

```
// Cache blocking; P&H p. 556
const int BLOCKSIZE = 32;
```

```
void do_block(int n, int si, int sj, int sk, double *A, double *B, double *C) {
    for (int i=si; i<si+BLOCKSIZE; i+=UNROLL*4)
        for (int j=sj; j<sj+BLOCKSIZE; j++) {
            __m256d c[4];
            for (int x=0; x<UNROLL; x++)
                c[x] = _mm256_load_pd(C+i+x*4+j*n);
            for (int k=sk; k<sk+BLOCKSIZE; k++) {
                __m256d b = _mm256_broadcast_sd(B+k+j*n);
                for (int x=0; x<UNROLL; x++)
                    c[x] = _mm256_add_pd(c[x],
                                         _mm256_mul_pd(_mm256_load_pd(A+n*k+x*4+i), b));
            }
            for (int x=0; x<UNROLL; x++)
                _mm256_store_pd(C+i+x*4+j*n, c[x]);
        }
}
```

```
void dgemm_block(int n, double* A, double* B, double* C) {
    for(int sj=0; sj<n; sj+=BLOCKSIZE)
        for(int si=0; si<n; si+=BLOCKSIZE)
            for (int sk=0; sk<n; sk += BLOCKSIZE)
                do_block(n, si, sj, sk, A, B, C);
}
```

Performance

| N | Gflops | | | |
|-----|--------|------|--------|----------|
| | scalar | avx | unroll | blocking |
| 32 | 1.30 | 4.56 | 12.95 | 13.80 |
| 160 | 1.30 | 5.47 | 19.70 | 21.79 |
| 480 | 1.32 | 5.27 | 14.50 | 20.17 |
| 960 | 0.91 | 3.64 | 6.91 | 15.82 |

Agenda

- 61C – the big picture
- Parallel processing
- Single instruction, multiple data
- SIMD matrix multiplication
- Amdahl's law
- Loop unrolling
- Memory access strategy - blocking
- And in Conclusion, ...

And in Conclusion, ...

- Approaches to Parallelism
 - SISD, SIMD, MIMD (next lecture)
- SIMD
 - One instruction operates on multiple operands simultaneously
- Example: matrix multiplication
 - Floating point heavy → exploit Moore's law to make fast
- Amdahl's Law:
 - Serial sections limit speedup
 - Cache
 - Blocking
 - Hazards
 - Loop unrolling