

CS 61C: Great Ideas in Computer Architecture

Lecture 19: *Thread-Level Parallel Processing*

Bernhard Boser & Randy Katz

<http://inst.eecs.berkeley.edu/~cs61c>

Agenda

- **MIMD - multiple programs simultaneously**
- Threads
- Parallel programming: OpenMP
- Synchronization primitives
- Synchronization in OpenMP
- And, in Conclusion ...

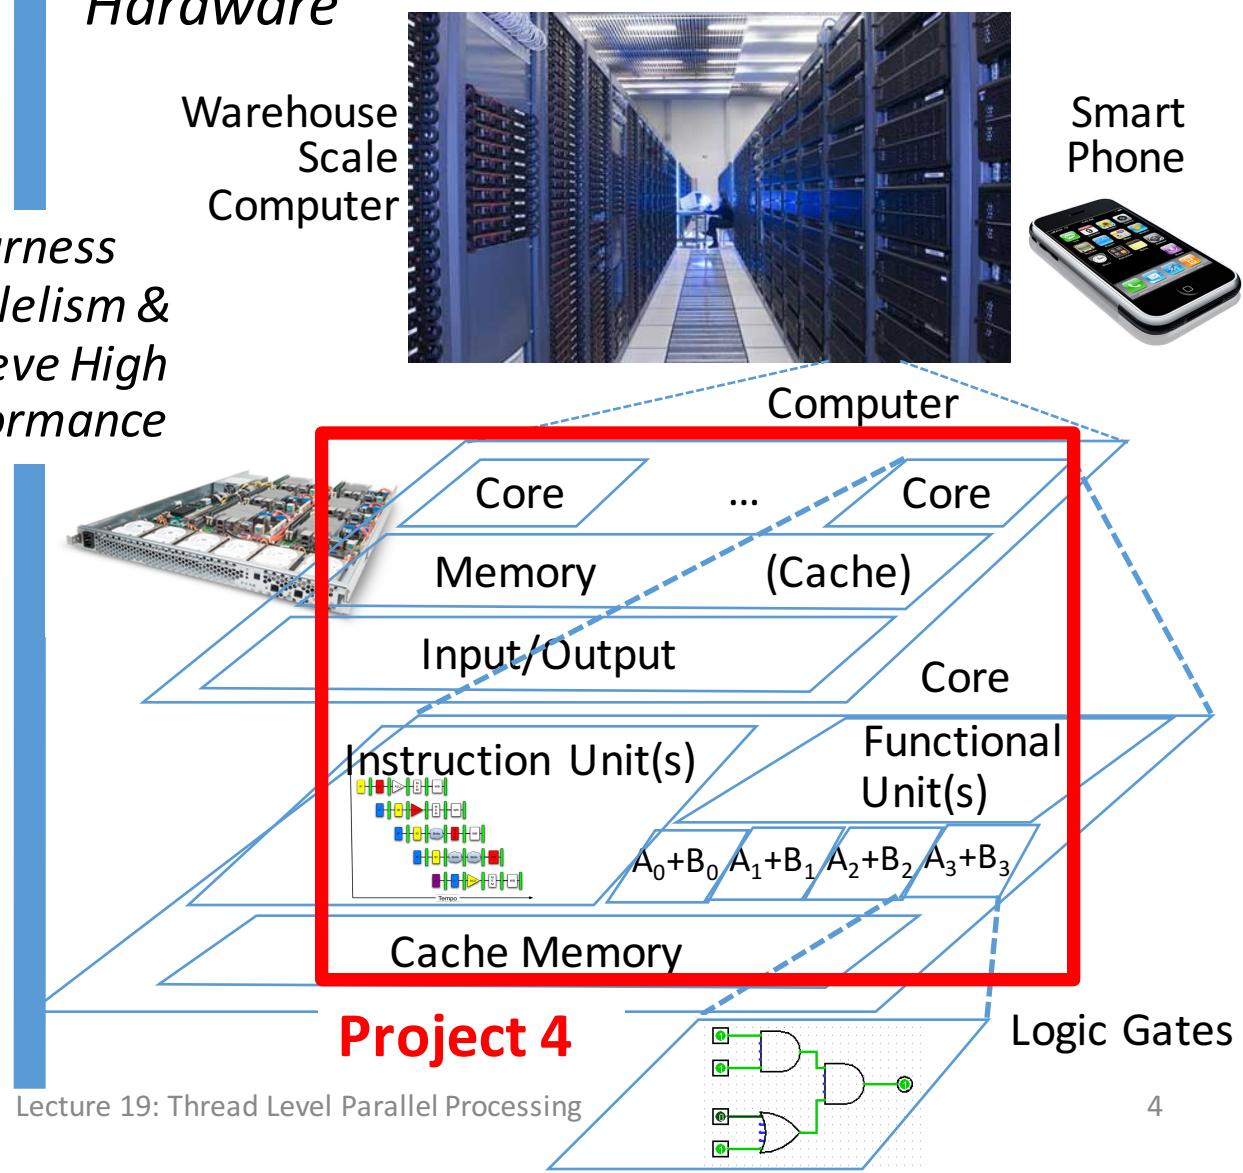
Improving Performance

1. Increase clock rate f_s
 - Reached practical maximum for today's technology
 - < 5GHz for general purpose computers
2. Lower CPI (cycles per instruction)
 - SIMD, "instruction level parallelism"
3. Perform multiple tasks simultaneously
 - Multiple CPUs, each executing different program
 - Tasks may be related
 - E.g. each CPU performs part of a big matrix multiplication
 - or unrelated
 - E.g. distribute different web http requests over different computers
 - E.g. run ppt (view lecture slides) and browser (youtube) simultaneously
4. Do all of the above:
 - High f_s , SIMD, multiple parallel tasks

Today's
lecture

New-School Machine Structures (It's a bit more complicated!)

- | Software | Hardware |
|--|--|
| • Parallel Requests
Assigned to computer
e.g., Search "Katz" | Warehouse Scale Computer |
| • Parallel Threads
Assigned to core
e.g., Lookup, Ads | Harness Parallelism & Achieve High Performance |
| • Parallel Instructions
>1 instruction @ one time
e.g., 5 pipelined instructions | Computer |
| • Parallel Data
>1 data item @ one time
e.g., Add of 4 pairs of words | Core
Memory
(Cache)
Input/Output
Instruction Unit(s)
Functional Unit(s) |
| • Hardware descriptions
All gates @ one time | Cache Memory |
| • Programming Languages | Logic Gates |



Parallel Computer Architectures



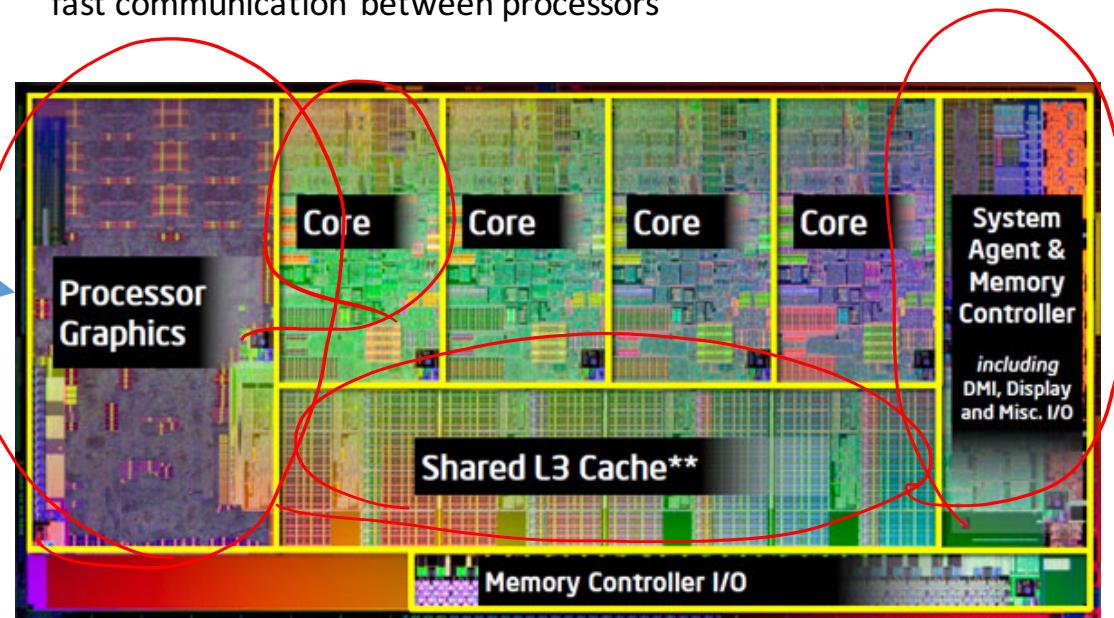
Several separate computers,
some means for communication
(e.g. Ethernet)

GPU “graphics processing unit”

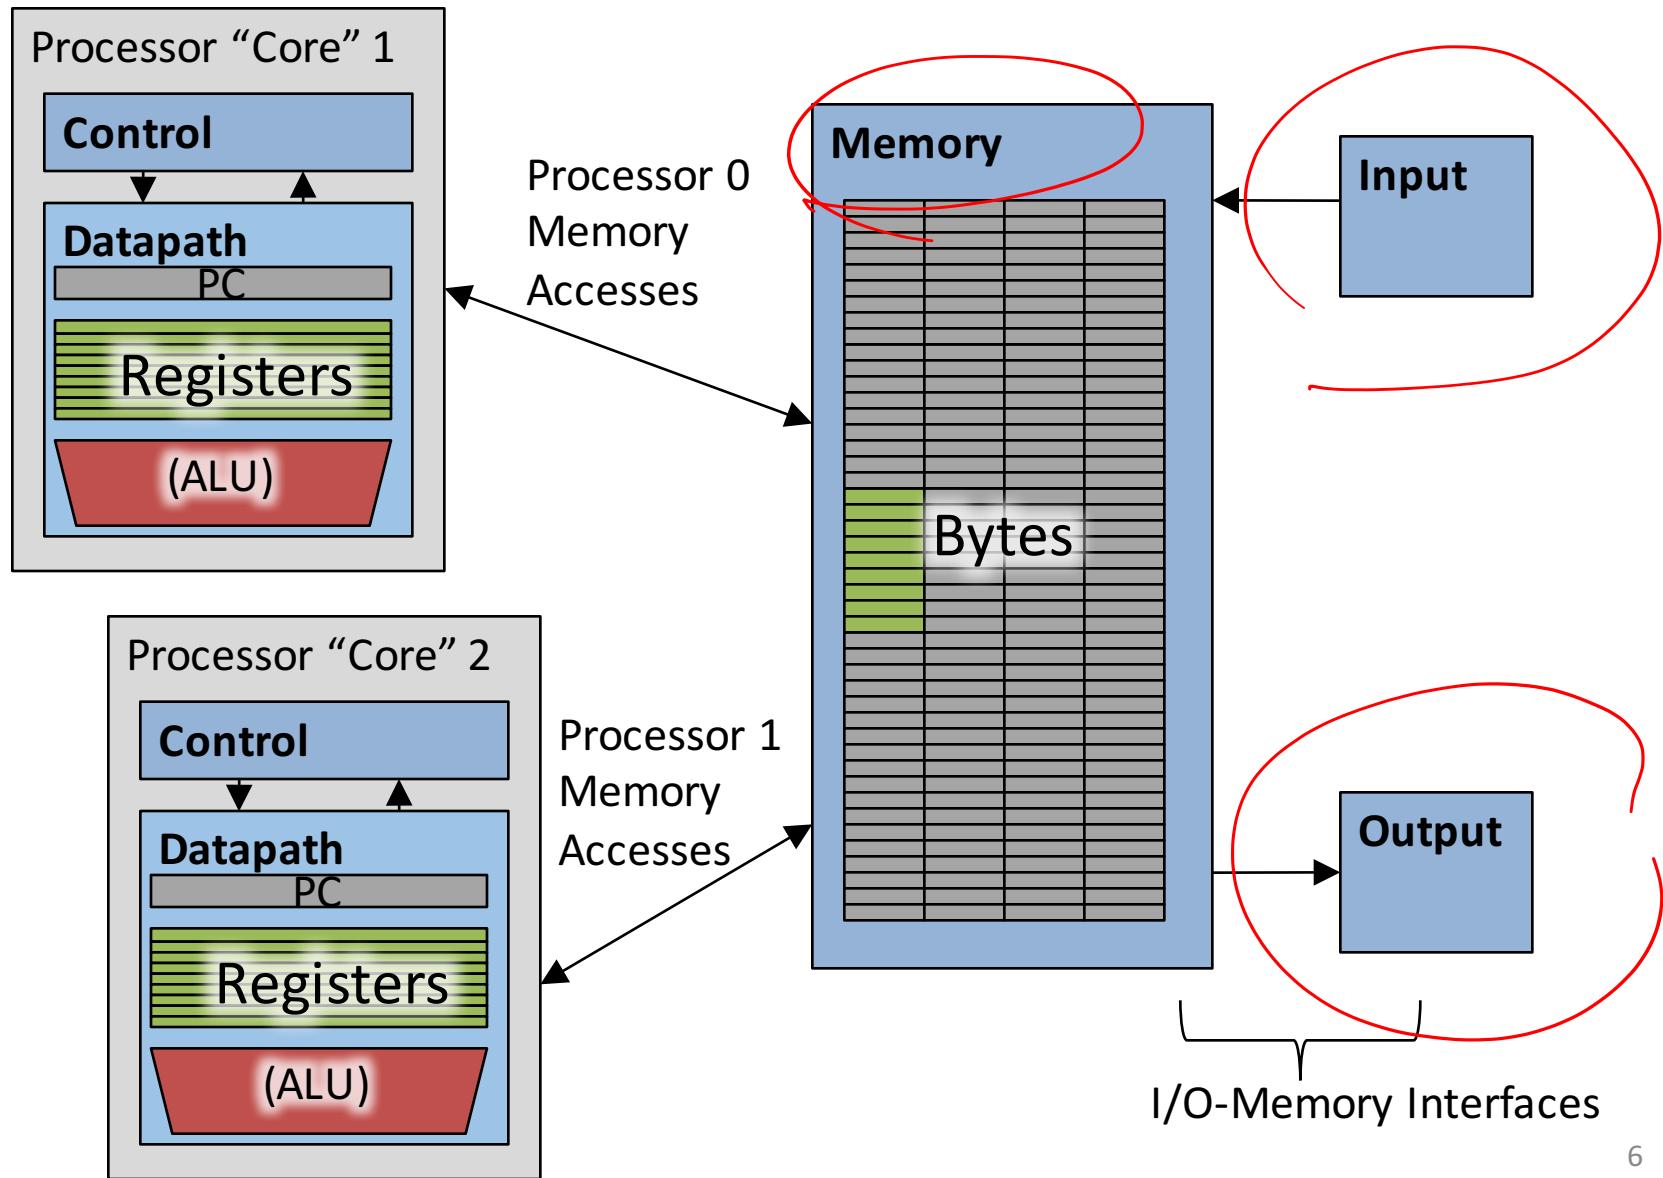
Multi-core CPU:
1 datapath in single chip
share L3 cache, memory, peripherals
Example: Hive machines



Massive array of computers,
fast communication between processors



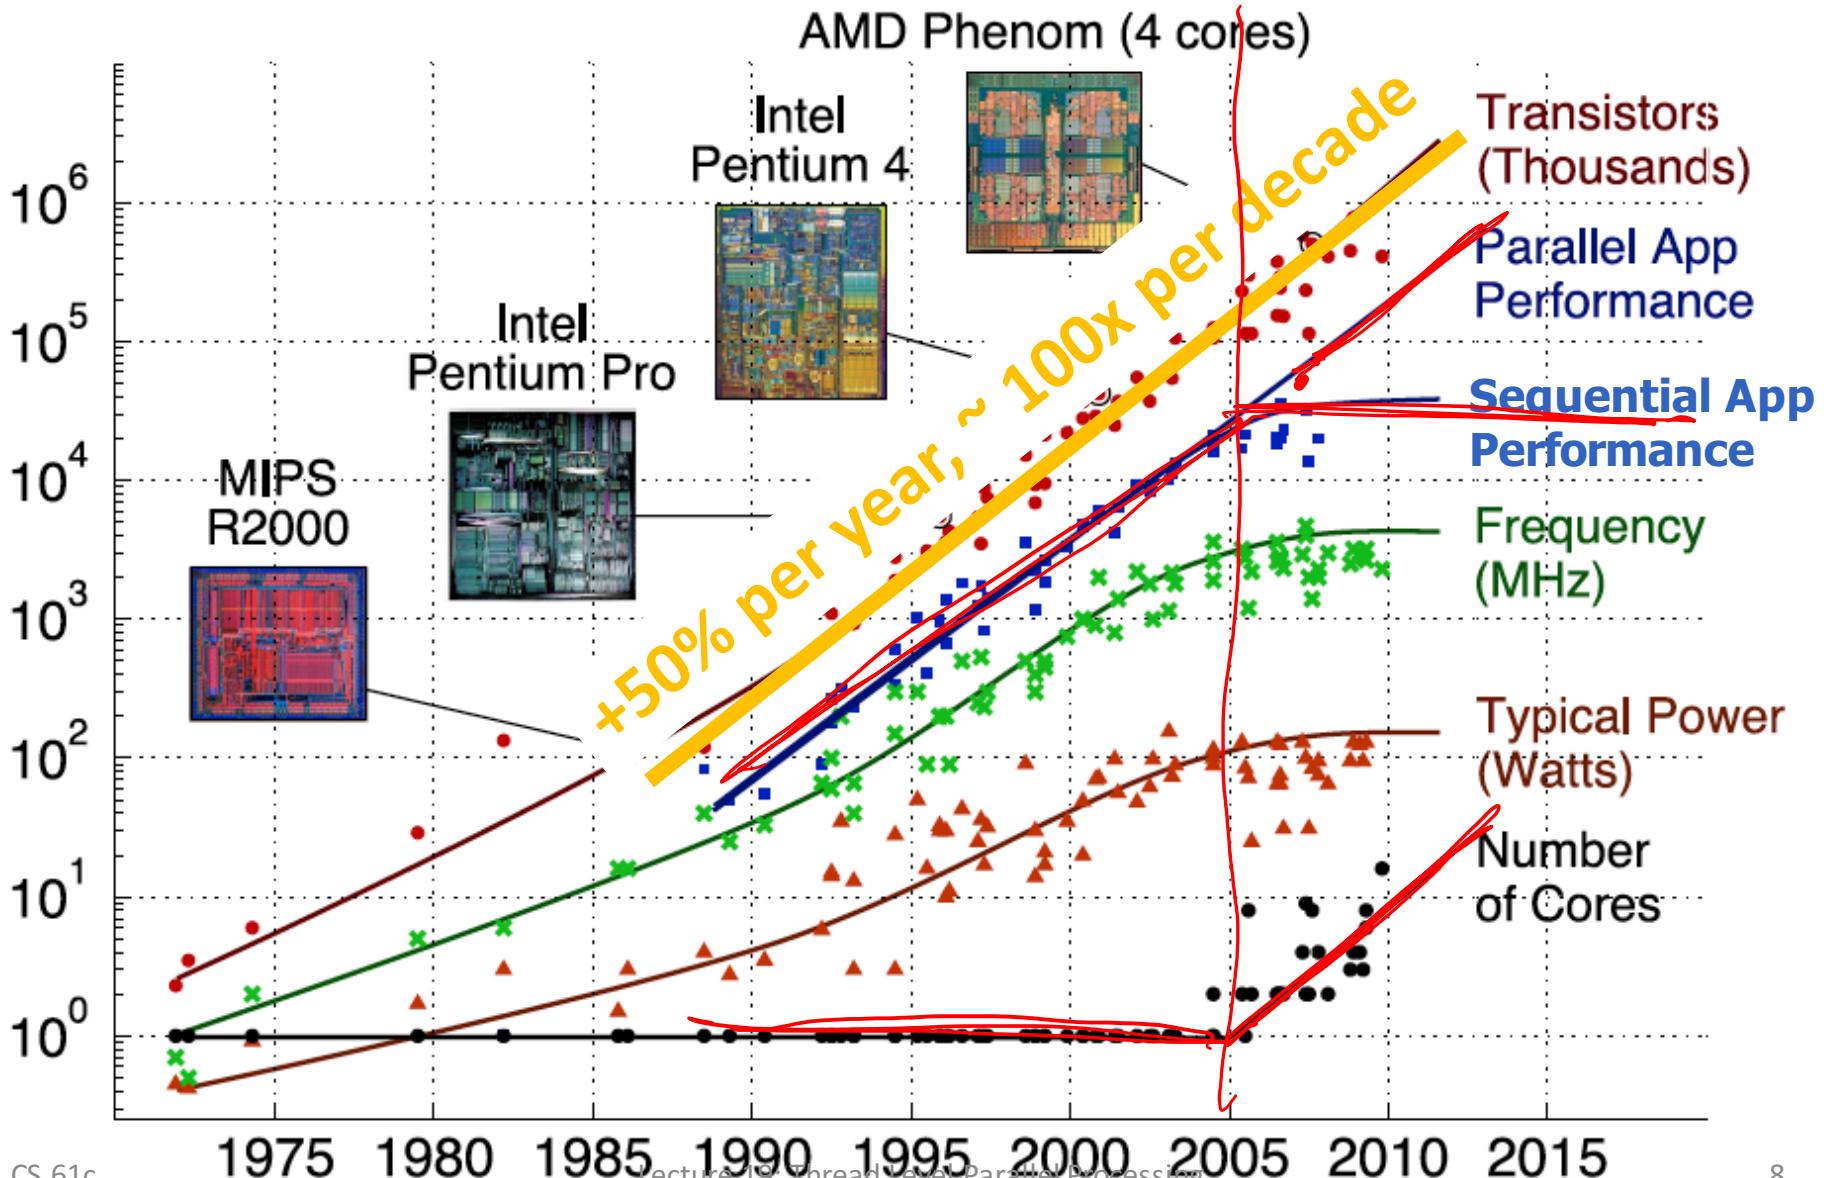
Example: CPU with 2 Cores



Multiprocessor Execution Model

- Each processor (core) executes its own instructions
- Separate resources (not shared)
 - Datapath (PC, registers, ALU)
 - Highest level caches (e.g. 1st and 2nd)
- Shared resources
 - Memory (DRAM)
 - Often 3rd level cache
 - Often on same silicon chip
 - But not a requirement
- Nomenclature
 - “Multiprocessor Microprocessor”
 - Multicore processor
 - E.g. 4 core CPU (central processing unit)
 - Executes 4 different instruction streams simultaneously

Transition to Multicore



Multiprocessor Execution Model

- Shared memory
 - Each “core” has access to the entire memory in the processor
 - Special hardware keeps caches consistent
 - Advantages:
 - Simplifies communication in program via shared variables
 - Drawbacks:
 - Does not scale well:
 - “Slow” memory shared by many “customers” (cores)
 - May become bottleneck (Amdahl’s Law)
- Two ways to use a multiprocessor:
 - Job-level parallelism
 - Processors work on unrelated problems
 - No communication between programs
 - Partition work of single task between several cores
 - E.g. each performs part of large matrix multiplication

Parallel Processing

- It's difficult!
- It's inevitable
 - Only path to increase performance
 - Only path to lower energy consumption (improve battery life)
- In mobile systems (e.g. smart phones, tablets)
 - Multiple cores
 - Dedicated processors, e.g.
 - motion processor in iPhone
 - GPU (graphics processing unit)
- Warehouse-scale computers
 - multiple “nodes”
 - “boxes” with several CPUs, disks per box
 - MIMD (multi-core) and SIMD (e.g. AVX) in each node

Potential Parallel Performance (assuming software can use it)

Year	Cores	SIMD bits /Core	Core * SIMD bits	Total, e.g. FLOPs/Cycle
2003	MIMD 2	SIMD 128	256	MIMD 4
2005	+2/ 4	2X/ 128	512	& SIMD 8
2007	2yrs 6	4yrs 128	768	12
2009	8	128	1024	16
2011	10	256	2560	40
2013	12	256	3072	48
2015	2.5X 14	8X 512	7168	112
2017	16	512	8192	128
2019	18	1024	18432	288
2021	20	1024	20480	320

12 years

20 x in 12 years
 $20^{1/12} = 1.28 \times \rightarrow 28\% \text{ per year or } 2x \text{ every 3 years!}$
 IF (!) we can use it

Agenda

- MIMD - multiple programs simultaneously
- **Threads**
- Parallel programming: OpenMP
- Synchronization primitives
- Synchronization in OpenMP
- And, in Conclusion ...

Programs Running on my Computer

```
PID TTY      TIME CMD
 220 ??      0:04.34 /usr/libexec/UserEventAgent (Aqua)
 222 ??      0:10.60 /usr/sbin/distnoted agent
 224 ??      0:09.11 /usr/sbin/cfprefsd agent
 229 ??      0:04.71 /usr/sbin/usernoted
 230 ??      0:02.35 /usr/libexec/nsurlsessiond
 232 ??      0:28.68 /System/Library/PrivateFrameworks/CalendarAgent.framework/Executables/CalendarAgent
 234 ??      0:04.36 /System/Library/PrivateFrameworks/GameCenterFoundation.framework/Versions/A/gamed
 235 ??      0:01.90 /System/Library/CoreServices/cloudphotosd.app/Contents/MacOS/cloudphotosd
 236 ??      0:49.72 /usr/libexec/secinitd
 239 ??      0:01.66 /System/Library/PrivateFrameworks/TCC.framework/Resources/tccd
 240 ??      0:12.68 /System/Library/Frameworks/Accounts.framework/Versions/A/Support/accountsd
 241 ??      0:09.56 /usr/libexec/SafariCloudHistoryPushAgent
 242 ??      0:00.27 /System/Library/PrivateFrameworks/CallHistory.framework/Support/CallHistorySyncHelper
 243 ??      0:00.74 /System/Library/CoreServices/mapspushd
 244 ??      0:00.79 /usr/libexec/fmfd
 246 ??      0:00.09 /System/Library/PrivateFrameworks/AskPermission.framework/Versions/A/Resources/askpermissiond
 248 ??      0:01.03 /System/Library/PrivateFrameworks/CloudDocsDaemon.framework/Versions/A/Support/bird
 249 ??      0:02.50 /System/Library/PrivateFrameworks/IDS.framework/identityservicesd.app/Contents/MacOS/identityservicesd
 250 ??      0:04.81 /usr/libexec/secd
 254 ??      0:24.01 /System/Library/PrivateFrameworks/CloudKitDaemon.framework/Support/cloudd
 258 ??      0:04.73 /System/Library/PrivateFrameworks/TelephonyUtilities.framework/callservicesd
 267 ??      0:02.15 /System/Library/CoreServices/AirPlayUIAgent.app/Contents/MacOS/AirPlayUIAgent --launchd
 271 ??      0:03.91 /usr/libexec/nsurlstoraged
 274 ??      0:00.90 /System/Library/PrivateFrameworks/CommerceKit.framework/Versions/A/Resources/storeaccountd
 282 ??      0:00.09 /usr/sbin/pboard
 283 ??      0:00.90
/System/Library/PrivateFrameworks/InternetAccounts.framework/Versions/A/XPCServices/com.apple.internetaccounts.xpc/Contents/MacOS/com.apple.internetaccounts
 285 ??      0:04.72 /System/Library/Frameworks/ApplicationServices.framework/Frameworks/ATS.framework/Support/fontd
 291 ??      0:00.25 /System/Library/Frameworks/Security.framework/Versions/A/Resources/CloudKeychainProxy.bundle/Contents/MacOS/CloudKeychainProxy
 292 ??      0:09.54 /System/Library/CoreServices/CoreServicesUIAgent.app/Contents/MacOS/CoreServicesUIAgent
 293 ??      0:00.29
/System/Library/PrivateFrameworks/CloudPhotoServices.framework/Versions/A/Frameworks/CloudPhotoServicesConfiguration.framework/Versions/A/XPCServices/com.apple.CloudPhotosConfiguration.xpc/Contents/MacOS/com.apple.CloudPhotosConfiguration
 297 ??      0:00.84 /System/Library/PrivateFrameworks/CloudServices.framework/Resources/com.apple.sbd
 302 ??      0:26.11 /System/Library/CoreServices/Dock.app/Contents/MacOS/Dock
 303 ??      0:09.55 /System/Library/CoreServices/SystemUIServer.app/Contents/MacOS/SystemUIServer
```

... 156 total at this moment
How does my laptop do this?
Imagine doing 156 assignments all at the same time!

Threads

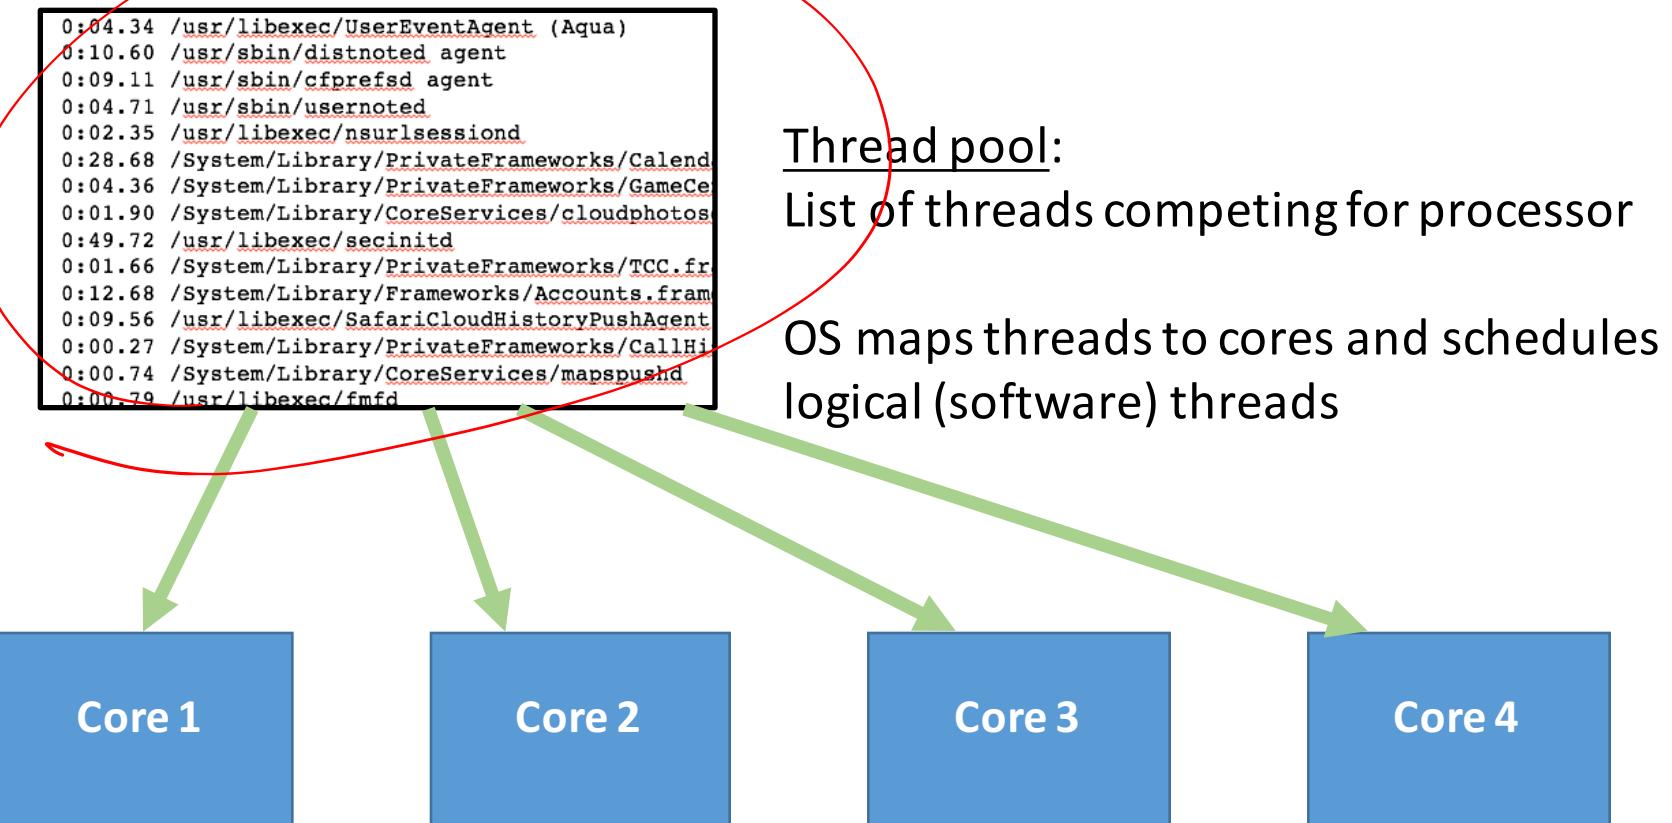
- Sequential flow of instructions that performs some task
 - Up to now we just called this a “program”
- Each thread has a
 - Dedicated PC (program counter)
 - Separate registers
 - Accesses the shared memory
- Each processor provides one (or more)
 - *hardware threads* (or *harts*) that actively execute instructions
 - Each core executes one “*hardware thread*”
- Operating system multiplexes multiple
 - *software threads* onto the available hardware threads
 - all threads except those mapped to hardware threads are waiting

Operating System Threads

Give illusion of many “simultaneously” active threads

1. Multiplex software threads onto hardware threads:
 - a) Switch out blocked threads (e.g. cache miss, user input, network access)
 - b) Timer (e.g. switch active thread every 1ms)
2. Remove a software thread from a hardware thread by
 - i. interrupting its execution
 - ii. saving its registers and PC to memory
3. Start executing a different software thread by
 - i. loading its previously saved registers into a hardware thread’s registers
 - ii. jumping to its saved PC

Example: 4 Cores

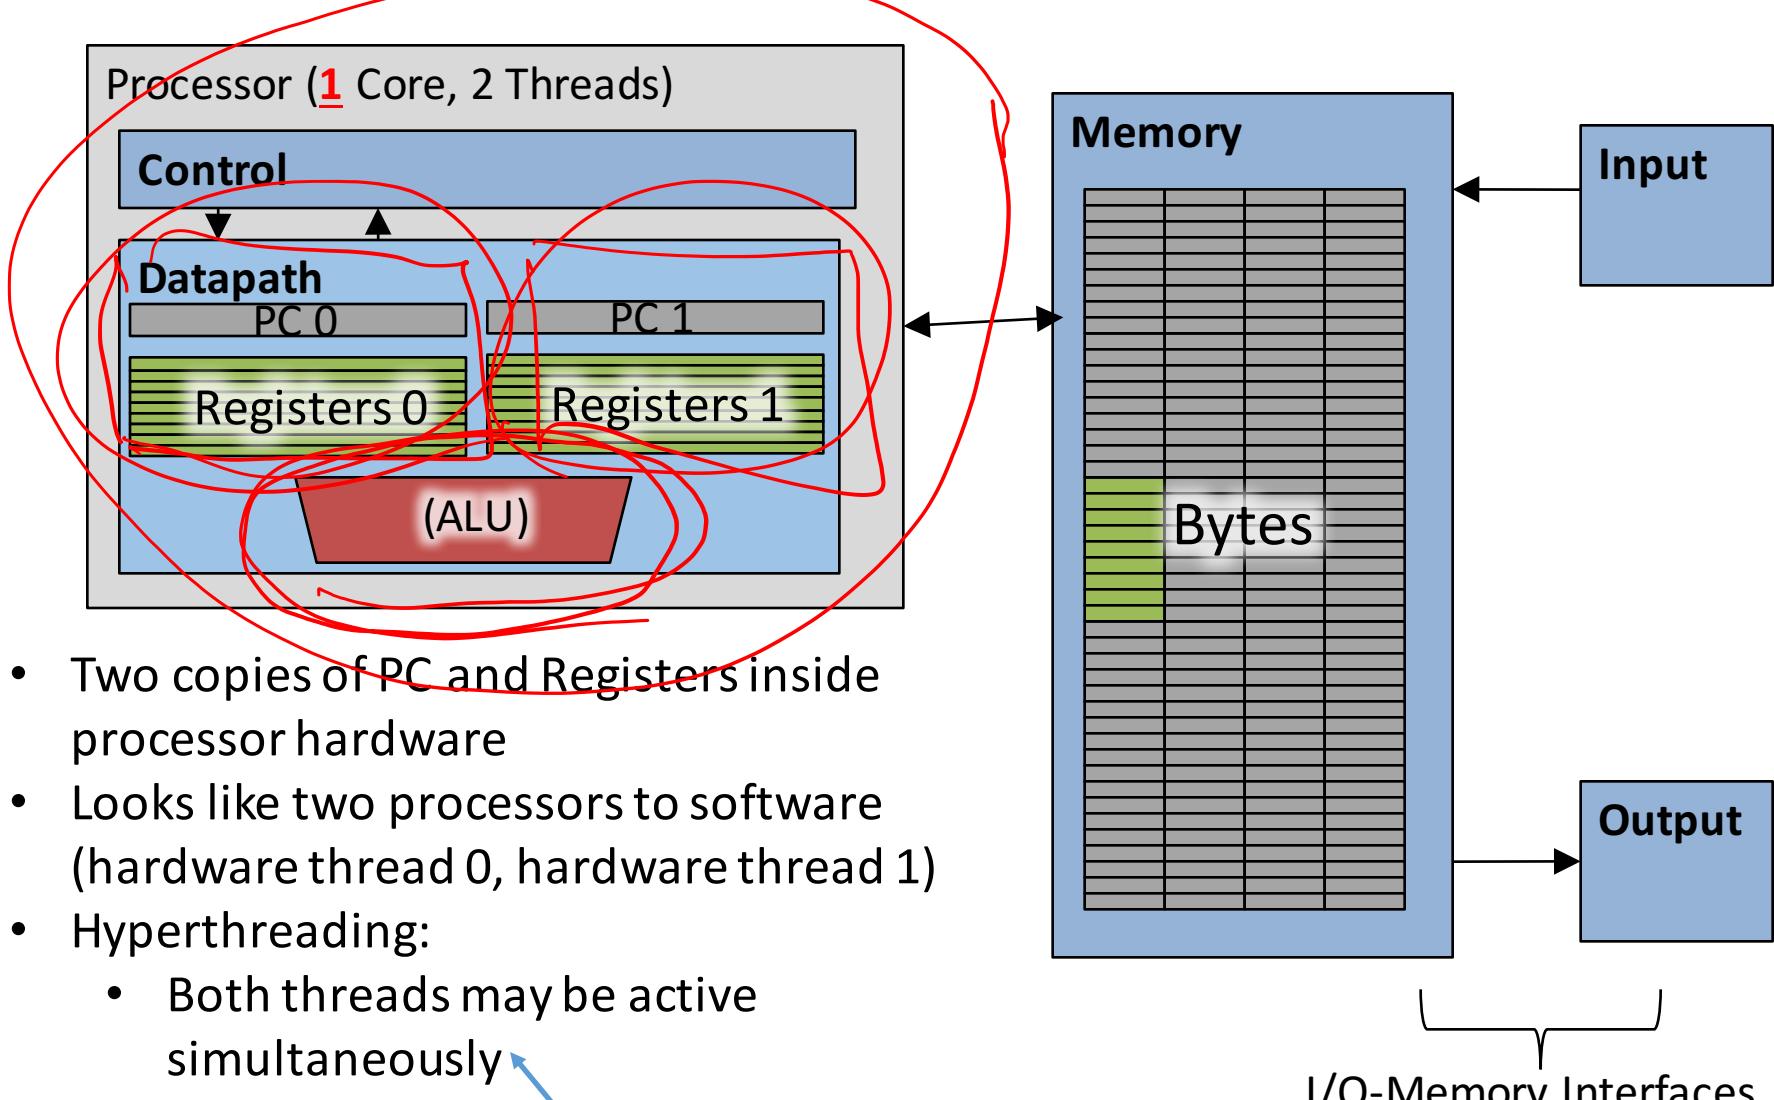


Each “Core” actively runs 1 program at a time

Multithreading

- Typical scenario:
 - Active thread encounters cache miss
 - Active thread waits ~ 1000 cycles for data from DRAM
 - \rightarrow switch out and run different thread until data available
- Problem
 - Must save current thread state and load new thread state
 - PC, all registers (could be many, e.g. AVX)
 - \rightarrow must perform switch in $\ll 1000$ cycles
- Can hardware help?
 - Moore's law: transistors are plenty

Hardware assisted Software Multithreading



Note: presented incorrectly in the lecture

Multithreading

- Logical threads
 - $\approx 1\%$ more hardware, $\approx 10\% (?)$ better performance
 - Separate registers
 - Share datapath, ALU(s), caches
- Multicore
 - \Rightarrow Duplicate Processors
 - $\approx 50\%$ more hardware, $\approx 2X$ better performance?
- Modern machines do both
 - Multiple cores with multiple threads per core

Bernhard's Laptop

```
$ sysctl -a | grep hw
```

hw. <u>physicalcpu</u> :	2
hw. <u>logicalcpu</u> :	4
hw.l1icachesize:	32,768
hw.l1dcachesize:	32,768
hw.l2cachesize:	262,144
hw.l3cachesize:	3,145,728

- 2 Cores
- 4 Threads total

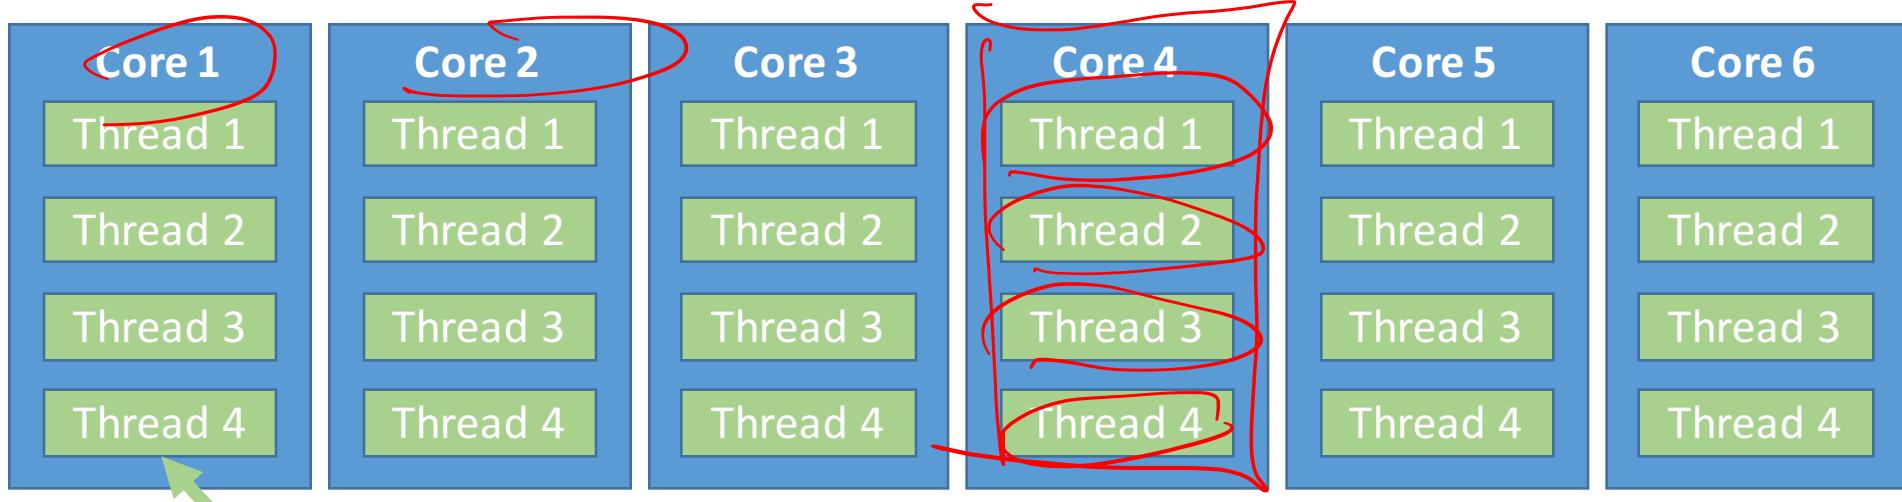
Example: 6 Cores, 24 Logical Threads

```
0:04.34 /usr/libexec/UserEventAgent (Aqua)
0:10.60 /usr/sbin/distnoded agent
0:09.11 /usr/sbin/cfprefsd agent
0:04.71 /usr/sbin/usernoted
0:02.35 /usr/libexec/nsurlsessiond
0:28.68 /System/Library/PrivateFrameworks/Calend...
0:04.36 /System/Library/PrivateFrameworks/GameCe...
0:01.90 /System/Library/CoreServices/cloudphotos...
0:49.72 /usr/libexec/secinitd
0:01.66 /System/Library/PrivateFrameworks/TCC.fr...
0:12.68 /System/Library/Frameworks/Accounts.fram...
0:09.56 /usr/libexec/SafariCloudHistoryPushAgent
0:00.27 /System/Library/PrivateFrameworks/CallHi...
0:00.74 /System/Library/CoreServices/mapspushd
0:00.79 /usr/libexec/tmfd
```

Thread pool:

List of threads competing for processor

OS maps threads to cores and schedules logical (software) threads



4 Logical threads per core (hardware) thread

Agenda

- MIMD - multiple programs simultaneously
- Threads
- **Parallel programming: OpenMP**
- Synchronization primitives
- Synchronization in OpenMP
- And, in Conclusion ...

Languages supporting Parallel Programming

ActorScript	Concurrent Pascal	JoCaml	Orc
Ada	Concurrent ML	Join	Oz
Afnix	Concurrent Haskell	Java	Pict
Alef	Curry	Joule	Reia
Alice	CUDA	Joyce	SALSA
APL	E	LabVIEW	Scala
Axum	Eiffel	Limbo	SISAL
Chapel	Erlang	Linda	SR
Cilk	Fortan 90	MultiLisp	Stackless Python
Clean	Go	Modula-3	SuperPascal
Clojure	Io	Occam	VHDL
Concurrent C	Janus	occam-π	XC

Which one to pick?

Why so many parallel programming languages?

- Piazza question:
 - Why “intrinsics”?
 - TO Intel: fix your #()&\$! Compiler!
- It's happening ... but
 - SIMD features are continually added to compilers (Intel, gcc)
 - Intense area of research
 - Research progress:
 - 20+ years to translate C into good (fast!) assembly
 - How long to translate C into good (fast!) parallel code?
 - General problem is very hard to solve
 - Present state: specialized solutions for specific cases
 - **Your opportunity to become famous!**

Parallel Programming Languages

- Number of choices is indication of
 - No universal solution
 - Needs are very problem specific
 - E.g.
 - Scientific computing (matrix multiply)
 - Webserver: handle many unrelated requests simultaneously
 - Input / output: it's all happening simultaneously!
- Specialized languages for different tasks
 - Some are easier to use (for some problems)
 - None is particularly "easy" to use
- 61C
 - Parallel language examples for high-performance computing
 - OpenMP

Parallel Loops

- Serial execution:

```
for (int i=0; i<100; i++) {  
    ...  
}
```

- Parallel Execution:

<code>for (int i=0; i<25; i++) { ... }</code>	<code>for (int i=25; i<50; i++) { ... }</code>	<code>for (int i=50; i<75; i++) { ... }</code>	<code>for (int i=75; i<100; i++) { ... }</code>
--	---	---	--

Parallel **for** in OpenMP

→ **#include <omp.h>**

→ **#pragma omp parallel for**
for (int i=0; i<100; i++) {
 ...
}

OpenMP Example

```
#include <stdio.h>
#include <omp.h>

// gcc-5 -fopenmp for.c

int main(void) {
    omp_set_num_threads(4);
    int a[] = { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10 };
    int N = sizeof(a)/sizeof(int);

    #pragma omp parallel for
    for (int i=0; i<N; i++) {
        printf("thread %d, i = %2d\n",
               omp_get_thread_num(), i);
        a[i] = a[i] + 100*omp_get_thread_num();
    }

    for (int i=0; i<N; i++) printf("%d ", a[i]);
    printf("\n");
}
```

\$ gcc-5 -fopenmp for.c; ./a.out

thread 0, i = 0 —

thread 1, i = 3

thread 2, i = 6

thread 3, i = 8 —

thread 0, i = 1 —

thread 1, i = 4

thread 2, i = 7

thread 3, i = 9 —

thread 0, i = 2 —

thread 1, i = 5 —

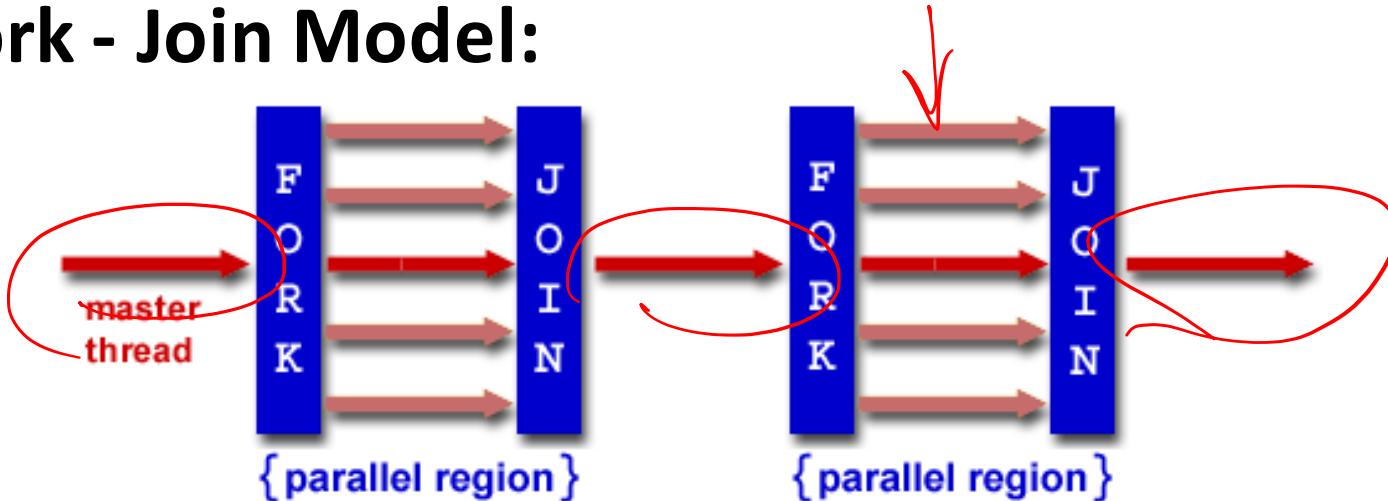
01 02 03 14 15 16 27 28 39 40

OpenMP

- C extension: no new language to learn
- Multi-threaded, shared-memory parallelism
 - Compiler Directives, **#pragma**
 - Runtime Library Routines, **#include <omp.h>**
- **#pragma**
 - Ignored by compilers unaware of OpenMP
 - Same source for multiple architectures
 - E.g. same program for 1 & 16 cores
- Only works with shared memory

OpenMP Programming Model

- **Fork - Join Model:**



- OpenMP programs begin as single process (*master thread*)
 - Sequential execution
- When parallel region is encountered
 - Master thread “forks” into team of parallel threads
 - Executed simultaneously
 - At end of parallel region, parallel threads “join”, leaving only master thread
- Process repeats for each parallel region
 - Amdahl’s law?

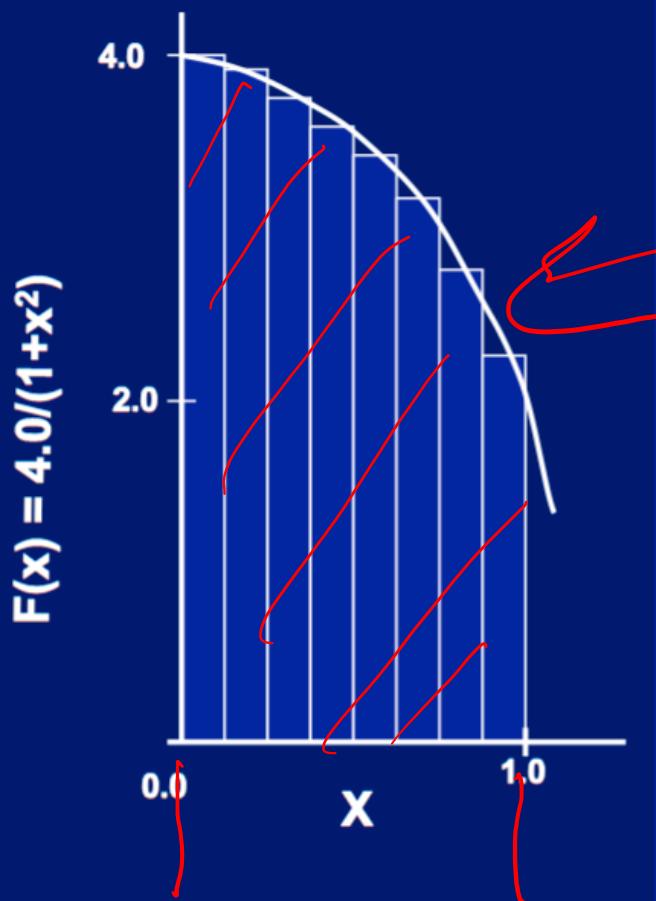
What Kind of Threads?

- OpenMP threads are operating system (software) threads.
- OS will multiplex requested OpenMP threads onto available hardware threads.
- Hopefully each gets a real hardware thread to run on, so no OS-level time-multiplexing.
- But other tasks on machine can also use hardware threads!
- Be “careful” (?) when timing results for project 4!
 - 5AM?
 - Job queue?

Example 2: computing π

Numerical Integration

Mathematically, we know that:



$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .

Sequential π

```
#include <stdio.h>

void main () {
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum = 0.0;
    for (int i=0; i<num_steps; i++) {
        double x = (i+0.5) *step;
        sum += 4.0*step/(1.0+x*x);
    }
    printf ("pi = %6.12f\n", sum);
}
```

pi = 3.142425985001

- Resembles π , but not very accurate
- Let's increase **num_steps** and parallelize

Parallelize (1) ...

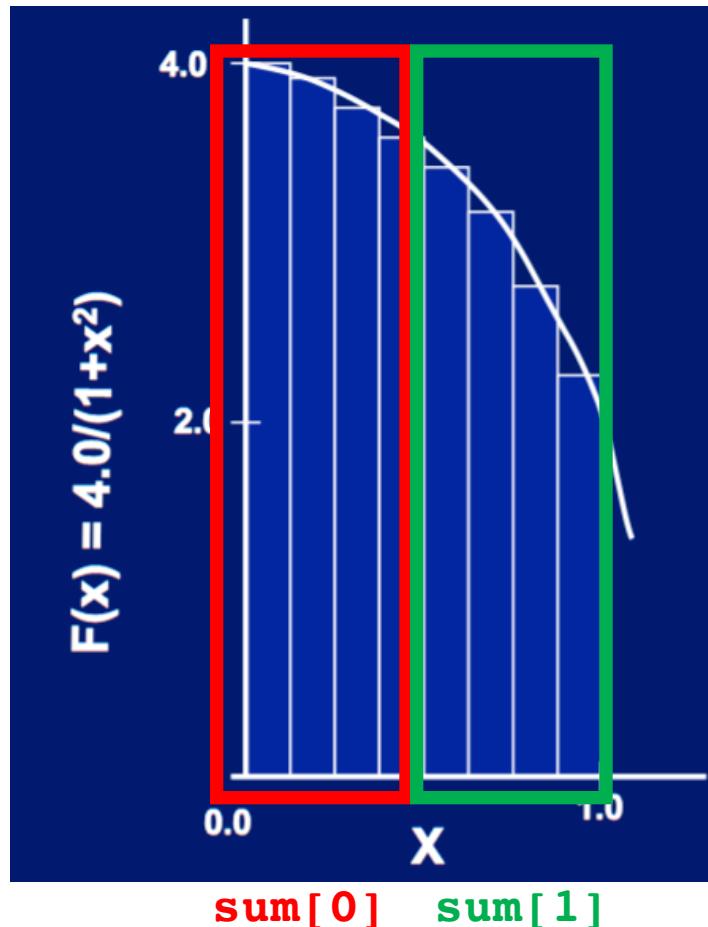
```
#include <stdio.h>

void main () {
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum = 0.0;
    #pragma parallel for
    for (int i=0; i<num_steps; i++) {
        double x = (i+0.5) *step;
        sum += 4.0*step/(1.0+x*x);
    }
    printf ("pi = %6.12f\n", sum);
}
```



- Problem: each threads needs access to the shared variable **sum**
- Code runs sequentially ...

Parallelize (2) ...



1. Compute
sum[0] and **sum[2]**
in parallel
2. Compute
sum = sum[0] + sum[1]
sequentially

Parallel π

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
        printf("i =%3d, id =%3d\n", i, id);
    }
}
double pi = 0;
for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
printf ("pi = %6.12f\n", pi);
```

Trial Run

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
        printf("i = %3d, id = %3d\n", i, id);
    }
}
double pi = 0;
for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
printf ("pi = %6.12f\n", pi);}
```

i = 1, id = 1	i = 0, id = 0	i = 2, id = 2	i = 3, id = 3	i = 5, id = 1	i = 4, id = 0	i = 6, id = 2	i = 7, id = 3	i = 9, id = 1	i = 8, id = 0	pi = 3.142425985001
---------------	---------------	---------------	---------------	---------------	---------------	---------------	---------------	---------------	---------------	---------------------

Scale up: num_steps = 10^6

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 1000000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
        // printf("i =%3d, id =%3d\n", i, id);
    }
}
double pi = 0;
for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
printf ("pi = %6.12f\n", pi);
```

pi = 3.141592653590

You verify how many digits
are correct...

Can we Parallelize Computing **sum**

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 1000;
    const long num_steps = 100000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    double pi = 0;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
    }
    pi += sum[id];
}
printf ("pi = %6.12f\n", pi);}
```

Always looking for ways to
beat Amdahl's Law ...

Summation inside parallel section

- Insignificant speedup in this example, but ...
- **pi = 3.138450662641**
- Wrong! And value changes between runs?!
- What's going on?

Your Turn

What are the possible values of `*($s0)` after executing this code by 2 *concurrent* threads?

```
# *($s0) = 100
lw    $t0,0($s0)
addi $t0,$t0,1
sw    $t0,0($s0)
```

Answer	<code>*(\$s0)</code>
A	100 or 101
B	101
C	101 or 102
D	100 or 101 or 102
E	100 or 101 or 102 or 103

Your Turn

What are the possible values of `*($s0)` after executing this code by 2 *concurrent* threads?

```
# *($s0) = 100
lw    $t0,0($s0)
addi $t0,$t0,1
sw    $t0,0($s0)
```

Answer	<code>*(\$s0)</code>
C	101 or 102

- 102 if the threads enter code section sequentially
- 101 if both execute `lw` before either runs `sw`
- one thread sees “stale” data

What's going on?

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 1000;
    const long num_steps = 100000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    double pi = 0;
    omp_set_num_threads(NUM_THREADS);

#pragma omp parallel
{
    int id = omp_get_thread_num();
    for (int i=id; i<num_steps; i+=NUM_THREADS) {
        double x = (i+0.5) *step;
        sum[id] += 4.0*step/(1.0+x*x);
    }
    pi += sum[id];
}
printf ("pi = %6.12f\n", pi);
```

- Operation is really $\text{pi} = \text{pi} + \text{sum}[id]$
- What if >1 threads reads current (same) value of **pi**, computes the sum, and stores the result back to **pi**?
- Each processor reads same intermediate value of **pi**!
- Result depends on who gets there when
 - A “race” → result is not deterministic

Agenda

- MIMD - multiple programs simultaneously
- Threads
- Parallel programming: OpenMP
- **Synchronization primitives**
- Synchronization in OpenMP
- And, in Conclusion ...

Synchronization

- Problem:
 - Limit access to shared resource to 1 actor at a time
 - E.g. only 1 person permitted to edit a file at a time
 - otherwise changes by several people get all mixed up
 - Solution:
 - Take turns:
 - Only one person gets the microphone & talks at a time
 - Also good practice for classrooms, btw ...
- 
- A photograph showing multiple hands reaching up towards a single microphone, symbolizing synchronization or taking turns. The image is watermarked with "gettyimages" and "Don Bayley". A small number "155149636" is visible in the bottom left corner of the image.

Locks

- Computers use locks to control access to shared resources
 - Serves purpose of microphone in example
 - Also referred to as “semaphore”
- Usually implemented with a variable
 - **int lock;**
 - 0 for unlocked
 - 1 for locked

Synchronization with locks

```
// wait for lock released  
while (lock != 0) ;  
// lock == 0 now (unlocked)
```

```
// set lock  
lock = 1;
```

```
// access shared resource ...  
// e.g. pi  
// sequential execution! (Amdahl ...)
```

```
// release lock  
lock = 0;
```

Lock Synchronization

Thread 1

`while (lock !=`

`lock = 1;`

`// critical section`

`lock`



`= 0) ;`

: set,

they got

and set the lock!

**Try as you want, this problem has no solution,
not even at the assembly level.**

Unless we introduce new instructions, that is!

Hardware Synchronization

- Solution:
 - Atomic read/write
 - Read & write in single instruction
 - No other access permitted between read and write
 - Note:
 - Must use *shared memory* (multiprocessing)
- Common implementations:
 - Atomic swap of register \leftrightarrow memory
 - Pair of instructions for “linked” read and write
 - write fails if memory location has been “tampered” with after linked read
 - MIPS uses this solution

MIPS Synchronization Instructions

- *Load linked:*

ll \$rt, off(\$rs)



- Reads memory location (like **lw**)
- Also sets (hidden) “link bit”
- Link bit is reset if memory location (**off(\$rs)**) is accessed

- *Store conditional:*

sc \$rt, off(\$rs)



- Stores **off(\$rs) = \$rt** (like **sw**)
- Sets **\$rt=1** (success) if link bit is set
- i.e. no (other) process accessed **off(\$rs)** since **ll**
- Sets **\$rt=0** (failure) otherwise
- Note: **sc** clobbers **\$rt**, i.e. changes its value

Lock Synchronization

\$t0 = 1 before calling ll: minimize time between ll and sc

Broken Synchronization

```
while (lock != 0) ;
```

```
lock = 1;
```

```
// critical section
```

```
lock = 0;
```

Fix (lock is at location \$s1)

```
Try: addiu $t0,$zero,1  
      ll    $t1,0($s1)  
      bne $t1,$zero,Try  
      sc    $t0,0($s1)  
      beq $t0,$zero,Try
```

Locked:

critical section

Unlock:

```
sw $zero,0($s1)
```

Agenda

- MIMD - multiple programs simultaneously
- Threads
- Parallel programming: OpenMP
- Synchronization primitives
- **Synchronization in OpenMP**
- And, in Conclusion ...

OpenMP Locks

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main(void) {
    omp_lock_t lock;
    omp_init_lock(&lock);

#pragma omp parallel
{
    int id = omp_get_thread_num();

    // parallel section
    // ...

    omp_set_lock(&lock);
    // start sequential section
    // ...
    printf("id = %d\n", id);

    // end sequential section
    omp_unset_lock(&lock);

    // parallel section
    // ...

}

    omp_destroy_lock(&lock);
}
```

Synchronization in OpenMP

- Typically are used in libraries of higher level parallel programming constructs
- E.g. OpenMP offers **\$pragmas** for common cases:
 - critical
 - atomic
 - barrier
 - ordered
- OpenMP offers many more features
 - see online documentation
 - or tutorial at
 - <http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>

OpenMP critical

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 1000;
    const long num_steps = 100000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    double pi = 0;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        for (int i=id; i<num_steps; i+=NUM_THREADS) {
            double x = (i+0.5) *step;
            sum[id] += 4.0*step/(1.0+x*x);
        }
        #pragma omp critical
        pi += sum[id];
    }
    printf ("pi = %6.12f\n", pi);
}
```

The Trouble with Locks ...

- ... is ***dead-locks***
- Consider 2 cooks sharing a kitchen
 - Each cooks a meal that requires salt and pepper (locks)
 - Cook 1 grabs salt
 - Cook 2 grabs pepper
 - Cook 1 notices s/he needs pepper
 - it's not there, so s/he waits
 - Cook 2 realizes s/he needs salt
 - it's not there, so s/he waits
- A not so common cause of cook starvation
 - But deadlocks are possible in parallel programs
 - Very difficult to debug
 - **malloc / free** is easy ...

Agenda

- MIMD - multiple programs simultaneously
- Threads
- Parallel programming: OpenMP
- Synchronization primitives
- Synchronization in OpenMP
- **And, in Conclusion ...**

And in Conclusion, ...

- Sequential software execution speed is limited
- Parallel processing is the only path to higher performance
 - SIMD: instruction level parallelism
 - Implemented in all high performance CPUs today (x86, ARM, ...)
 - Partially supported by compilers
 - MIMD: thread level parallelism
 - Multicore processors
 - Supported by Operating Systems (OS)
 - Requires programmer intervention to exploit at single program level
 - E.g. OpenMP
 - SIMD & MIMD for maximum performance
- Synchronization
 - Requires hardware support: specialized assembly instructions
 - Typically use higher-level support
 - Beware of deadlocks