

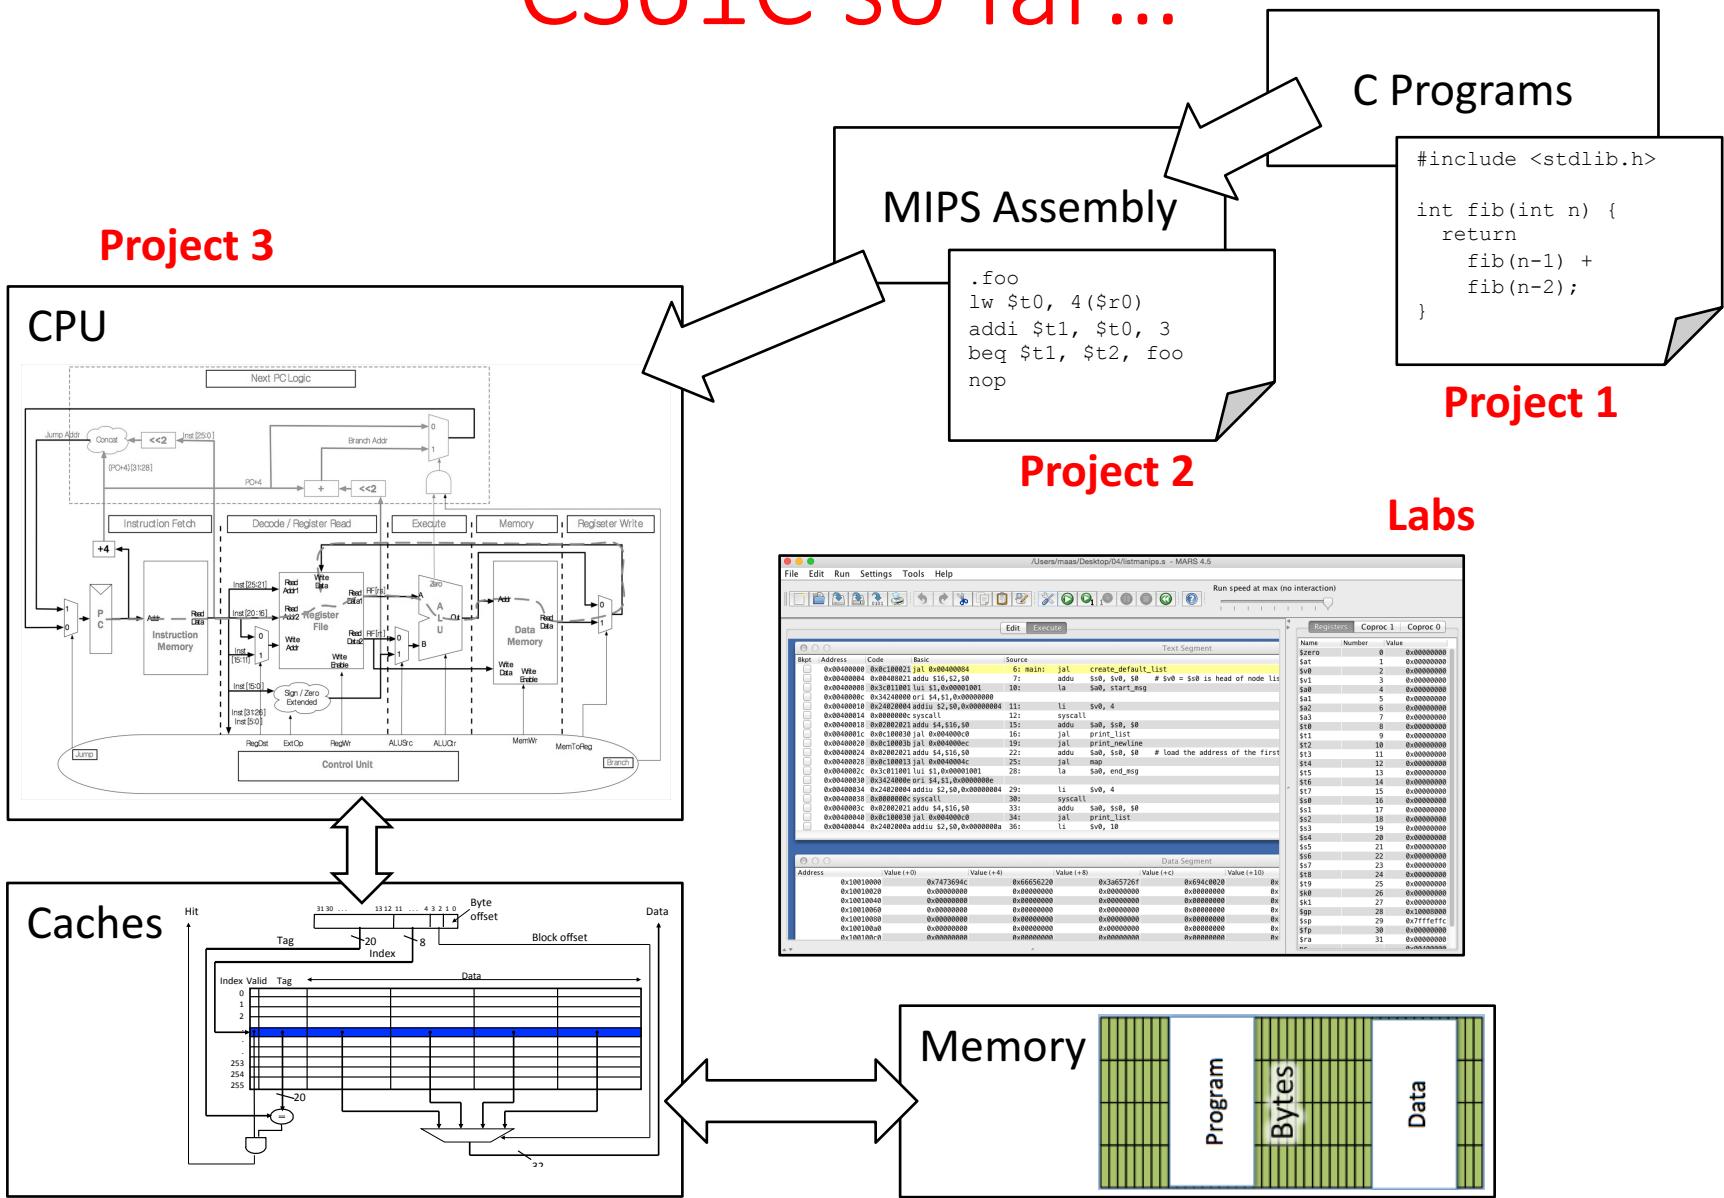
CS 61C: Great Ideas in Computer Architecture

Lecture 22: *Operating System*

Bernhard Boser & Randy Katz

<http://inst.eecs.berkeley.edu/~cs61c>

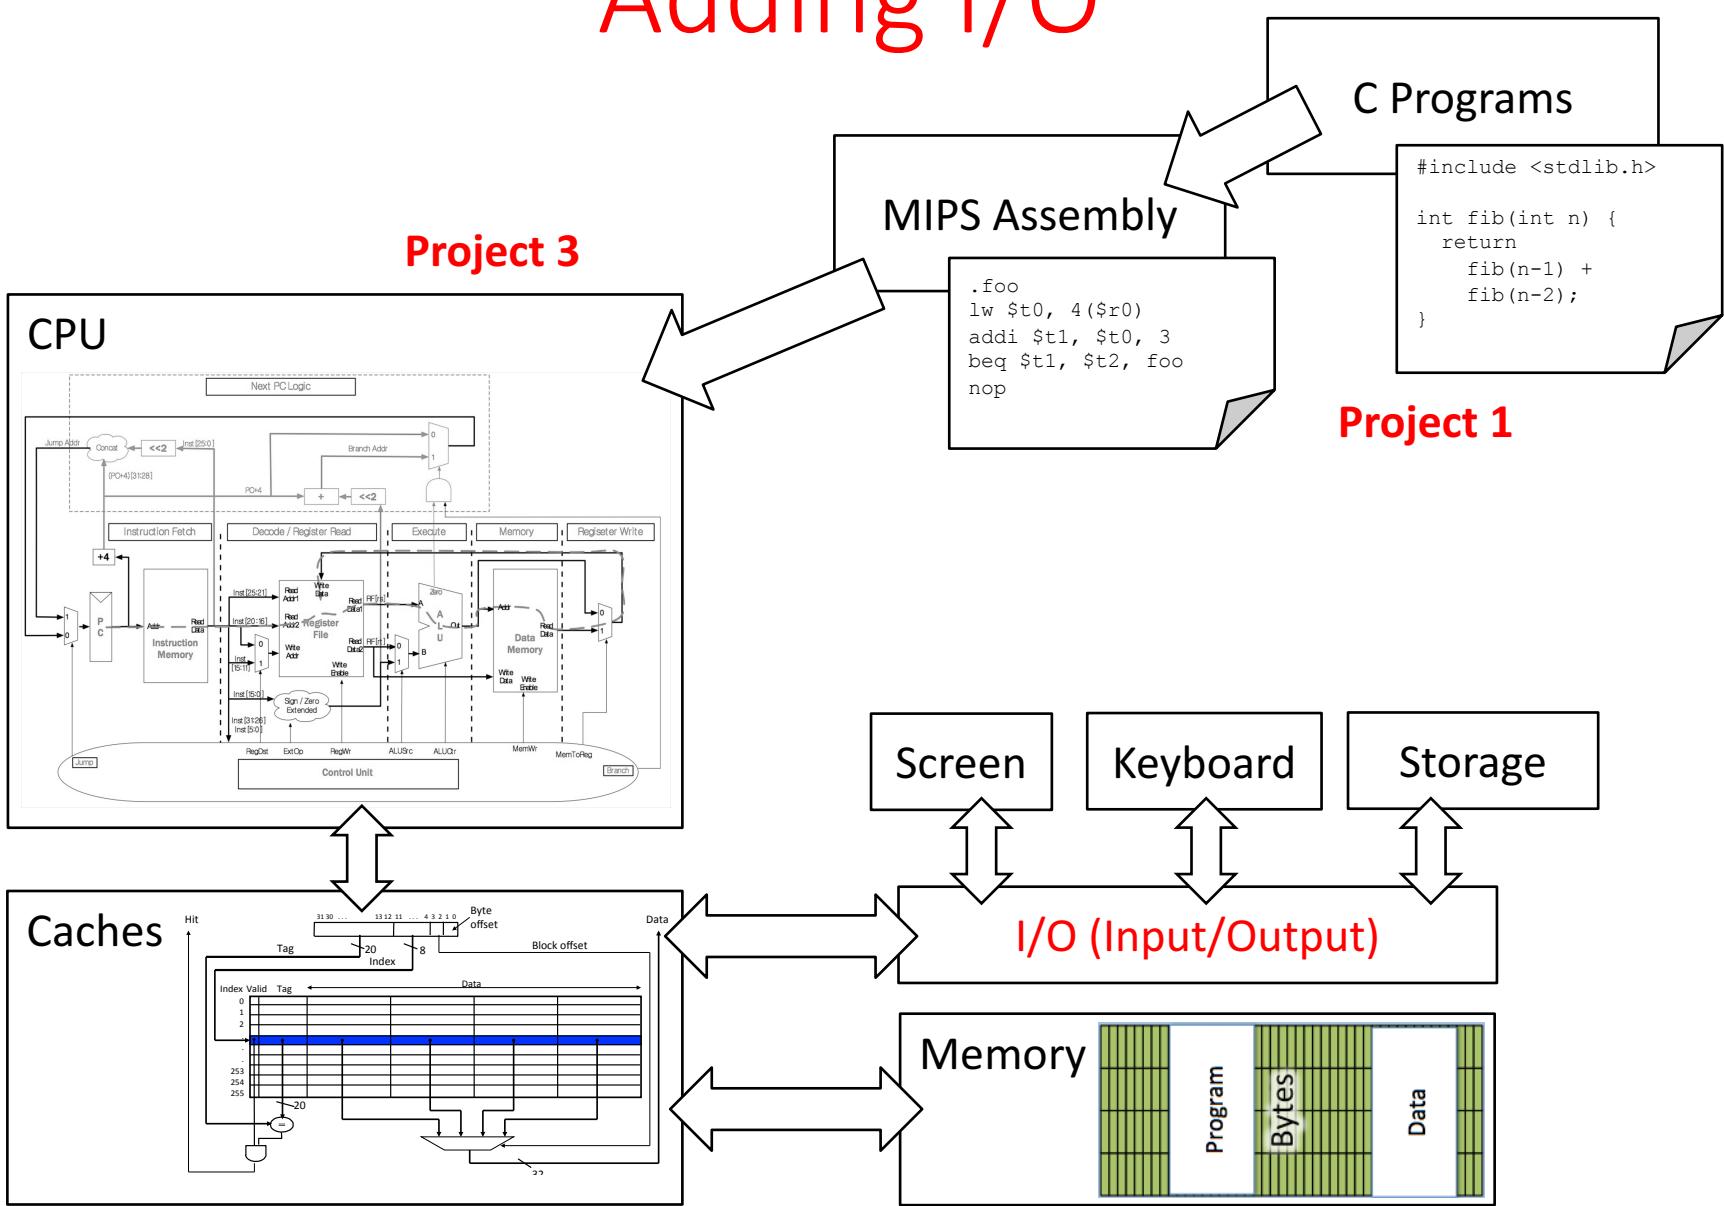
CS61C so far...



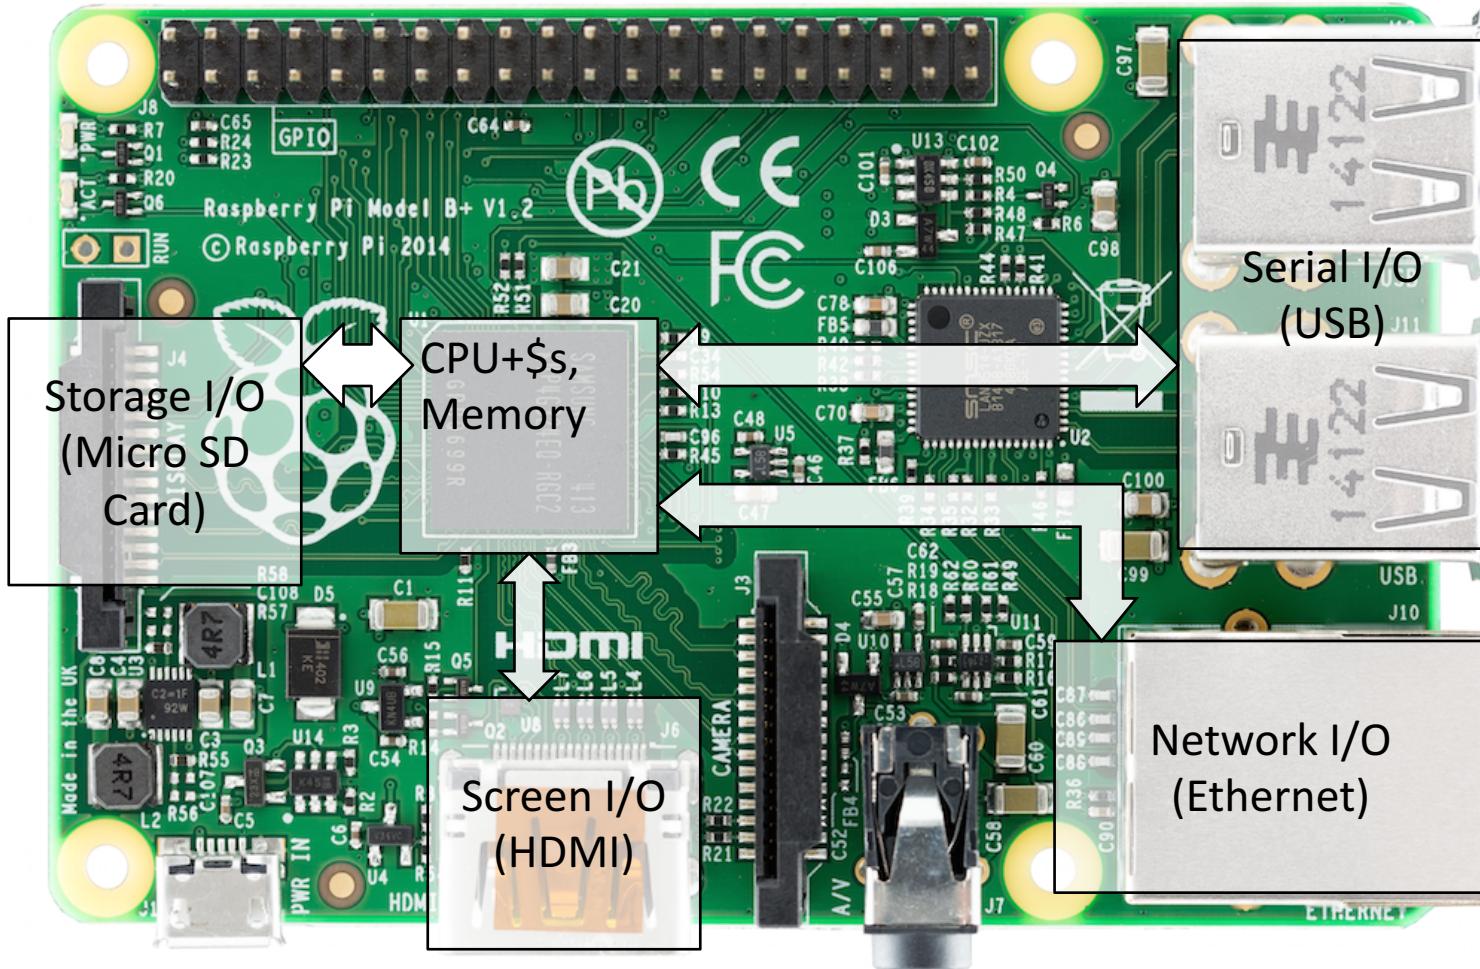
So how is this any different?



Adding I/O



Raspberry Pi (\$40 on Amazon)

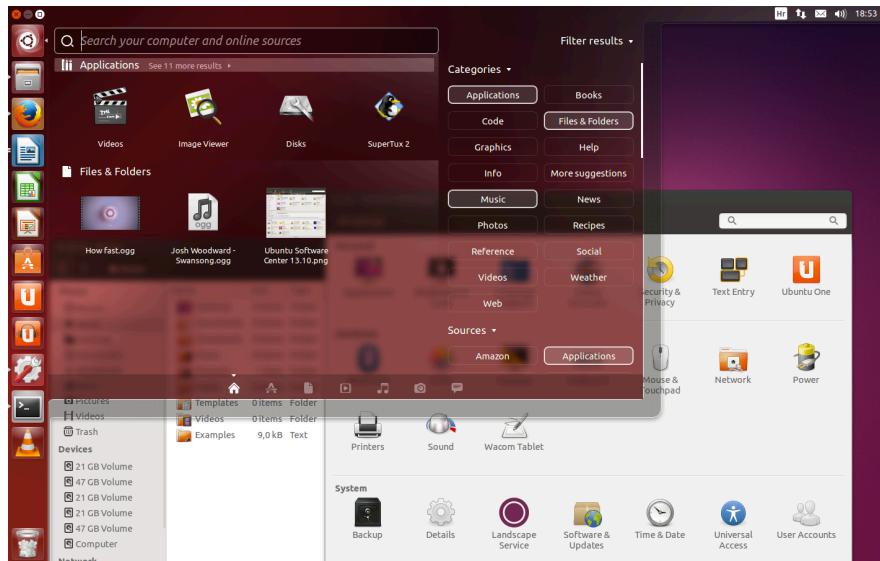


It's a real computer!



But wait...

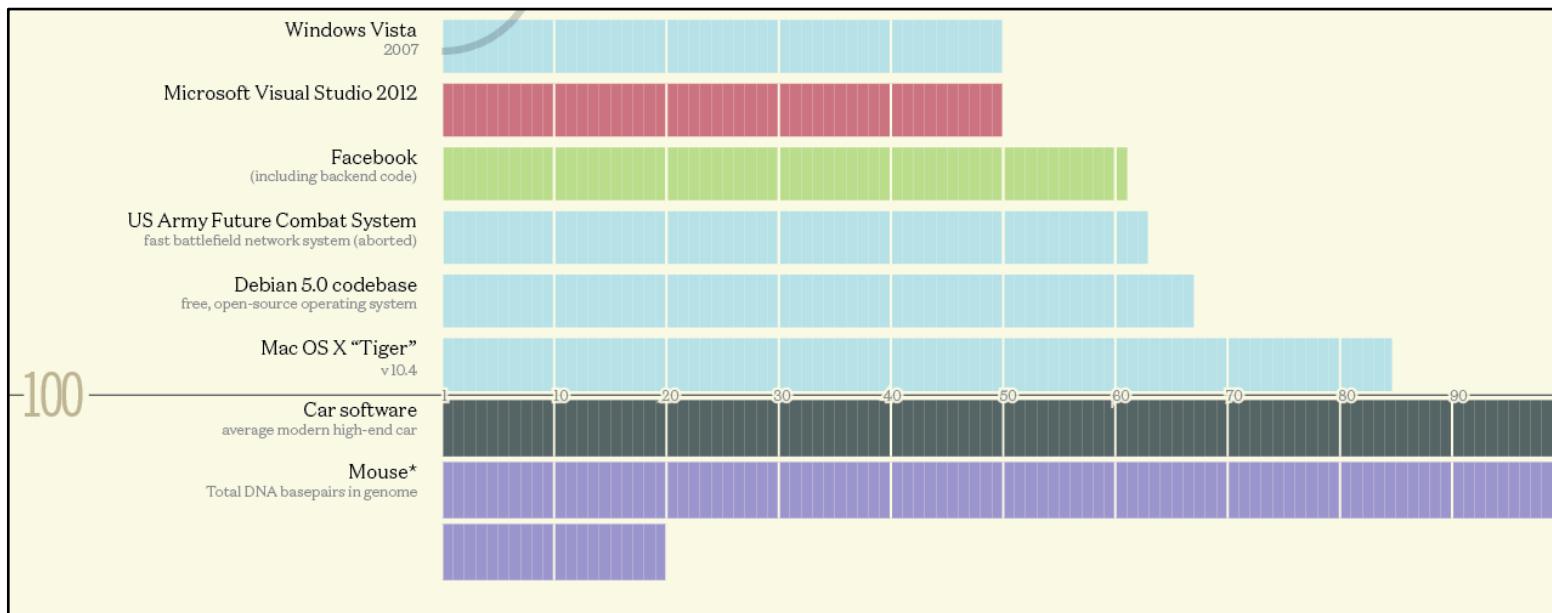
- That's not the same! When we run MARS, it only executes one program and then stops.
- When I switch on my computer, I get this:



Yes, but that's *just* software! **The Operating System (OS)**

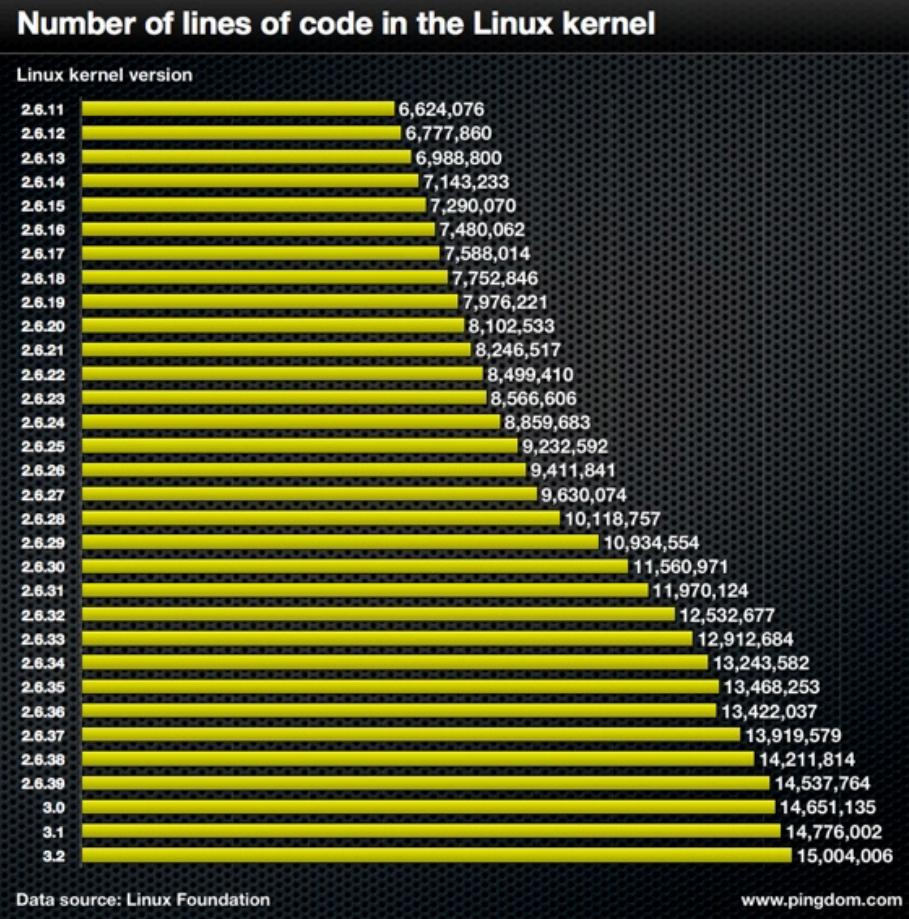
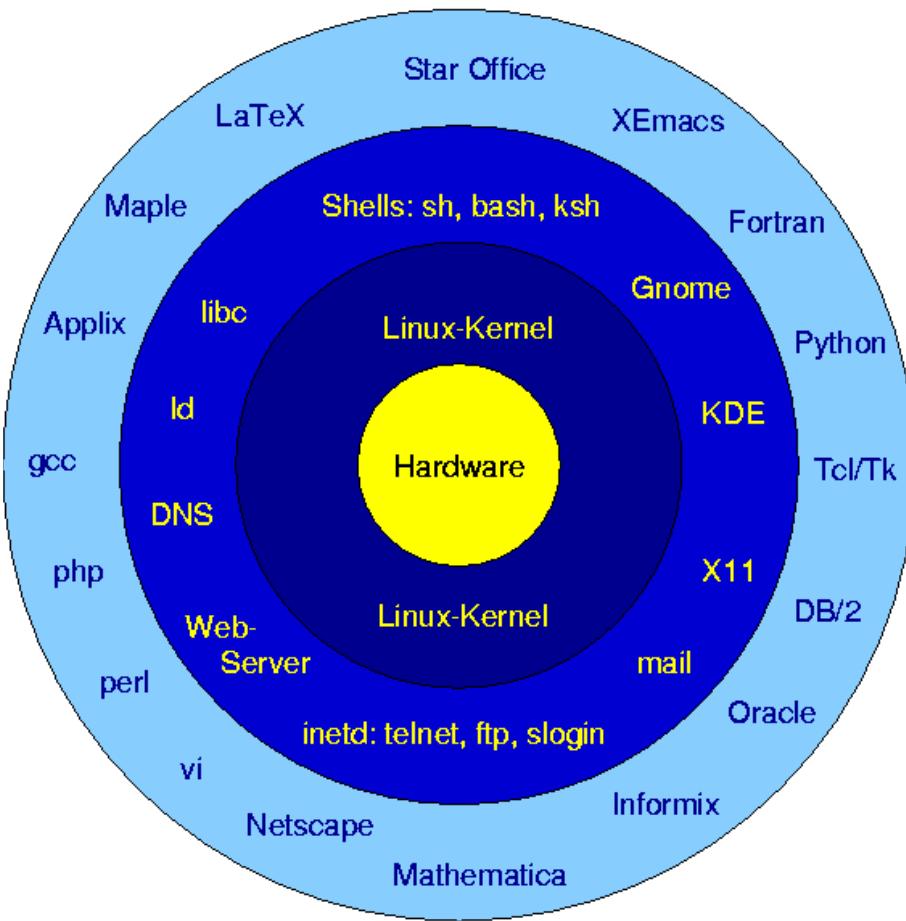
Well, “just software”

- The biggest piece of software on your machine?
- How many lines of code? These are guesstimates:



Codebases (in millions of lines of code). CC BY-NC 3.0 — David McCandless © 2013
<http://www.informationisbeautiful.net/visualizations/million-lines-of-code/>

Operating System



What does the OS do?

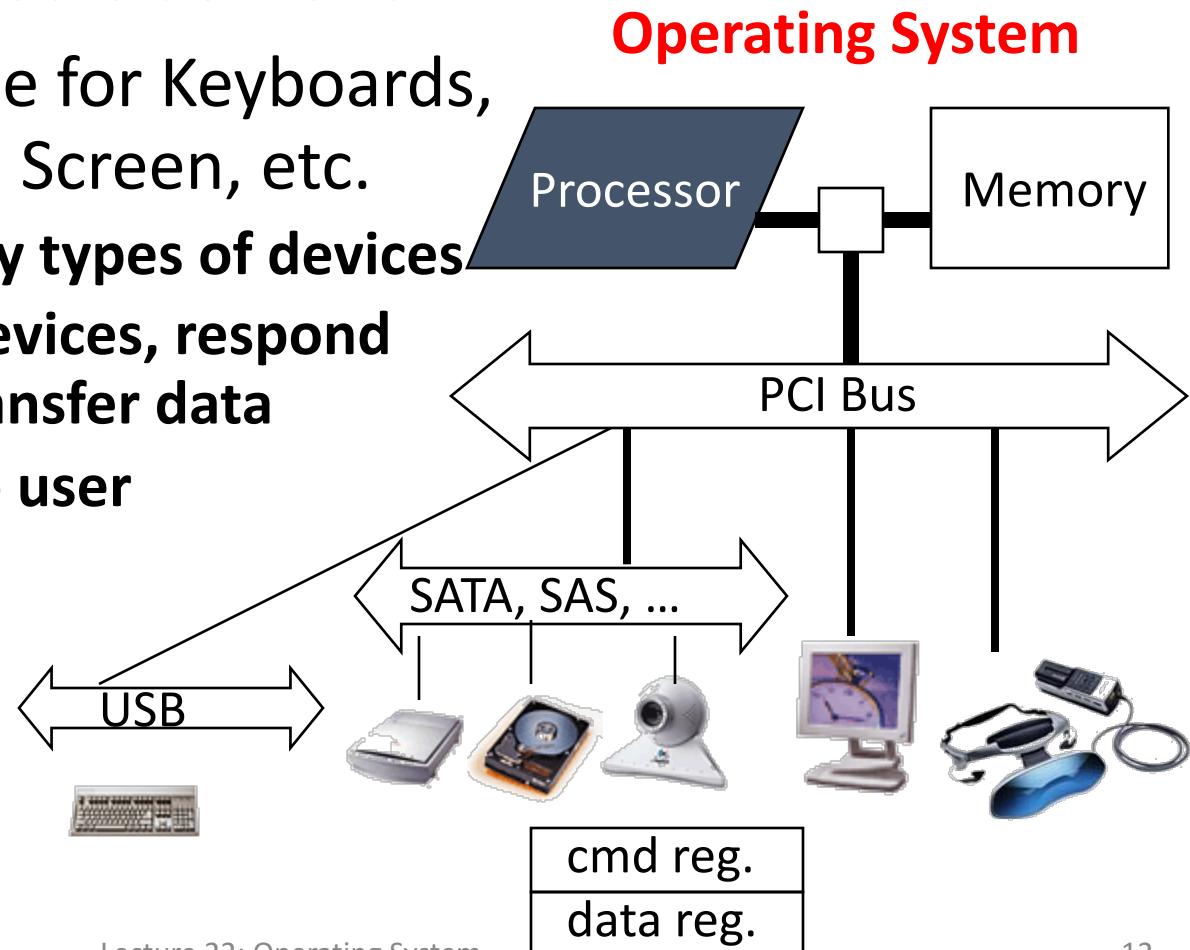
- OS is first thing that runs when computer starts
- Finds and controls all devices in the machine in a general way
 - Relying on hardware specific “device drivers”
- Starts services (100+)
 - File system,
 - Network stack (Ethernet, WiFi, Bluetooth, ...),
 - TTY (keyboard),
 - ...
- Loads, runs and manages programs:
 - Multiple programs at the same time (time-sharing)
 - Isolate programs from each other (isolation)
 - Multiplex resources between applications (e.g., devices)

Agenda

- Devices and I/O
- Polling
- Interrupts
- OS Boot Sequence
- Multiprogramming/time-sharing

How to interact with devices?

- Assume a program running on a CPU. How does it interact with the outside world?
- Need I/O interface for Keyboards, Network, Mouse, Screen, etc.
 - Connect to many types of devices
 - Control these devices, respond to them, and transfer data
 - Present them to user programs so they are useful

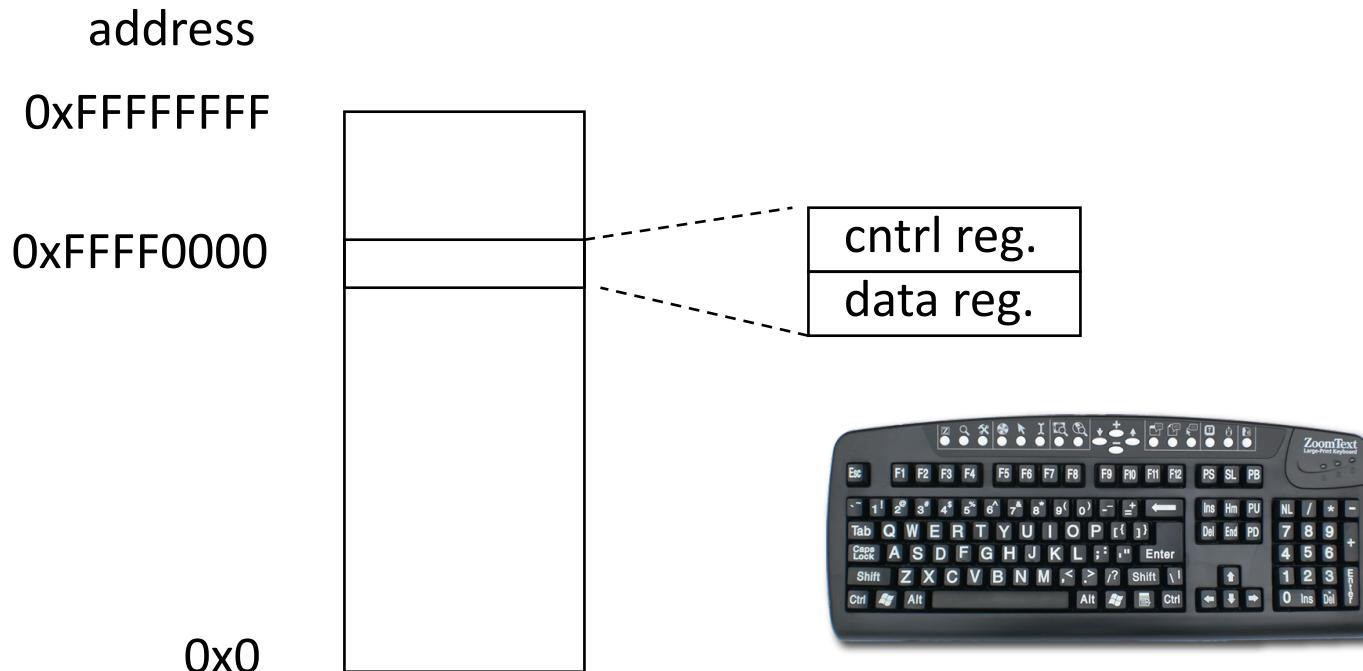


Instruction Set Architecture for I/O

- What must the processor do for I/O?
 - Input: read a sequence of bytes
 - Output: write a sequence of bytes
- Interface options
 - a) Special input/output instructions & hardware
 - b) Memory mapped I/O
 - Portion of address space dedicated to I/O
 - I/O device registers there (no memory)
 - Use normal load/store instructions, e.g. **lw / sw**
 - Very common, used by MIPS

Memory Mapped I/O

- Certain addresses are not regular memory
- Instead, they correspond to registers in I/O devices



Processor-I/O Speed Mismatch

- 1GHz microprocessor I/O throughput:
 - 4 Gi-B/s (**1w/sw**)
 - Typical I/O data rates:
 - 10 B/s (keyboard)
 - 100 Ki-B/s (Bluetooth)
 - 60 Mi-B/s (USB 2)
 - 100 Mi-B/s (Wifi, depends on standard)
 - 125 Mi-B/s (G-bit Ethernet)
 - 550 Mi-B/s (cutting edge SSD)
 - 1.25 Gi-B/s (USB 3.1 Gen 2)
 - 6.4 GiB/s (DDR3 DRAM)
 - These are peak rates – actual throughput is lower
- Common I/O devices neither deliver nor accept data matching processor speed



No Endorsement. Just a break from the lecture.

Agenda

- Devices and I/O
- **Polling**
- Interrupts
- OS Boot Sequence
- Multiprogramming/time-sharing

Processor Checks Status before Acting

- Device registers generally serve 2 functions:
 - **Control Register**, says it's OK to read/write (I/O ready) [think of a flagman on a road]
 - **Data Register**, contains data
- Processor reads from Control Register in loop
 - waiting for device to set **Ready** bit in Control reg ($0 \rightarrow 1$)
 - Indicates “data available” or “ready to accept data”
- Processor then loads from (input) or writes to (output) data register
 - I/O device resets control register bit ($1 \rightarrow 0$)
- Procedure called “**Polling**”

I/O Example (polling)

- Input: Read from keyboard into \$v0

```
lui    $t0, 0xffff #ffff0000 (io addr)
Waitloop:      lw     $t1, 0($t0) #read control
                  andi  $t1,$t1,0x1 #ready bit
                  beq   $t1,$zero, Waitloop
                  lw     $v0, 4($t0) #data
```

- Output: Write to display from \$a0

```
lui    $t0, 0xffff #ffff0000
Waitloop:      lw     $t1, 8($t0) #write control
                  andi  $t1,$t1,0x1 #ready bit
                  beq   $t1,$zero, Waitloop
                  sw     $a0,12($t0) #data
```

“Ready” bit is from processor’s point of view!

Cost of Polling?

- Assume for a processor with
 - 1GHz clock rate
 - Taking 400 clock cycles for a polling operation
 - Call polling routine
 - Check device (e.g. keyboard or wifi input available)
 - Return
 - What's the percentage of processor time spent polling?
- Example:
 - Mouse
 - Poll 30 times per second
 - Set by requirement not to miss any mouse motion (which would lead to choppy motion of the cursor on the screen)

Percent Processor Time to Poll Mouse

- Mouse Polling [instructions/sec]
 - = 30 [polls/s] * 400 [instructions/poll]
 - = 12K [instructions/s]
- % Processor for polling:
 - Total processor throughput: $T = 1 \times 10^9$ [instructions/s]
 - Polling: $P = 12 \times 10^3$ [instructions/s]
 - Ratio: $P/T = 0.0012\%$
- Negligible:
 - Polling mouse has little impact on processor throughput

Clicker Time

Hard disk: transfers data in 16-Byte chunks and can transfer at 16 MB/second. No transfer can be missed. What percentage of processor time is spent in polling (assume 1GHz clock)?

- A: 2%
- B: 4%
- C: 20%
- D: 40%
- E: 80%

% Processor time needed poll disk

- Frequency of Polling Disk
 $= 16 \text{ [MB/s]} / 16 \text{ [B/poll]} = 1\text{M} \text{ [polls/s]}$
- Disk Polling, instructions/sec
 $= 1\text{M} \text{ [polls/s]} * 400 \text{ [instructions/poll]}$
 $= 400\text{M} \text{ [instructions/s]}$
- % Processor for polling:
 $400 * 10^6 \text{ [instructions/s]} / 1 * 10^9 \text{ [instructions/s]} = 40\%$
→ Unacceptable

What is the alternative to polling?

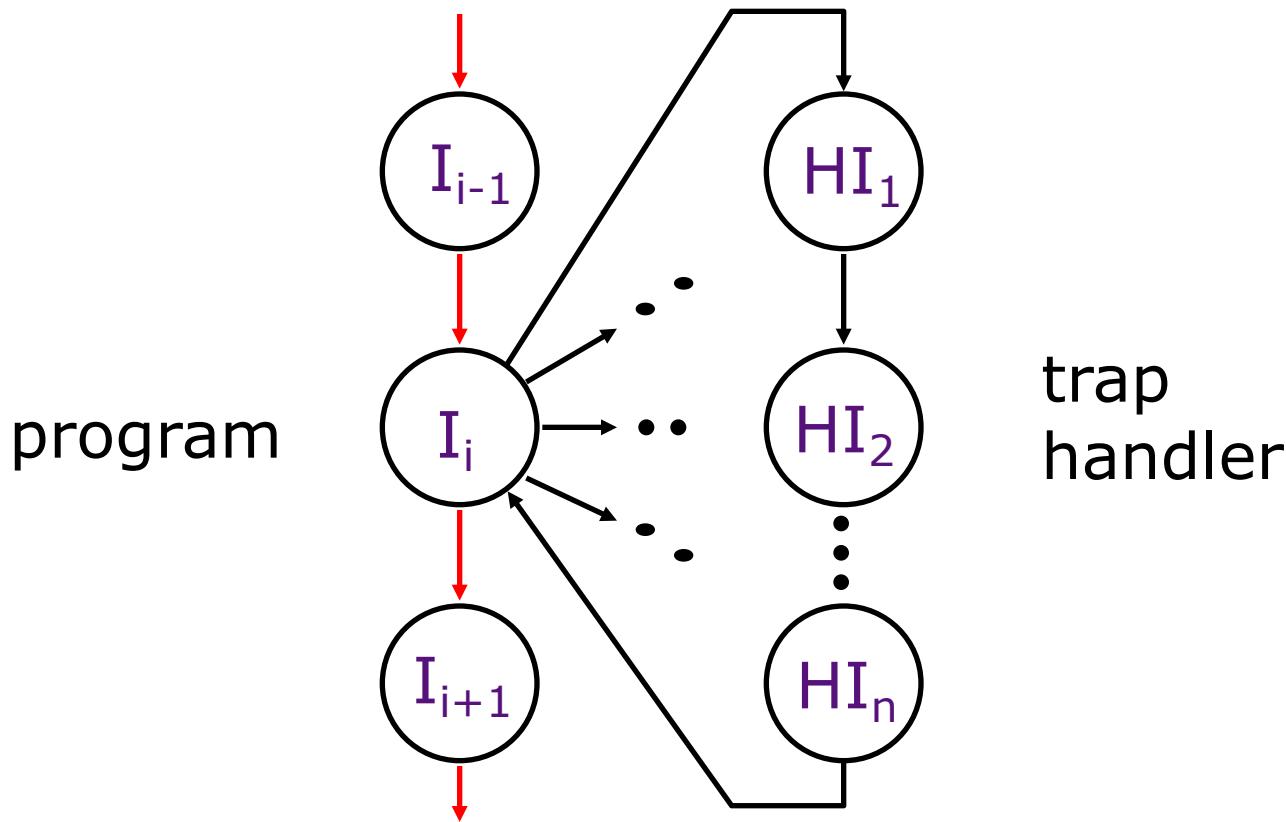
- Polling wastes processor resources
- Akin to waiting at the door for guests to show up
 - What about a bell?
- Computer lingo for bell:
 - **Interrupt**
 - Occurs when I/O is ready or needs attention
 - Interrupt current program
 - Transfer control to special code “**interrupt handler**”

Agenda

- Devices and I/O
- Polling
- **Interrupts**
- OS Boot Sequence
- Multiprogramming/time-sharing

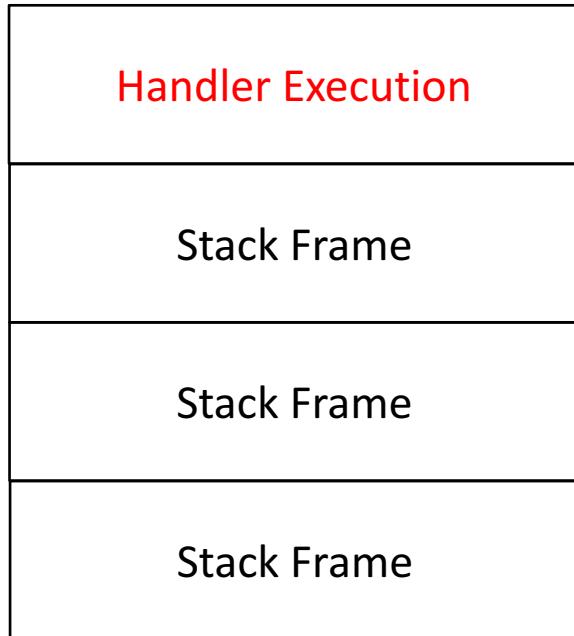
Traps/Interrupts/Exceptions:

altering the normal flow of control (P&H Section 4.9)



An *external or internal event* that needs to be processed - by another program – the OS. The event is often unexpected from original program's point of view.

Interrupt-driven I/O



```
Label: sll $t1,$s3,2  
       addu $t1,$t1,$s5  
       lw    $t1,0($t1) ←  
       or    $s1,$s1,$t1  
       addu $s3,$s3,$s4  
       bne  $s3,$s2,Label
```

1. Incoming interrupt suspends instruction stream
2. Looks up the vector (function address) of a handler in an interrupt vector table stored within the CPU
3. Perform a jal to the handler (save PC in EPC* register)
4. Handler run on current stack and returns on finish (thread doesn't notice that a handler was run)

```
handler: lui  $t0, 0xffff  
         lw   $t1, 0($t0)  
         andi $t1,$t1,0x1  
         lw   $v0, 4($t0)  
         sw   $t1, 8($t0)  
         ret
```

Code for saving/restoring \$t0, \$t1, \$v0 not shown

Interrupt(SPI0)

CPU Vector Interrupt Table

SPI0	handler
...	...

*EPC: Exception program counter

Terminology

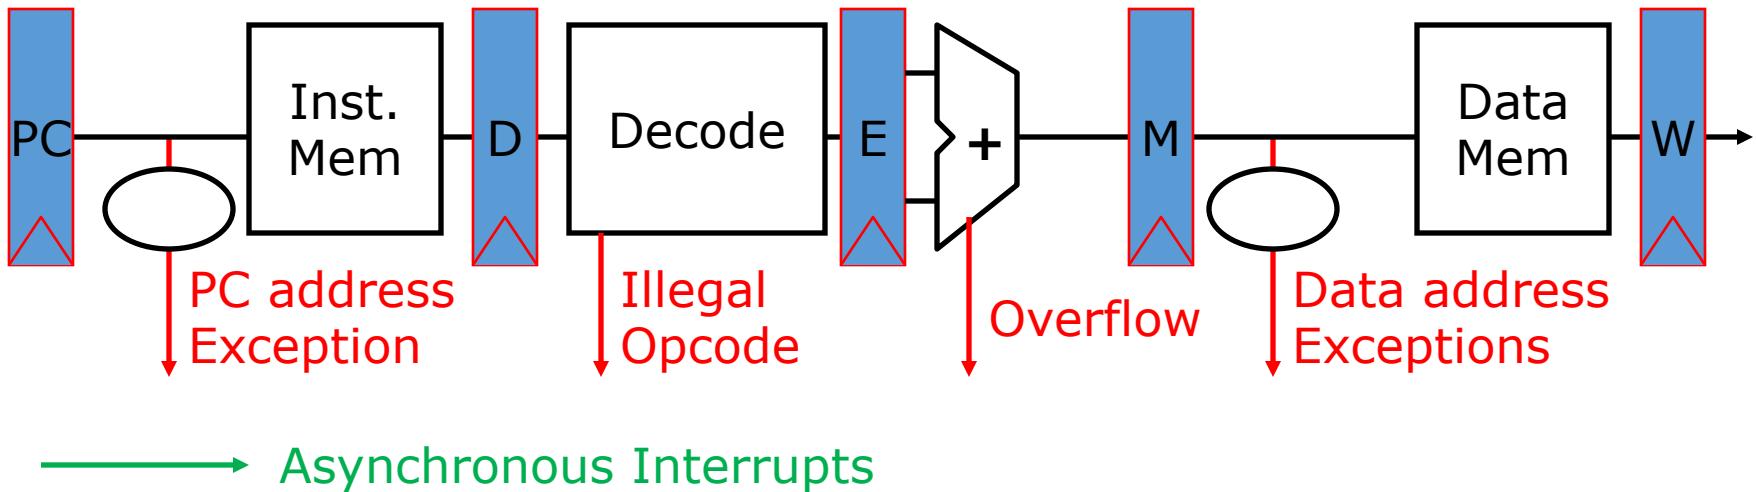
In CS61C (other definitions in use elsewhere):

- **Interrupt** – caused by an event *external* to current running program
 - E.g. key press, disk
 - Asynchronous to current program
 - Can handle interrupt on any convenient instruction
 - “Whenever it’s convenient, just don’t wait too long”
- **Exception** – caused by some event *during* execution of one instruction of current running program
 - E.g., overflow, bus error, illegal instruction
 - Synchronous
 - Must handle exception *precisely* on instruction that causes exception
 - “Drop whatever you are doing and act now”
- **Trap** – action of servicing interrupt or exception by hardware jump to “interrupt or trap handler” code

Precise Traps

- *Trap handler's view of machine state is that every instruction prior to the trapped one (e.g. overflow) has completed, and no instruction after the trap has executed.*
- Implies that handler can return from an interrupt by restoring user registers and jumping back to interrupted instruction
 - Interrupt handler software doesn't need to understand the pipeline of the machine, or what program was doing!
 - More complex to handle trap caused by an exception than interrupt
- Providing precise traps is tricky in a pipelined superscalar out-of-order processor!
 - But a requirement e.g. for
 - Virtual memory to function properly (see next lecture)

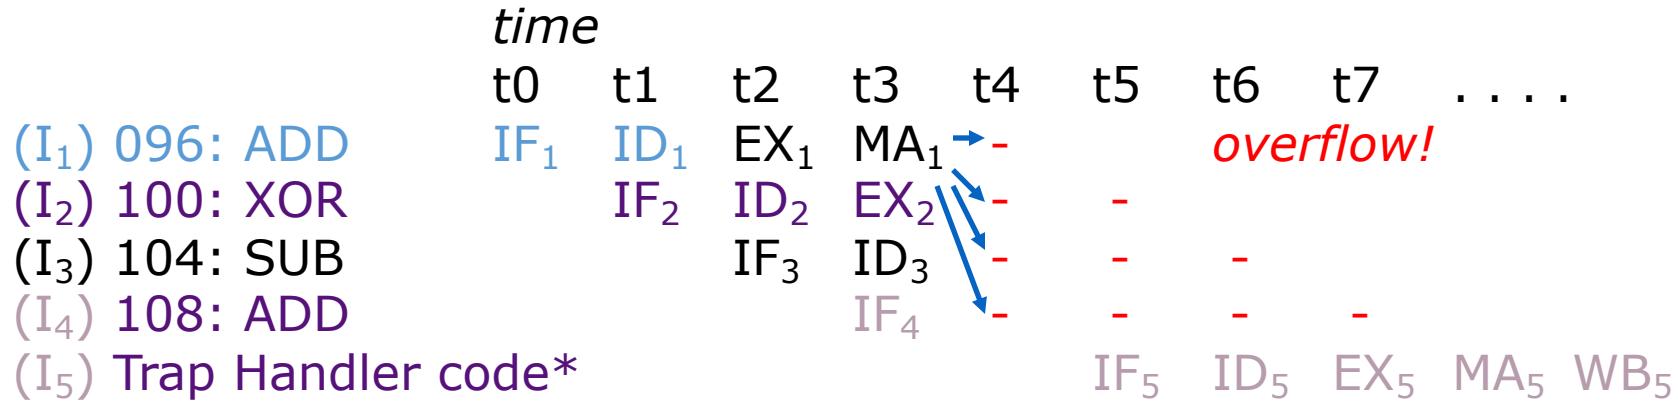
Trap Handling in 5-Stage Pipeline



Exceptions are handled *like pipeline hazards*

- 1) Complete execution of instructions before exception occurred
- 2) Flush instructions currently in pipeline
(convert to **nops** or “bubbles”)
- 3) Optionally store exception cause in status register (e.g. MIPS)
 - Indicate type of exception
 - **Note: several exceptions can occur in a single clock cycle!**
- 4) Transfer execution to trap handler

Trap Pipeline Diagram



*EPC = 100 (instruction following offending ADD)

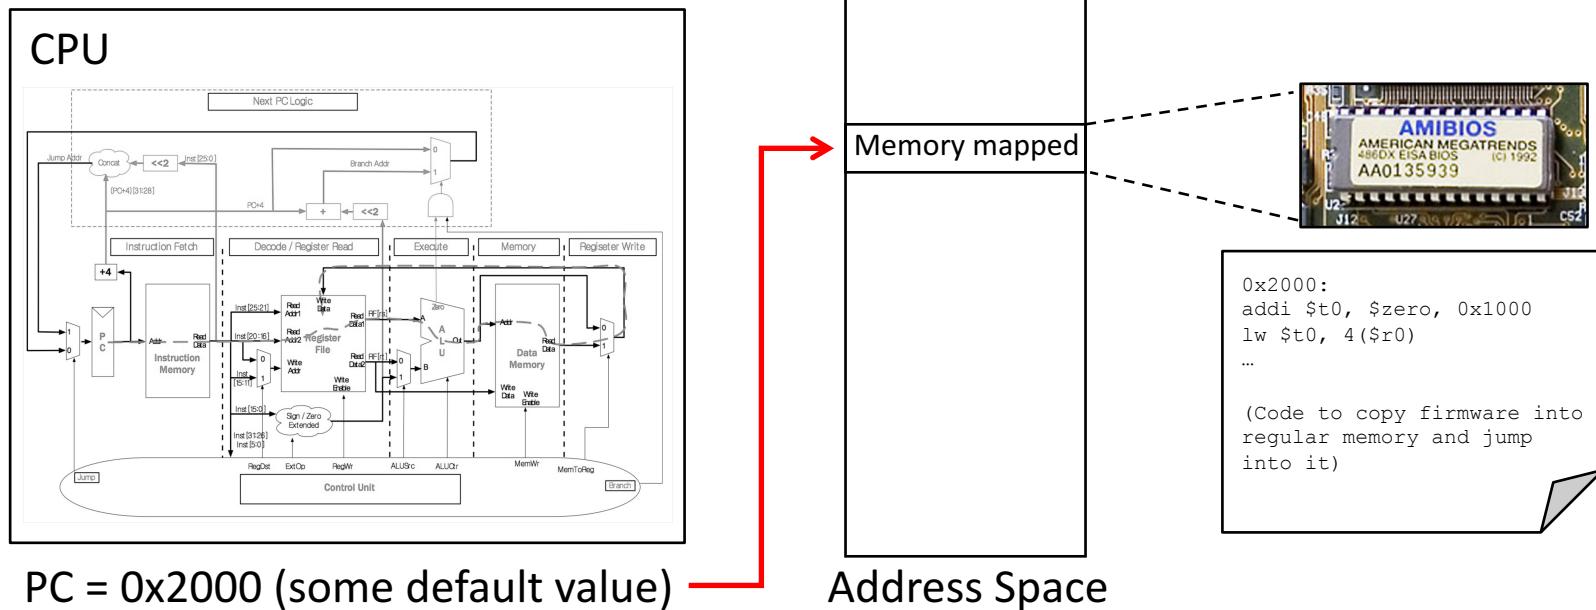


Agenda

- Devices and I/O
- Polling
- Interrupts
- **OS Boot Sequence**
- Multiprogramming/time-sharing

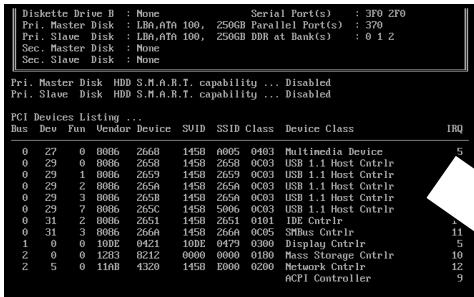
What happens at boot?

- When the computer switches on, it does the same as MARS: the CPU executes instructions from some start address (stored in Flash ROM)



What happens at boot?

1. BIOS*: Find a storage device and load first sector (block of data)



2. Bootloader (stored on, e.g., disk): Load the OS *kernel* from disk into a location in memory and jump into it.

Detailed description: A terminal window titled 'QUESTION 3:' shows a user navigating through a configuration file named 'answers.txt'. The user runs commands like 'make', 'make cmm', and 'make: 'cmm' is up to date.' The terminal also shows the path '/src/proj3/proj3_starter \$' and the command '/src/proj3/proj3_starter'.

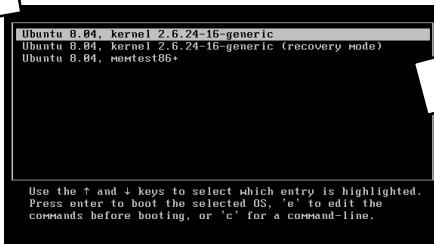
```
QUESTION 3:
conv: <speedup> x
relu: <speedup> x
pool: <speedup> x
fc: <speedup> x
softmax: <speedup> x

Which layer should we optimize? [cnn] >
<which layers>

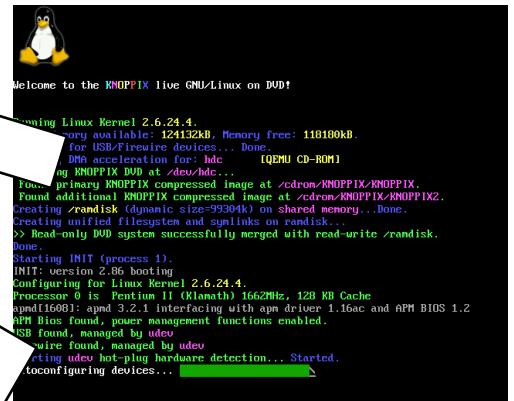
[23:04:09 Wed Apr 15 2015 cs61c-ti@hive22 Linux x86_64]
~/src/proj3/proj3_starter $ cd src/
cnn.c main.c python.c util.c

[23:04:16 Wed Apr 15 2015 cs61c-ti@hive22 Linux x86_64]
~/src/proj3/proj3_starter $ make cmm
make: 'cmm' is up to date.

[23:04:20 Wed Apr 15 2015 cs61c-ti@hive22 Linux x86_64]
~/src/proj3/proj3_starter $
```



4. Init: Launch an application that waits for input in loop (e.g., Terminal/Desktop/...)



3. OS Boot: Initialize services, drivers, etc.

*BIOS: Basic Input Output System

Launching Applications

- Applications are called “processes” in most OSs.
 - Thread: shared memory
 - Process: separate memory
 - Both threads and processes run (pseudo) simultaneously
- Apps are started by another process (e.g. shell) calling an OS routine (using a “syscall”)
 - Depends on OS, but Linux uses **fork** to create a new process, and **execve** to load application.
- Loads executable file from disk (using the file system service) and puts instructions & data into memory (.text, .data sections), prepares stack and heap.
- Set argc and argv, jump to start of main.
- Shell waits for main to return (**join**)

Supervisor Mode

- If something goes wrong in an application, it could crash the entire machine. And what about malware, etc.?
- The OS enforces resource constraints to applications (e.g., access to memory, devices).
- To help protect the OS from the application, CPUs have a **supervisor mode** (set e.g. by a status bit).
 - A process can only access a subset of instructions and (physical) memory when not in supervisor mode (user mode).
 - Process can change out of supervisor mode using a special instruction, but not into it directly – only using an interrupt.
 - Supervisory mode is a bit like “superuser”
 - But used much more sparingly
(most of OS code does *not* run in supervisory mode)
 - Errors in supervisory mode often catastrophic
(blue “screen of death”, or “I just corrupted your disk”)

Syscalls

- What if we want to call an OS routine? E.g.,
 - to read a file,
 - launch a new process,
 - ask for more memory (malloc),
 - send data, etc.
- Need to perform a **syscall**:
 - set up function arguments in registers,
 - raise **software interrupt (with special assembly instruction)**
- OS will perform the operation and return to user mode
- This way, the OS can mediate access to all resources, and devices.

Agenda

- Devices and I/O
- Polling
- Interrupts
- OS Boot Sequence
- **Multiprogramming/time-sharing**

Multiprogramming

- The OS runs multiple applications at the same time.
- But not really (unless you have a core per process)
- Switches between processes very quickly (on human time scale). This is called a “context switch”.
- When jumping into process, set timer interrupt.
 - When it expires, store PC, registers, etc. (process state).
 - Pick a different process to run and load its state.
 - Set timer, change to user mode, jump to the new PC.
- Deciding what process to run is called **scheduling**.

Protection, Translation, Paging

- Supervisor mode alone is not sufficient to fully isolate applications from each other or from the OS.
 - Application could overwrite another application's memory.
 - Typically programs start at some fixed address, e.g. 0x1000
 - How can 100's of programs share memory at location 0x1000?
 - Also, may want to address more memory than we actually have (e.g., for sparse data structures).
- Solution: **Virtual Memory**.
 - Gives each process the *illusion* of a full memory address space that it has completely for itself.

And, in Conclusion, ...

- Basic machine (datapath, memory, IO devices) are application agnostic
- Same concepts / processor architecture apply to large variety of applications. E.g.
 - OS with command line and graphical interface (Linux, ...)
 - Embedded processor in network switch, car engine control, ...
- Input / output (I/O)
 - Memory mapped: appears like “special kind of memory”
 - Access with usual load/store instructions (e.g. **lw**, **sw**)
- Exceptions
 - Notify processor of special events, e.g. overflow, page fault (next lecture)
 - “precise” handling: immediately at offending instruction
- Interrupts
 - Notification of external events, e.g. keyboard input or Ethernet traffic
- Multiprogramming and supervisory mode
 - Enables and isolates multiple programs
- Take CS162 to learn more about operating systems